



UNIVERSITÄT ZU LÜBECK
INSTITUTE FOR SOFTWARE ENGINEERING
AND PROGRAMMING LANGUAGES

An Asynchronous Evaluation Engine for Stream Based Specifications

Asynchrone Evaluierung von Strombasierten Spezifikation

Masterarbeit

im Rahmen des Studiengangs

Informatik

der Universität zu Lübeck

vorgelegt von

Alexander Schramm

*ausgegeben und
betreut von*

Prof. Dr. Martin Leucker

mit Unterstützung von

Cesar Sanchez, Ph.D.

Die Arbeit ist im Rahmen einer Tätigkeit beim IMDEA Software Institute entstanden.

Lübeck, den 20. November 2016

Alexander Schramm

An Asynchronous Evaluation Engine for Stream Based Specifications

Masterarbeit, 20. November 2016

Reviewers: Prof. Dr. Martin Leucker

Supervisors: Cesar Sanchez, Ph.D.

University of Luebeck

Institute For Software Engineering and Programming Languages

Ratzeburger Allee 160

23562 Luebeck

Abstract

Abstract (Deutsch)

Acknowledgement

Contents

1	Introduction	1
1.1	Motivation and Problem Statement	1
1.2	Results	3
1.3	Thesis Structure	4
2	Related Work	5
2.1	RV Techniques for C Programs	5
2.1.1	Copilot	6
2.1.2	RMOR	7
2.2	Distributed Verification Techniques	7
2.3	Stream Based Specification Techniques	8
2.3.1	LOLA	8
2.3.2	BeepBeep 3	10
2.4	TESSLA	10
2.5	Trace Data	12
2.5.1	General Benchmarks for RV Tools	13
2.5.2	Tools to Generate Traces	14
2.5.3	CIL	14
2.5.4	Google XRay	14
2.5.5	DTrace	15
2.5.6	LLVM	16
3	Preliminaries	19
3.1	Time	19
3.2	Transducers	19
3.3	Timed Transducers	22
3.4	Events	27
3.5	Streams	27
3.6	Functions	28
3.7	Nodes	29
3.8	TESSLA Evaluation Engine	29
3.9	TESSLA Functions	30
3.9.1	Complete Functions	31
3.9.2	Output Complete Functions	31

3.9.3	Input Complete Functions	31
3.9.4	Incomplete Functions	32
3.9.5	Timing Functions	35
3.10	State and History	36
3.11	Transitions	37
3.12	Run	38
4	Behaviours of Evaluation Engines	45
4.1	Schedules Without Timing Functions	46
4.1.1	Greedy Evaluation Engines	47
4.1.2	Fair Evaluation Engines	49
4.2	Equivalence of Different Schedules Without Timing Functions	50
4.2.1	Equivalence of Greedy Systems	50
4.2.2	Equivalence of Greedy and Fair Evaluation Engines	55
4.3	Timing functions	56
5	Implementation Details	59
5.1	TesslaServer	59
5.1.1	Erlang and Elixir	59
5.1.2	Architecture	61
5.1.3	Synthesis of the Evaluation Engine	62
5.1.4	Node Implementation	63
5.1.5	TesslaServer V1: Stream passing	65
5.1.6	TesslaServer V2: Event passing	67
5.2	Instrumentation Pass	68
6	Evaluation	71
6.1	Runtime Benchmarks	71
6.1.1	Number of Processors	71
6.1.2	Number of Events	73
6.1.3	Number of Nodes	74
6.2	Instrumentation Benchmarks	77
6.2.1	Performance Comparison with non Instrumented Code and Compiler Optimizations	78
6.2.2	Performance Impact in Regard to Instrumentation Percentage	79
6.3	Practical Examples	79
7	Conclusion	81
7.1	TesslaServer	81
7.2	Instrumentation Pass	83
7.3	Further Work	84
7.3.1	Implementation for the JVM	85
7.3.2	Composition of Evaluation Engines	85

7.3.3	Parameterized Streams	86
7.3.4	Easier Definition of New Node Types	86
7.3.5	Runtime Optimizations	87
7.3.6	Error prevention	88
8	Appendix	89
8.1	Runtime Benchmark Data	89
8.1.1	Execution Time in Regard to Used Processor Cores	89
8.1.2	Execution Time in Regard to Number of Input Events	91
8.1.3	Ram Usage with Respect to Number of Input Events	93
8.1.4	Execution Time in Regard to Number of Nodes	95
8.2	Ringbuffer Code	97
8.3	Instrumentation Benchmark Data	100
8.4	Example TESSLA Specifications	105
	Bibliography	107
	Glossary	115

Introduction

In this Chapter we will look at the challenges that motivate the works of this thesis and how the results solve these challenges.

1.1 Motivation and Problem Statement

Software verification is an important tool to harden critical systems against faults and exploits. Due to the raising importance of computer based systems, verification has become a big field of research in computer science.

While pure verification approaches try to proof the correct behaviour of a system under all possible executions, Runtime Verification (RV) limits itself to single, finite runs of a system. The goal is to proof that a run conforms to a given specification under specific conditions, like input sequences or scheduling. Specifications can be given in various ways, including Temporal Logic (TL) formulas or in specification languages that are specifically developed for RV. Examples for this are RMOR [Hav08], LOLA [DAn+05] and others [Zhe+15; Pik+10; MB15], which we will look at more closely in Chapter 2.

The project TESSLA[DTH16] presents ways to specify and evaluate properties over streams of events including timing information. To achieve this it introduces a language to expressively describe the conditions one or more streams should fulfill by applying transformations on them. The evaluation of TESSLA a specification is done in two steps: first the specification is compiled by a compiler written at the Institute for Software Engineering and Programming Languages (ISP) of the University of Lübeck. The output is a canonical representation of the transformations on the streams in the specification. In the second step the compiled specification is connected with a system that produces traces that are treated as the input streams of the specification.

The second step can be done in different ways: online or offline, interweaving the monitors into the monitored program (like for example done in [Hav08]) or by executing them standalone. These different approaches lead to different ways the monitored program has to be altered, for example manipulating its original code to log status informations or to invoke the monitoring code.

Interweaved monitors can alter the original system and produce new errors or even suppress others. Standalone monitors on the other hand will have a much smaller impact on the monitored system. But as a consequence there will be a bigger delay between the occurrence of events in the program and their evaluation in the monitor. Furthermore interweaved monitors can optionally react to detected errors. They could change the control flow of the original system or alert a third party and eliminate cascading errors. Standalone monitors can't directly modify the program but can still produce warnings and alerts that can then be reacted to.

While online monitoring can be used to actively react to error conditions, either automatically or by notification of a third party, offline monitoring can be thought of as an extension to software testing [DAn+05].

At the beginning of this thesis there was one implementation of a runtime for TESSLA specifications that is based on Field Programmable Gate Arrays (FPGAs) that have to be manually reconfigured for each new specification. While this is a very performant approach for actual monitoring it isn't feasible for testing and prototyping. This leads to the wish to implement a software based TESSLA runtime which can be executed independent of hardware restrictions.

Furthermore most RV approaches are specific to one programming language or environment and combine ways of generating the data, which is used for monitoring, and the monitoring itself. TESSLA specifications themselves are independent of any implementation details of the monitored system, working only on streams of data, which can be gathered in any way. This can be used to implement a runtime that is also independent of the monitored system and how traces of it are collected.

During the thesis it is proven that the actual approach of the runtime, a functional, actor based, asynchronous system, will generate the same observations on input traces as a synchronous evaluation of the specification. While TESSLA specifications can work on all kinds of streams, especially on traces on all levels of a program, including instruction counters or spawning processes, in this thesis we will mainly focus on the level of function calls and variable reads/writes. Other applications of the system could easily extend it to use traces representing drastically different fields, for example health data, temperatures, battery levels, web services and more.

To test the software based runtime different specifications will be tested on multiple traces. Some of the traces are generated by actually running a program which was instrumented by hand or automatically to generate traces. Others were generated or modified by hand to deliberately introduce bugs which should be detected by the system.

1.2 Results

The main contributions of this thesis consist of three parts. The first is a theoretical approach to asynchronously evaluate timed specifications over streams. The second is an implementation that can synthesize systems to evaluate such specifications based on the theoretical approach. And the third is a proof of concept implementation of a system that can instrument code which is compiled with Low Level Virtual Machine (LLVM), mainly targeted at C and C++.

The theoretical evaluation approach aims to solve the challenges of RV in a way that facilitates the current state of computation like parallelism and distribution. While this is sensible to enable the efficient usage of resources in an implementation, it introduces new challenges which needs to be tackled. Therefore a theoretic basis is introduced which enables the reasoning about correctness in the context of the asynchronous evaluation approach.

The implementation of the system to synthesize evaluation engines for specifications is an attempt to translate the theoretical evaluation approach directly into software. To do so the Erlang platform was chosen, which provides abstractions that can be used to implement an asynchronous and distributed system in a straightforward way. The implementation can evaluate a multitude of specifications written in TESSLA, scales well with the size of the specification, the number of events in a trace and the number of cores of the hardware it is executed with. The whole system implemented in a modular way, enabling a user to switch some parts, for example the part of the system that parses a TESSLA specification could be exchanged with one that parses a specification in another language and the rest of the system would not be affected. Finally the implementation makes it easy to implement new functions, making the system open to extensions of TESSLA.

To generate test data for the implementation the third part was developed. For testing purposes a proof of concept implementation was sufficient. This proof of concept was successfully used to instrument multiple programs written in C and one in Swift. Instrumented programs emit trace data about calls of specified functions without any manual work required. At this point the instrumentation is only able to instrument function calls and isn't optimized in any way, leading to a big performance impact if it is used excessively. Nonetheless it shows great potential of the underlying system that it uses to perform the instrumentation and is therefore an interesting and important part of this thesis.

1.3 Thesis Structure

As the whole evaluation engine is built on top of different technical and theoretical ideas, it is structured to show the reasoning behind the decisions that were made during the development. Furthermore it will proof equalitys of different kinds of systems in multiple steps that build on one another. In the following a quick overview of the different parts of the thesis is given.

Chapter 2

In this chapter the theoretical foundation for the system is explained. Furthermore multiple approaches solving similar problems are shown and it is highlighted which concepts of them were used in the new system and which were disregarded and why.

Chapter 3

Building on the theoretical and practical findings of the previous chapters new definitions are presented, which are needed to reason about the implemented system.

Chapter 4

The work from the previous chapter is put to work to reason about the semantics of the implemented runtime and to show its correctness.

Chapter 5

This chapter highlights technical details of the implementation. It will present alternative implementation approaches and the reasoning why specific choices were made during the development of the system.

Chapter 6

To show the value of the implemented system it is thoroughly tested and benchmarked with fabricated and real world examples and traces. The results of this testing is used to evaluate the implementation.

Chapter 7

On the basis of the evaluation in the conclusion the results of the thesis are summarized. Furthermore it is highlighted what remains to do and which future challenges exist.

Related Work

As Runtime Monitoring and Verification is a widely researched field, multiple approaches have been developed and new approaches are presented all the time.

As stated in [Hav08] most approaches are geared towards software written in Java, while many critical systems are written in C. Additionally there are countless other systems that could benefit from monitoring and verification written in other programming languages. TESSLA is a specification language over streams, which has no assumptions on the environment of the system that produces the streams. Using TESSLA as the base for our monitoring approach, we recognized the possibility to decouple the monitoring platform from the monitored program. This means that the runtime developed for TESSLA is not restricted to monitor programs written in a specific language but can monitor anything that can produce streams of data.

To show that the runtime is valuable in the context of existing approaches we will look at ways to generate traces from systems during their execution. Based on the generated traces we benchmarked out runtime to evaluate how it scales with respect to different characteristics of specifications and the platform it is run on.

The following chapter will summarize the approaches that are available in the field of RV and how they influenced TESSLA and the runtime we implemented.

2.1 RV Techniques for C Programs

Most RV approaches are developed for Java and the Java Virtual Machine (JVM) and therefore they are not usable for a large class of software not written in Java. In particular C and C++ are used in a wide variety of safety critical systems, including aircrafts and energy plants. While TESSLA as a language and the implemented runtime is independent of the platform of the monitored program, the domain of C programs and especially embedded software has a special focus in this thesis. Therefore in this section we will look at some existing approaches to monitor C programs and in Section 2.5 we will discuss ways to generate trace data from such programs.

2.1.1 Copilot

The realtime runtime monitor system Copilot was introduced in [Pik+10]. Copilot is designed to overcome the shortcomings of existing RV tools in regards to hard-realtime software written in C.

To accomplish this goal Copilot defines characteristics that a monitoring approach has to fulfill to be considered valuable for this domain. The four principles are:

Functionality Monitors cannot change the functionality of the observed program unless a failure is observed.

Schedulability Monitors cannot alter the schedule of the observed program.

Certifiability Monitors must minimize the difficulty in re-validating the observed program. In particular, monitoring must be accomplished without modifying the observed programs source code.

SWaP overhead Monitors must minimize the additional overhead required including size, weight, and power (SWaP).

The monitors follow a sampling based approach, where at specified steps the values of global variables are observed and the monitors are evaluated on the observed values. While sampling based approaches are widely disregarded in RV because they can lead to both false positives and false negatives, [Pik+10] argues:

In a hard real-time context, sampling is a suitable strategy. Under the assumption that the monitor and the observed program share a global clock and a static periodic schedule, while false positives are possible, false negatives are not.

A special detail of Copilot is that monitors are not inlined into the program but can be scheduled as independent processes. The implementation of the TESSLA runtime in this thesis follows a similar approach. Our runtime is a totally independent program and therefore also enjoys some of the gains with respect to the specified four characteristics. Because the runtime works with all kinds of traces, it is insignificant how these traces are produced. Our runtime can work with traces based on sampling, working in a similar fashion as Copilot, or by actually instrumenting code to generate traces, which might alter the semantics of the monitored program.

2.1.2 RMOR

RMOR [Hav08] is another approach for monitoring C programs. RMOR transforms C code into an *armored* version, which includes monitors to check conformance to a specification.

Specifications are given as a textual representation of state machines which is strongly influenced by RCAT [SH08]. The specifications are then interweaved into the program using C Intermediate Language (CIL) [Nec+02].

The specifications considered in RMOR work on the level of function calls and state properties like *write may never be called before open was called*. Because software developers are often working at the same abstraction level (in contrast to assembler or machine instructions), they can define specifications without having to learn new concepts. The TESSLA runtime supports the definition of traces at the same abstraction level (function calls, variable reads and writes). This is used in most of the tests in Section 6.3.

Because RMOR specifications are interweaved into the program, their observations can not only be reported but also used to recover the program or even to prevent errors by calling specified functions when a critical condition is encountered. The TESSLA runtime does not support this functionality out of the box as its primary purpose is testing and offline monitoring. In Section 7.3.6 we will look at possible extensions to enable communication from the monitor back to the monitored program.

2.2 Distributed Verification Techniques

Many works in the field of RV do not consider parallelism and distributed systems. There are two aspects to this challenge: monitoring programs that are run in a distributed fashion on the one hand, and using parallelism and distribution to implement monitors on the other.

Monitoring distributed programs carries an inherent problem: events are no longer globally ordered. This led to algorithms like Lamport timestamps [Lam78] and vector clocks [Fid88]. The distribution of the monitored system, and therefore of the events, must be expressible in the language that is used to write specifications about the system. Two examples for such a specification language are presented in [Sen+04] and [EC00]. [MB15] presents a way to monitor distributed programs that uses Linear Temporal Logic (LTL) as the specification language, but requires the presence of a global state which is constructed using Lamport timestamps. [MB15]

presents a method to implement distributed monitors that have to communicate with one another.

While the TESSLA specification language does not include mechanisms to explicitly reason about distributed properties, the TESSLA runtime does not care about the environment of the monitored program, so it does not distinguish between traces from distributed and non distributed programs. As we will see in Section 5.1 this means that TESSLA can be adopted to monitor at least some characteristics of a distributed system. But more importantly, the runtime takes parallelism and distribution as one of its core concepts as we will explain in the next chapters.

The other challenge of parallelism and distribution is how these features can be used to implement monitoring algorithms. Since modern systems often contain many processors or cores it is important to study how to exploit parallelism when building monitors, especially when performance of the monitor is important.

The work in [AF16] uses the Erlang platform to implement a highly distributed monitoring algorithm for a branching time logic called Monitorable Hennessy-Milner Logic (MHML). The approach works by synthesizing formulas into *submonitors* which can then be run in parallel. The verdicts from these monitors are combined to form a global verdict. For example the formula $\Phi \vee \Psi$ can be synthesized as two *submonitors*, one for Ψ and one for Φ .

The TESSLA runtime follows a similar approach where each part of the specification is implemented as an independent actor. These actors can then be scheduled by the environment in a parallel fashion. In Section 5.1 we will look at this architecture more closely.

2.3 Stream Based Specification Techniques

Specifications in the field of RV are defined over finite words, meaning a finite sequence of events. But some works take another perspective on the problem and look at the sequence of properties as streams of data. As TESSLA takes this approach we will look at other work that promoted this view and how it can be used to incorporate techniques from other fields into RV.

2.3.1 LOLA

The concepts of LOLA [DAn+05] are very similar to the ones of TESSLA. The biggest difference between these two approaches is that streams in LOLA are based on a

```
s = s[-1, 0] + ite((a ∧ ¬b), 1, 0) + ite((b ∧ ¬a), 1, 0)
trigger(s ≤ 0)
```

Listing 2.1: A LOLA specification describing the property that the number of *a*'s in a stream shall never be less than the number of *b*'s

discrete model of time while TESSLA uses a continuous timing model highlighted in Section 2.4.

The specification language of LOLA is very small (expressions are built upon three operators) but the expressiveness surpasses Temporal Logics and many other formalisms for finite traces [DAn+05]. Expressions in LOLA are built by defining streams from other streams. Therefore streams depend on other streams and they can be arranged in a weighted dependency graph, where the weight describes the amount of steps a generated stream is delayed compared to the stream in its definition. In contrast to TESSLA streams in LOLA can depend on themselves and therefore the dependency graph can contain cycles. Based on this graph a notion of efficiently monitorable properties is given in [DAn+05] which also presents a monitoring algorithm.

Listing 2.1 shows a small specification written in LOLA. This example shows the usage of integer streams, which allows LOLA the use of counters and therefore enables it to express context-free properties.

LOLA was recently extended in [Fay+16] with two new features: template stream expressions, which allow input data to be carried along the stream, and dynamic stream generation, where new monitors can be invoked during the monitoring process for the monitoring of new subtasks on their own time scale. These extensions add the concept of *slicing* with template streams, which enables the generation of parameterized streams. A parameterized stream, also called template stream, is a mechanism to generate dynamically concrete streams from one template stream by *slicing* it with respect to some data. For example, a template stream could be used to monitor logins of individual users, where one *slice* is generated for each user. Adding parameterization to TESSLA was discussed but is not currently implemented yet.

TESSLA takes concepts of LOLA and applies them to a continuous model of time and introduces a language and a rich set of functions that can be applied to streams. The dependency graph is a core concept of TESSLA and is used to check if specifications are valid and is also the core concept to evaluate specifications over traces in this thesis.

2.3.2 BeepBeep 3

The project BeepBeep aims to introduce concepts from the field of Complex Event Processign (CEP) to RV. The third version of BeepBeep is presented in [Hal16] where different concepts from CEP are applied to problems of RV. BeepBeep implements a system for evaluating specifications over synchronous event streams. It exposes two ways to express specifications: an Application Programming Interface (API) in Java and the language Event Stream Query Language (eSQL).

The evaluation model of BeepBeep is very similiar to the one implemented for TESSLA in this thesis: events enter the system through sources, flow through multiple processors which perform transformations and in turn can lead to the generation of an output. The central unit of computation are called processors: processors transform multiple input streams into multiple output streams. A processor can only perform a transformation when there is at least one event buffered at each input stream. The combination of the first event on each buffer is called the *front*. Whenever a front can be formed the processor is invoked to perform its work and generate new outputs based on the front We will reuse this terminology in Section 5.1 when we talk about the way computations in TESSLA work.

2.4 TESSLA

The implemented runtime and the theoretic work of this thesis is built upon the TESSLA project from [DTH16], that defines a syntax and a formal semantic of the TESSLA specification language.

Specifications in TESSLA are based on streams of data. Streams are the representation of data over time, for example the value of a variable in a program or the temperature of a processor. To model streams TESSLA defines a timing model. That model is based on timestamps that are isomorphic to real numbers \mathbb{R} . Figure 2.1 illustrates how streams behave over time.

The syntax of TESSLA is concise, but can be used to define complex functions and specifications:

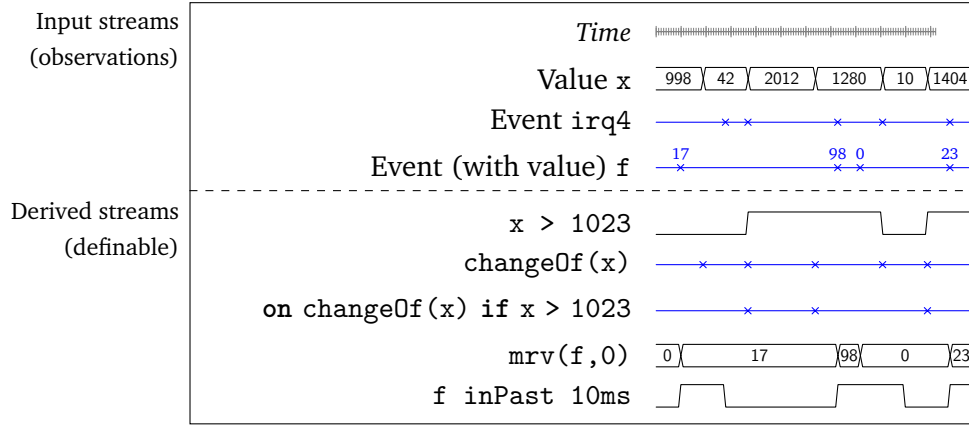


Fig. 2.1: Visualization of TeSSLA stream model, taken from [DTH16]

```

spec ::= define name[: stype] := texpr
      out texprspec spec
texpr := expr[: type]
expr  := name | literal | name(texpr(, texpr)*)
type  := btype | stype
stype := Signal<btype> | Events<btype>

```

One of the main contributions of TeSSLA is the syntax which mimics modern programming languages and diverges from more classical approaches in RV that use more formal specification languages typically based on logics or automata. This is an important step to enable practitioners without a strong theoretical background to adopt RV techniques in their workflow. While RV has a lot of mechanisms to express specifications they often lack the ability to be intuitively understood.

There are many specification languages like LTL [Pnu77], Regular Linear Temporal Logic (RLTL) [LS07], Computational Tree Logic (CTL) [CE82] and many others that are geared heavily toward scientific work and theoretic reasoning. When introducing new concepts, like realtime, this trend keeps up and formalism like Timed Linear Temporal Logic (TLTL) from [RS97], Signal Temporal Logic (STL) from [MN04] and Metric Temporal Logic (MTL) from [Koy90] provide theoretical foundations to reason about realtime properties but formulas using these logics are even harder to understand than their non realtime counterparts.

One approach to make RV more usage friendly is Structured Assertion Language for Temporal Logic (SALT) presented in [BLS06] which acts as a frontend language to the more formal specification languages and can be translated into them. SALT unifies many different mechanisms, like specification patterns, nested scopes, excep-

```

1 define max_twice_at_floor_before_open(i) := always (
    occurring[<=2] atfloor_$i$ between inclusive optional
    call_$i$ , exclusive optional open_$i$)
define max_60s_before_open(i) := always (call_$i$ implies
    eventually timed[<=60.0] open_$i$)
assert allof enumerate[1..3] as floor in
    max_twice_at_floor_before_open(floor) and
    max_60s_before_open(floor)

```

Listing 2.2: An example specification in the SALT language taken from [BLS06] defining behaviour of an elevator.

tions, regular expressions and realtime. Listing 2.2 shows an example specification in SALT taken and adapted from [DAC99] which specifies that on all three floors in a building, calling the elevator at floor i implies that it may pass at most two times at that floor without opening its doors, and that it must finally open its doors at that floor within 60 seconds.

This specification shows, how SALT specifications are more intuitively understandable than logic formulas, for example by allowing to split the formula in multiple parts and assign meaningful identifiers to subformulas.

TESSLA and the runtime implemented in this thesis aim to combine many of the aspects that were presented in this Chapter: An understandable specification language like SALT that is able to express realtime properties, using streams of data as a central element like LOLA, incorporate techniques from CEP like BeepBeep, and distribute the monitoring to many processors using Erlang and the actor model as shown in [AF16].

2.5 Trace Data

As in RV we talk about monitoring of properties over traces, another important aspect is how traces can be extracted from a running program. Many RV tools solve this problem by using a technique called *interweaving*: The code of the monitored program is changed in a way that the monitor becomes part of it, therefore the monitor becomes part of the program and has explicit access to the state of the running system. Examples for this approach are AspectJ [Kic+01] and DiSL [MZA12] for Java, and RMOR [Hav08] and Runtime Time-triggered Heterogeneous Monitoring (RiTHM) [Nav+13] for C. This approach is not feasible for the goals of the TESSLA runtime, because we seek the portability that enables users to monitor multiple kinds of systems.

Therefore we are looking for ways to manipulate programs in a general way to emit trace data.

As a first step we sought other projects to extract traces of data that can be used to evaluate the implemented runtime. While there are many benchmarks available to test monitoring tools we were not able to find any that satisfied all of the characteristics that are needed for an evaluation of TESSLA. The next sections list some of the benchmarks and evaluation tools that were surveyed and concludes with a tool that enabled the generation of suitable trace data.

2.5.1 General Benchmarks for RV Tools

The field of RV lacks a common specification language and can be expanded to a wide variety of different properties to verify. Therefore it is no surprise that there is no common benchmark that is applicable to all tools. Nonetheless some benchmarks are occasionally used to compare the expressiveness of a new RV tool.

The DaCapo benchmark [Bla+06] was introduced as a general purpose benchmark for the Java language. It turned out that it can be used to benchmark monitor implementations and trace collection tools [Wu+16; Che07; MZA12]. Since Java was not the target architecture of our runtime the DaCapo benchmark was not used for the evaluation in this thesis.

The works in [DAC99] can be seen as a benchmark for the expressiveness of a specification language. It categorizes commonly monitored properties into so called patterns like *precedence*, *absence* and *response*. While it would be valuable to test TESSLA and the implemented runtimes against these patterns this would be a two step process. First each pattern would have to be checked, if and how it could be expressed as a TESSLA specification. As a second step, if the pattern can be expressed, test trace data would have to be generated and only then the runtime could be evaluated using that pattern. Since this thesis does not work on the TESSLA language definition but only on the runtime we decided not to pursue this goal.

Another interesting project for evaluation of offline monitoring tools is TraceBench presented in [Zho+14]. TraceBench is a big set of traces collected from a distributed system running Hadoop Distributed Filesystem (HDFS). During the collection errors of different categories were deliberately introduced, like network timeouts or data corruption. The generated traces are organized in a hierarchy: if an event is produced as the effect of another event, for example a function call that leads to another function call, the second event is a child of the first. Furthermore the traces contain the beginning and end of each event as a timestamp, making the traces suitable for TESSLA. Unfortunately it seems that the generated trace data has a problem: since events are produced in a distributed system the timestamps

are faulty, for example some events started before their parent event. Furthermore TESSLA doesn't have a mechanism to express nested events as of now, meaning the traces would need to be manipulated before they could be used.

As a final benchmark we evaluated the Competition on Runtime Verification (CRV) [RHF16]. CRV has collected a set of benchmarks especially developed for the usage in monitoring algorithms. It consists of three tracks: Java, C and offline. While the C and offline track would have been very interesting the benchmarks don't contain any realtime properties. Since the realtime fragment is an important part of TESSLA we decided that the benchmark wasn't appropriate either.

2.5.2 Tools to Generate Traces

Since no suitable benchmark could be found the next step was to search for a tool that can be used to generate appropriate traces from programs during execution. These tools can be categorized into static and dynamic instrumentation tools: static tools transform the source code of a program during compilation to include logging statements, dynamic ones only work at runtime and are not involved in the compilation process.

2.5.3 CIL

CIL [Nec+02] is a tool to write source-to-source transformations for C programs, therefore it is a static instrumentation tool. CIL implements "a highly-structured clean subset of C" in the OCaml language. CIL transforms a C program into an intermediate representation in OCaml. CIL will then apply transformations that are supplied by the programmer to the intermediate representation. When all transformations are applied the program is written back as a normal C program.

Since CIL is able to represent the complete C90 standard and also extensions that are commonly used like the ones from GNU C, this tool can be used to write an instrumentation pass for the use case of trace generation. The main reason that CIL was not used is that CIL only supports C and in this thesis we seek a tool that could be used on a variety of programming languages.

2.5.4 Google XRay

Google XRay [Ber+16] is a function call tracing system for C and C++. It is mainly a static instrumentation tool but has some aspects of a dynamic tool. While XRay requires the original source code to be instrumented, the tracing functionality can

be turned off at runtime which minimizes overhead. XRay works by inserting a series of no-ops after function entry points and before return points. At runtime a library that is part of XRay can then patch these no-ops with instructions to call a log function if tracing is enabled. This rather complex mechanism is chosen to enable a minimal overhead which was a main requirement for the development of XRay.

XRay is at the moment implemented as a set of patches onto GNU Compiler Collection (GCC) but is planned to be migrated to LLVM. In typical usage scenarios XRay leads to an overhead of around 20% to 40%.

While XRay looks like a promising all-in-one solution to trace collection it only supports the instrumentation of function entries and exits. For our instrumentation we wanted to explore the possibility to also trace other events and maybe even allow to attach conditions when an event is actually logged. As an example consider logging an event whenever a variable is assigned more than once in a single function call.

2.5.5 DTrace

[CSL04] presents DTrace, a “facility for dynamic instrumentation of production systems”. DTrace relies on support from the Kernel of the operating system and does require specific compilation settings for some instrumentations. DTrace includes a language to specify instrumentations called D.

To use DTrace one has to specify *probes* in the D language. A *probe* consists of probe descriptions, a predicate and action statements. The probe descriptions specifies the kind of instrumentation it provides and the modules and functions the probe should be applied to. The predicate can be used to skip the invocation of the probe if certain conditions are not met. Action statements contain code that is executed when the probe is called.

Listing 2.3 shows a DTrace specification with two probes. The first probe is invoked everytime the instrumented program calls the `read` function provided by the operating system and assigns the current timestamp to a variable. The second probe is invoked everytime the `read` function returns and logs the time used by the `read` function before returning. The predicate of the second probe checks if the variable `t` has been assigned.

DTrace offers a variety of instrumentation types, called *providers* in DTrace, `syscall` from the previous example being one of them. Interesting for our research is the *pid* provider which enables the instrumentation of arbitrary instructions in a specified

```

syscall::read:entry {
    self->t = timestamp;
3 }

syscall::read:return
6 /self->t/ {
    printf("%d/%d spent %d nsecs in read\n",
        pid, tid, timestamp - self->t);
9 }

```

Listing 2.3: A DTrace specification in the D language specifying two probes

process. The *pid* provider can be used to log function entries and exits of every function in a program.

[RSS16] builds upon DTrace to build a RV framework. It produces automaton based monitors in the D language from specifications written in LTL. To do so it connects atomic propositions of the formula with observable events that are represented as probes in DTrace. Whenever a probe is invoked it updates the state of the automaton as specified with the action statements of that probe.

2.5.6 LLVM

LLVM [LA04] is a compiler framework that allows program analysis and transformation. The compilation process of LLVM is separated into three parts: a front-end, a middle-end and a back-end. A front-end is responsible for translating a source language, for example C or C++, into LLVM Intermediate Representation (IR), a strongly typed Reduced Instruction Set Computing (RISC) instruction set which is independent from the target platform. The middle-end performs source-to-source transformations on the IR that can perform analysis or optimizations. Such transformations are called *compiler passes* and are independent of the source language. The back-end is then used to transform the IR into native machine code for the target platform.

The *compiler pass* on top of IR provides a good abstraction as a base for an instrumentation to generate traces. Since a pass works on IR and not on the source language itself the instrumentation can be used for every language that has an LLVM front-end. At the time of writing there is a large collection of such frontends for many languages, including C, C++, Objective-C, Swift, Haskell, Ruby and many more¹.

¹<http://llvm.org/Projects/WithLLVM/>

Furthermore a compiler pass can work on all parts of a program: whole modules (think classes in C), function definitions, variable reads and writes, memory allocation and others. The building block of IR programs in LLVM are the so called *instructions*². Each statement from a source language is represented as such an instruction. A *compiler pass* is able to examine instructions, change or reorder them and generate new instructions and insert them appropriately.

Due to the great flexibility that a *compiler pass* offers we chose this approach for the instrumentation pass. The implementation details can be seen in Section 5.2.

²http://llvm.org/docs/doxygen/html/classllvm_1_1Instruction.html

Preliminaries

In this chapter we will define concepts that are used in Chapter 4 to reason about the implemented runtime.

While the TESSLA specification itself defines a set of semantics, for this thesis we will slightly alter some of it and add some new definitions based on them. This is necessary to reason about the specifics how the runtime is built (Note that TESSLA doesn't define an operational semantic, therefore we will define our own) and how it behaves.

3.1 Time

TESSLA has a model of continuous time, where timestamps $\pi \in \mathbb{T}$ are used to represent a certain point in time and \mathbb{T} has to be isomorphic to \mathbb{R} .

3.2 Transducers

Fundamentally TESSLA is a special kind of a transducer. Therefore in this section we will define a model of transducers which can be used to reason about the evaluation of a TESSLA specification.

A transducer is a system, which consumes an input and produces an output. Let Φ, Γ be two alphabets and ϵ the empty word.

Definition 1: Transducer.

A transducer t is a relation $t \subseteq \Phi^ \times \Gamma^*$, Φ is called the input alphabet, Γ the output alphabet.*

TESSLA specifications are deterministic for any input, meaning they should produce the same output for the same input.

Definition 2: Deterministic Transducer.

A deterministic transducer relates each input to at most one output.

Example 1: Deterministic and Nondeterministic Transducers.

$t_d = \{(a, 1), (b, 2), (ab, 12), (ba, 21)\}$ is a deterministic transducer; $t_{nd} = \{(a, 1), (a, 2)\}$ is nondeterministic, because it relates input a to both outputs 1 and 2.

Transducers can furthermore be categorized as synchronous, asynchronous, causal and clairvoyant transducers: synchronosity is a property over the behaviour of a transducer when it's consuming input per element. If it is synchronous, it will produce an output element for each input element.

Definition 3: Synchronous Transducer.

Let $\vec{i} \in \Phi^*$, $i \in \Phi$, $\vec{o} \in \Gamma^*$, $o \in \Gamma$. A transducer t is called synchronous, when it satisfies, that: if $(\vec{i} \circ i, \vec{o} \circ o) \in t$ then $(\vec{i}, \vec{o}) \in t$

An asynchronous transducer can produce zero, one or many outputs for each input it consumes.

Definition 4: Asynchronous Transducer.

Let $\vec{i} \in \Phi^*$, $i \in \Phi$, $\vec{o} \in \Gamma^*$. A transducer t is called asynchronous when it satisfies the formula: if $(\vec{i} \circ i, \vec{o}) \in t$ then $\exists \vec{o}', \vec{o}'' \in \Gamma^*$ so that $\vec{o} = \vec{o}' \circ \vec{o}''$ and $(\vec{i}, \vec{o}') \in t$

Example 2: Synchronous and Asynchronous Transducers.

$t_s = \{(a, 1), (b, 2), (ab, 12), (ba, 21)\}$ is a synchronous transducer, $t_{as} = \{(a, \epsilon), (ab, 12)\}$ is asynchronous.

A causal transducer is one, where the output depends only on consumed inputs and not on future inputs:

Definition 5: Causal and Clairvoyant Transducers.

A transducer t is called causal, when it satisfies, that: if $(\vec{i}, \vec{o}) \in t$ then $\forall \vec{i}' \in \Phi^*$ with $(\vec{i} \circ \vec{i}', \vec{o}) \in t$ it holds, that $\vec{o} \sqsubseteq \vec{o}'$

A transducer that isn't causal is called clairvoyant.

Example 3: Causal and Clairvoyant Transducers.

$t_{cl} = \{(a, 1), (b, 2), (ab, 12), (ba, 21)\}$ is a causal transducer, because each output only depends on the inputs seen upto that point, $t_{cl} = \{(a, 1), (ab, 22), (aa, 11)\}$ is clairvoyant, because the output when the letter a is seen depends on the next input.

When talking about transducers, it is interesting to know if two transducers are equivalent. There are multiple possible definitions for equivalence of transducers, we will look at two, which are interesting for this thesis. In the following σ_i is used to get the element at position i and $\sigma_{[i,j]}$ to get the infix of σ which starts at position i and ends at position j (With 0 as the index of the first element).

Definition 6: Asynchronous equivalence of Transducers.

Let t_1, t_2 be two asynchronous transducers from Φ^* to Γ^* . They are called asynchronous equivalent, written $t_1 \equiv_a t_2$, if they satisfy:

$\forall \sigma \in \Phi^*$:

- $\forall (\sigma_{[0,k]}, \vec{o}) \in t_1: \exists k' \geq k$ with $(\sigma_{[0,k']}, \vec{o}) \in t_2$ and $\vec{o} \sqsubseteq \vec{o}'$
- and $\forall (\sigma_{[0,k]}, \vec{o}) \in t_2: \exists k' \geq k$ with $(\sigma_{[0,k']}, \vec{o}) \in t_1$ and $\vec{o} \sqsubseteq \vec{o}'$

Lemma 1: Asynchronous equivalence is an equivalence Relation.

Asynchronous equivalence is symmetric, reflexive and transitive.

Proof. Symmetry is trivial, since the second part of the definition is requiring it.

Reflexivity is also trivial, for $(\sigma_{[0,k]}, \vec{o})$ select $k' = k$.

For transitivity:

Let $t_1 \equiv_a t_2, t_2 \equiv_a t_3$.

First case:

Since $t_1 \equiv_a t_2 : \forall (\sigma_{[0,k_1]}, \vec{o}_1) \in t_1 :$

$\exists k_2$ such, that $(\sigma_{[0,k_2]}, \vec{o}_2) \in t_2$ with $\vec{o}_1 \sqsubseteq \vec{o}_2$

and since $t_2 \equiv_a t_3$

$\exists k_3$ such, that $(\sigma_{[0,k_3]}, \vec{o}_3) \in t_3$ with $\vec{o}_2 \sqsubseteq \vec{o}_3$

With $\vec{o}_1 \sqsubseteq \vec{o}_2 \sqsubseteq \vec{o}_3$ it follows, that $t_1 \equiv_a t_3$

The second case works the same, just change t_1 and t_3 .

□

Example 4: Asynchronous equivalence of Transducers.

Let $\Phi = \{a\}, \Gamma = \{1\}$ and

$t_1 = \{$	$(a, \epsilon),$	$(aa, \epsilon),$	$(aaa, 111)$	$\}$
$t_2 = \{$	$(a, 1),$	$(aa, 1),$	$(aaa, 111)$	$\}$
$t_3 = \{$	$(a, \epsilon),$	$(aa, 1),$	$(aaa, 11)$	$\}$

All three transducers are asynchronous and causal. Let's see which ones are asynchronous equivalent:

$t_1 \stackrel{?}{\equiv}_a t_2$

(a, ϵ)	$\in t_1, k = 1$	$\rightarrow k' = 1, (a, 1) \in t_2,$	$\epsilon \sqsubseteq 1$
(aa, ϵ)	$\in t_1, k = 2$	$\rightarrow k' = 2, (aa, 1) \in t_2,$	$\epsilon \sqsubseteq 1$
$(aaa, 111)$	$\in t_1, k = 3$	$\rightarrow k' = 3, (aaa, 111) \in t_2,$	$111 \sqsubseteq 111$
$(a, 1)$	$\in t_2, k = 1$	$\rightarrow k' = 3, (aaa, 111) \in t_1,$	$1 \sqsubseteq 111$
$(aa, 1)$	$\in t_2, k = 2$	$\rightarrow k' = 3, (aaa, 111) \in t_1,$	$1 \sqsubseteq 111$
$(aaa, 111)$	$\in t_2, k = 3$	$\rightarrow k' = 3, (aaa, 111) \in t_1,$	$111 \sqsubseteq 111$

$$\Rightarrow t_1 \equiv_a t_2$$

$$t_1 \stackrel{?}{\equiv}_a t_3$$

$$(aaa, 111) \in t_1, k = 3 \rightarrow \nexists k'$$

$$\Rightarrow t_1 \not\equiv_a t_3$$

Because of Lemma 1 $\Rightarrow t_2 \not\equiv_a t_3$.

3.3 Timed Transducers

For the second kind of equivalence we need to introduce *timed sequences*, originally introduced as timed words in [AD94], and *timed transducers*. Note that timed sequences don't have to be monotonically increasing like in the original definition. Quite on the contrary the unorderedness of outputs is an important key principle to much of the later work as you will see.

Let \mathbb{T} be a timing model that is isomorphic to \mathbb{R} . For the examples we will use \mathbb{R} for \mathbb{T} .

Definition 7: Timed Sequence.

A sequence is called *timed*, if every element of it is associated with a timestamp: $\sigma \in (\Gamma \times \mathbb{T})^*$. For brevity a timed sequence can be written with the timestamps as the index of the elements: $\sigma = e_0 e_{0.5} e_1$.

The function

$$timed : (\Gamma \times \mathbb{T})^* \rightarrow (\Gamma \times \mathbb{T})^*$$

reorders a timed sequence σ by its timestamps, such that:

$$\forall i, j \in \mathbb{N} : \text{ if } i < j \text{ then } \pi_i < \pi_j \text{ with } (o_i, \pi_i) = \sigma_i \text{ and } (o_j, \pi_j) = \sigma_j$$

The function

$$upto : \mathbb{T} \times (\Gamma \times \mathbb{T})^* \rightarrow (\Gamma \times \mathbb{T})^*$$

removes all elements from a timed sequence, that have a timestamp bigger than the first argument.

The function

$$maxTime : (\Gamma \times \mathbb{T})^* \rightarrow \mathbb{T}$$

returns the biggest Timestamp in a timed sequence.

Example 5: Functions on Timed Sequences.

Let $\sigma = a_1 a_{0.5} a_{1.5} a_0$.

Then is

$$\text{timed}(\sigma) = a_0 a_{0.5} a_1 a_{1.5}$$

$$\text{upto}(1.3, \sigma) = a_1 a_{0.5} a_0$$

$$\text{maxTime}(\sigma) = 1.5$$

Definition 8: Monotonicity of Timed Sequences.

A timed sequence σ with alphabet Φ is called *monotonic*, if $\text{timed}(\sigma) = \sigma$

Definition 9: Timed Transducer.

A timed transducer t with input alphabet Φ and output alphabet Γ works on monotonic, timed sequences as inputs and has timed sequences as outputs:

$$t \subset (\Phi \times \mathbb{T})^* \times (\Gamma \times \mathbb{T})^*$$

Example 6: Timed Transducers.

Let $\Phi = \{a\}, \Gamma = \{b\}$.

$t_{tsc} = \{(a_0, b_0), (a_0 a_1, b_0 b_1)\}$ is a timed, causal and synchronous transducer.

$t_{tac} = \{(a_0, \epsilon), (a_0 a_1, b_0 b_1)\}$ is a timed, causal and asynchronous transducer.

For later theoretic work we have to restrict timed transducers.

Definition 10: Boundedness of Timed Transducers.

A timed transducer t with input alphabet Φ and output alphabet Γ is called *bounded*, if it satisfies:

$\forall \sigma \in (\Phi \times \mathbb{T})^* :$

if $(\sigma_{[0,k]}, \vec{\sigma}) \in t$

then $\exists k' > k$ with

$$(\sigma_{[0,k']}, \vec{\sigma} \circ \vec{\sigma}') \in t$$

and $\forall k'' > k'$ with $(\sigma_{[0,k'']}, \vec{\sigma} \circ \vec{\sigma}' \circ \vec{\sigma}'') \in t$ it holds, that

$$\text{upto}(\text{maxTime}(\vec{\sigma}), \text{timed}(\vec{\sigma} \circ \vec{\sigma}')) = \text{upto}(\text{maxTime}(\vec{\sigma}), \text{timed}(\vec{\sigma} \circ \vec{\sigma}' \circ \vec{\sigma}''))$$

Based on the definitions we can define an equivalence relationship on bounded timed transducers.

Definition 11: Observational Equivalence.

Let t_1, t_2 be two bounded timed transducers with input alphabet Φ and output alphabet Γ . They are called observational equivalent, written $t_1 \equiv_o t_2$, if they satisfy:

$$\forall \sigma \in (\Phi \times \mathbb{T})^* :$$

$$\forall (\sigma_{[0,k]}, \vec{o}) \in t_1 : \exists k', k'' \geq k \text{ such that}$$

$$(\sigma_{[0,k']}, \vec{o} \circ \vec{o}') \in t_1$$

$$\text{and } (\sigma_{[0,k'']}, \vec{o}_2) \in t_2$$

$$\text{and } \text{timed}(\text{upto}(\text{maxTime}(\vec{o}), \vec{o} \circ \vec{o}')) = \text{timed}(\text{upto}(\text{maxTime}(\vec{o}), \vec{o}_2))$$

and the same for switched t_1, t_2 .

What does observational equivalence between two transducers intuitively mean? It means that two transducers eventually produce the same output values for the same timed inputs, maybe in a different order, but with the same timestamps, which is very important. Since the values are associated with timestamps the outputs can be reordered by them and therefore be exactly equal.

Lemma 2: Observational Equivalence is an Equivalence Relationship for Bounded Transducers.

\equiv_o is symmetric, reflexive and transitive for bounded timed transducers.

Proof. Let t_1, t_2, t_3 be bounded timed transducers. Symmetry follows directly from the definition.

Reflexivity: For $(\sigma_{[0,k]}, \vec{o})$ select $k' = k''$ as the k , for which the transducer is bounded for that input.

Transitivity:

- Let $t_1 \equiv_o t_2, t_2 \equiv_o t_3$.

- First case:

Since $t_1 \equiv_o t_2 : \forall (\sigma_{[0,k_1]}, \vec{o}_1) \in t_1 :$

$\exists k'_1, k_2 > k_1$ with $(\sigma_{[0,k'_1]}, \vec{o}_1 \circ \vec{o}_1') \in t_1$ and $(\sigma_{[0,k_2]}, \vec{o}_2) \in t_2$

with $\text{timed}(\text{upto}(\text{maxTime}(\vec{o}_1), \vec{o}_1 \circ \vec{o}_1'))$

$= \text{timed}(\text{upto}(\text{maxTime}(\vec{o}_1), \vec{o}_2))$

(*)

and since $t_2 \equiv_o t_3 : \exists k'_2, k_3 > k_2$ with $(\sigma_{[0,k'_2]}, \vec{o}_2 \circ \vec{o}_2') \in t_2$

and $(\sigma_{[0,k_3]}, \vec{o}_3) \in t_3$

with $\text{timed}(\text{upto}(\text{maxTime}(\vec{o}_2), \vec{o}_2 \circ \vec{o}_2'))$

$= \text{timed}(\text{upto}(\text{maxTime}(\vec{o}_2), \vec{o}_3))$

(**)

$\text{maxTime}(\vec{o}_1)$ has to be smaller than $\text{maxTime}(\vec{o}_2)$,

else (*) couldn't hold, therefore, combined with boundedness and (**) :

$\text{timed}(\text{upto}(\text{maxTime}(\vec{o}_1), \vec{o}_2))$

$= \text{timed}(\text{upto}(\text{maxTime}(\vec{o}_1), \vec{o}_3))$

which concludes $\text{timed}(\text{upto}(\text{maxTime}(\vec{o}_1), \vec{o}_1 \circ \vec{o}_1'))$

$= \text{timed}(\text{upto}(\text{maxTime}(\vec{o}_1), \vec{o}_3))$

- The second case works the same, just switch t_1 and t_3 .

□

Example 7: Observational Equivalence.

Let

$$\begin{aligned} t_1 &= \{ & (a_0, \epsilon), & (a_0 a_1, b_1), & (a_0 a_1 a_2, b_1 b_2 b_0) & \} \\ t_2 &= \{ & (a_0, \epsilon), & (a_0 a_1, \epsilon), & (a_0 a_1 a_2, b_2 b_1 b_0) & \} \\ t_3 &= \{ & (a_0, b_0), & (a_0 a_1, b_0), & (a_0 a_1 a_2, b_2 b_1) & \} \end{aligned}$$

All three are causal, asynchronous timed transducers.

Let's see which ones are observational equivalent:

$$t_1 \stackrel{?}{\equiv}_o t_2$$

$$(a_0, \epsilon) \in t_1, k = 1, \text{maxTime}(\epsilon) = 0$$

$$\rightarrow k' = 1, (a_0, \epsilon) \in t_1$$

$$\rightarrow k'' = 1, (a_0, \epsilon) \in t_2$$

$$(a_0 a_1, b_1) \in t_1, k = 2, \text{maxTime}(b_1) = 1$$

$$\rightarrow k' = 3, (a_0 a_1 a_2, b_1 b_2 b_0) \in t_1$$

$$\rightarrow k'' = 3, (a_0 a_1 a_2, b_2 b_1 b_0) \in t_2$$

$$\text{timed}(\text{upto}(1, b_1 b_2 b_0)) = b_0 b_1 = \text{timed}(\text{upto}(1, b_2 b_1 b_0))$$

$$(a_0 a_1 a_2, b_1 b_2 b_0) \in t_1, k = 3, \text{maxTime}(b_1 b_2 b_0) = 2$$

$$\rightarrow k' = 3, (a_0 a_1 a_2, b_1 b_2 b_0) \in t_1$$

$$\rightarrow k'' = 3, (a_0 a_1 a_2, b_2 b_1 b_0) \in t_2$$

$$\text{timed}(\text{upto}(2, b_1 b_2 b_0)) = b_0 b_1 b_2 = \text{timed}(\text{upto}(2, b_2 b_1 b_0))$$

$$(a_0, \epsilon) \in t_2, k = 1, \text{maxTime}(\epsilon) = 0$$

$$\rightarrow k' = 1, (a_0, \epsilon) \in t_2$$

$$\rightarrow k'' = 1, (a_0, \epsilon) \in t_1$$

$$(a_0 a_1, \epsilon) \in t_2, k = 2, \text{maxTime}(\epsilon) = 0$$

$$\rightarrow k' = 2, (a_0 a_1, \epsilon) \in t_2$$

$$\rightarrow k'' = 2, (a_0 a_1, b_1) \in t_1$$

$$\text{timed}(\text{upto}(0, \epsilon)) = \epsilon = \text{timed}(\text{upto}(0, b_1))$$

$$(a_0 a_1 a_2, b_2 b_1 b_0) \in t_2, k = 3, \text{maxTime}(b_2 b_1 b_0) = 2$$

$$\rightarrow k' = 3, (a_0 a_1 a_2, b_2 b_1 b_0) \in t_2$$

$$\rightarrow k'' = 3, (a_0 a_1 a_2, b_1 b_2 b_0) \in t_1$$

$$\text{timed}(\text{upto}(2, b_2 b_1 b_0)) = b_0 b_1 b_2 = \text{timed}(\text{upto}(2, b_1 b_2 b_0))$$

$$\Rightarrow t_1 \equiv_a t_2$$

$$t_1 \stackrel{?}{\equiv}_a t_3$$

$$(a_0 a_1 a_2, b_1 b_2 b_0) \in t_1, k = 3, \text{maxTime}(b_1 b_2 b_0) = 2$$

$$\rightarrow k' = 3, (a_0 a_1 a_2, b_1 b_2 b_0) \in t_1$$

$$\rightarrow \vec{A}(\vec{i}, \vec{o}) \in t_3 \text{ with } \exists n \in \mathbb{N} : \vec{o}_n = b_0$$

$$\rightarrow \vec{A}(\vec{i}, \vec{o}) \in t_3 \text{ with } \text{timed}(\text{upto}(2, b_1 b_2 b_0)) = b_0 b_1 b_2 = \text{timed}(\text{upto}(2, \vec{o}))$$

$$\Rightarrow t_1 \not\equiv_a t_3$$

If t_3 weren't bounded (and therefore not finite) there would be no way to know, if it was equivalent to t_1 , because it could always produce a missing event at a later time.

Because of Lemma 2 $\Rightarrow t_2 \not\equiv_a t_3$.

3.4 Events

Events are the atomic unit of information that all computations are based on. There are three types of events: external, output and internal events.

The set of all events is denoted as E . Each event carries a value, which can be *nothing* or a value of a type (types are formally defined in the TESSLA specification, but aren't important for this thesis), a timestamp and the stream it's perceived on (for example a function call of a specific function or the name of an output stream).

The value of an event can be queried with the function v , its timestamp with $time$ and its stream with $stream$.

$E_e \subset E$ is the set of all external events, their stream corresponds to a specific trace. $E_o \subset E$ is the set of all output events, their stream is specified by an output name of the TESSLA specification. $E_n \subset E$ is the set of all internal events. Internal events are mostly an implementation detail, which denote steps of computation inside the runtime. The stream of internal events is implicitly given by the node that produces the stream of the event. Note that E_e, E_o, E_n are pairwise disjoint and $E_e \cup E_o \cup E_n = E$.

3.5 Streams

Streams are a collection of events with specific characteristics. While events are the atomic unit of information, streams represent the sequence of related events over time.

There are two kind of streams: signals, which carry values at all times, and eventstreams, which only hold values at specific times. Eventstreams can be described by a sequence of events. Signals can be described by a sequence of changes, where a change denotes that the value of a signal changed at a specific timestamp. The only difference between a signal and an eventstreams is that signals always have a value while an eventstream may return \perp when queried for its value at a specific time, which denotes that no event happened at that time. Based on the similarity of signals and eventstreams in the following we will mainly reason about eventstreams, but most things can also be applied to signals.

Formally a stream σ can be represented as the product of a sequence of events $\langle e_1, \dots, e_n \rangle$ where $time(e_i) < time(e_{i+1})$, $\forall i < n \in \mathbb{N}$. The set of all streams Σ is defined as all possible finite sequences of events $\Sigma = \{\sigma | \sigma \in E^*\}$. An external stream

σ_e is a stream consisting only of external events, the set of all external streams is $\Sigma_e = \{\sigma_e | \sigma_e \in E_e^*\} \times \mathbb{T}$. Output and internal streams are defined analogous.

To get the event of a stream σ at a timestamp π it can be queried like a function: $\sigma(\pi) = e$ with $time(e) = \pi$. When working with signals, the function will return the latest event that happened at or before t while an eventstream may return \perp . The progress of a stream, which is the timestamp of the last event that happened on them, can be obtained with $progress(\sigma) = \pi \in \mathbb{T}$. Internal and output streams can be queried for the node that produced them with $node(\sigma) = n \in N$.

3.6 Functions

A TESSLA specification consists of functions over streams. Functions generate new streams by applying an operation on existing streams. TESSLA itself defines a syntax to write a specification, a set of types and a standard library of functions, but an implementation is free to choose the functions it supports.

An example function is $add(S_D, S_D) \rightarrow S_D$: It takes two signals, which have to hold values of some numerical type, and produces a signal which holds values of the same type. The produced stream can either be assigned to a named identifier (think: a variable) or consumed by another function (function composition).

Functions can be divided into three categories: pure, unpure and timing. Pure functions, also called stateless, are evaluated only on the values their inputs have at the timestamp they are evaluated, therefore they don't have to remember a state and will only return events. Unpure, or stateful, functions are evaluated over the values if its inputs at that timestamp and a state and will return not only new events but also an updated state. As an example a function *eventCount* has to *remember* how many events already happened on it's input stream and increment that counter on every new event. Timing functions are evaluated not only on the value of events but also on their timestamp and can also manipulate it: While non timing functions will consume events at a specific timestamp and emit events with that timestamp, timing functions can emit events with a changed timestamp. In this thesis we will only look at past time functions, meaning functions can only delay timestamps, therefore can't depend on future values.

Timing functions complicate the reasoning about schedules and causality and therefore aren't included in Section 4.1. In Section 4.3 the conclusions of earlier sections will be extended to include timing functions.

3.7 Nodes

Nodes are the atomic unit of computation for the evaluation of a TESSLA specification. A node implements a single function: there is an *AddNode* which takes two input signals and produces a new signal. Therefore a node is the concrete implementation of a function in a runtime for TESSLA specifications. The set of all nodes is called N . The function of a node $n \in N$ is written as f_n .

Each node has a set of inputs, which are either external or internal streams, and one output, which is either an internal or an output stream. Nodes which have at least one external stream as an input are called *sources*. Nodes have a state, described in Section 3.10, which contains First In First Out (FIFO) queues, provided by the Erlang platform, which buffer events from its inputs for later computation.

Every new event added to a queue has to have a bigger timestamp than the previous event added to the queue. This means a queue has a kind of progress timestamp, which denotes the timestamp of the latest event added to it and which is strictly increasing over time. Queues support the standard operators for lists like *hd*, *tl*, *++* to respectively get the head, the tail or to append to the end.

3.8 TESSLA Evaluation Engine

Because functions in TESSLA specifications depend on other functions, and these dependencies have to be cycle free, the specification can be represented as a Directed Acyclic Graph (DAG), where the functions are vertices and the relationship between functions are edges. This is exactly how the TESSLA compiler outputs a specification. One can now use the DAG of a TESSLA specification to synthesize a system to evaluate it over inputs: The vertices of the DAG become nodes representing the functions and the edges are the input and output streams between the nodes. We will call this synthesized system an *evaluation engine*.

When fed with inputs (or *traces*) the engine will produce outputs. The relationship between inputs and outputs that is produced can be seen as a timed transducer. The input to an evaluation engine has to have strictly increasing timestamps. This is needed to have a known progress which can be distributed through the system. If inputs weren't ordered by their timestamp for example the absence of input events on a specific stream couldn't be detected because events could be present at a later position of the input trace. Especially for offline monitoring this obviously is no problem because the traces can simply be reordered into a strictly increasing sequence, except when multiple input events are at the same timestamp. This can be solved in two ways: either increase the timing precision when generating the traces

or manipulate the timestamps in the traces by adding a minimal offset to them if they are equal to another timestamp.

To evaluate a specification over traces, the evaluation engine has process the events that were traced. To do so the nodes have to run their computations until no more events are present (or the specification found an error in the trace). This leads to the question in which order nodes should be scheduled to perform their computation. We will use the term *step* to denote that one node was scheduled and performed its computation. While some schedules are simply not rational (think of unfairness and causality), there are many different schedules that are feasible. It has to be proven that a chosen schedule produces the correct conclusions for a specification, else the evaluation engine is not valid.

An evaluation engine is run inside an environment. The environment has knowledge over the state of the nodes, most important which nodes are enabled. Based on that information the environment is responsible of feeding the trace data to the engine: only when no sources are enabled the next trace is added to the queue of an input node. This ensures that the consumption of input signals is strictly ordered by timestamp, which is important as we will see in Chapter 4.

3.9 TESSLA Functions

TESSLA puts no restrictions on the semantics of functions other than that they have to work on streams or constants and produce streams, but allows to restrict them for evaluation approaches. We are taking advantage of that to categorize functions based on how or if at all they can be encoded in our evaluation approach. The categorization is based upon the relationship between consumption and production of input and output events that a node representing the function in an evaluation engine produces. For this we will use the terms node and functions somewhat interchangeable in the following subsections.

Nodes can only be scheduled when they are enabled, meaning they have events on their inputs buffered. When a node is scheduled, it will compute the minimal timestamp of all buffered events. The function implemented by the node is then evaluated at that timestamp, which is called the *evaluation timestamp*. This is important to understand the completeness criteria: It means that when the function is evaluated there is at least one event with that timestamp on one input. The inverse of that statement shows why this is important: a function is never evaluated at a timestamp where no event is present, therefore the system can't arbitrarily produce new timestamps.

Nodes having signals as inputs always have to remember the last occurred change of them in their state. When such a node is scheduled, the function will work on the remembered value if no new change is present at the evaluation timestamp for the signal. If a new change is present at the evaluation timestamp, it will be used and the state of the node will be updated to remember the new change.

It is important to note that all functions always have to consume an event from at least one input queue, else an evaluation engine can enter a livelock, where new events are produced forever out of nowhere. Also all functions will only produce a finite amount of events at every evaluation. Especially only timed functions can produce more than one event at an evaluation timestamp.

3.9.1 Complete Functions

Complete functions will consume one event from every input and produce one event at every timestamp they are evaluated. Most complete functions are pretty simple and often have eventstreams as inputs or only have one input. The complete functions that are present in the implemented runtime are explained in Table 3.1.

The first four functions are sources which take an external event and format them for internal use. *variable_values* takes a string containing the name and the value of a variable, casts the value to an appropriate type and produces a signal holding that produced value. The other sources work in a similar way.

3.9.2 Output Complete Functions

Output complete functions will produce a new event everytime they are evaluated but only have to consume events from some inputs and not from all. Table 3.2 summarizes all input complete functions.

3.9.3 Input Complete Functions

Input complete function consume one events from every input but can produce zero or one output events everytime they are evaluated. Table 3.3 summarizes all input complete functions.

Name	Domain	Range	Explanation
<i>instruction_executions</i>	Events	Events	Converts a trace to an event that denotes the execution of a specific instruction in the monitored program.
<i>function_returns</i>	Events	Events	Converts a trace to an event that denotes the return from a function in the monitored program.
<i>function_calls</i>	Events	Events	Converts a trace to an event that denotes the call of a function in the monitored program.
<i>variable_values</i>	Events	Signal	Converts a trace to a change that denotes the value of a variable in a monitored program.
<i>signalAbs</i>	Signal	Signal	Computes the absolute value of a signal.
<i>eventAbs</i>	Events	Events	Computes the absolute value of an event.
<i>changeOf</i>	Signal	Events	Emits an event everytime the signal changes its value holding the new value.
<i>neg</i>	Signal	Signal	Emits the mathematical opposite of the value of a real signal.
<i>signalNot</i>	Signal	Signal	Emits the Boolean negation of a Boolean signal.
<i>eventNot</i>	Events	Events	Emits the Boolean negation of a Boolean event.
<i>eventCount</i>	Events	Signal	Emits a signal holding the number of times an event occurred on the input.
<i>timestamps</i>	Events	Events	Emits an event holding the timestamp of an input event everytime one occurs.
<i>sma</i>	Events	Events	Emits an event holding the simple moving average over the last specified number of events that occurred.

Tab. 3.1: List of complete functions

3.9.4 Incomplete Functions

Incomplete functions always consume at least one input from any input and will produce zero or one event. Table 3.4 lists all supported incomplete functions. If no explanation is given, why the function is incomplete it is the following: The function is not input complete, because it only consumes events or changes that have the timestamp at which it is evaluated, if one input only has events with bigger timestamps no event or change is removed from them and the remembered last change of them is used as a base for computation if it is a signal. Also it is

Name	Domain	Range	Explanation
<i>merge</i>	Events \times Events	Events	Merges two eventstreams. When an event is present on the first input, will emit an event with the same value, else with the value from the event on the second input. Not input complete because if an event on the second input occurs at a timestamp where no event of the first input occurs no event of the first input is removed.
<i>occurAny</i>	Events \times Events	Events	Emits an event without a value everytime an event occurs on any input. Not input complete because events are only removed from both inputs if they have the same timestamp.

Tab. 3.2: List of output complete functions

Name	Domain	Range	Explanation
<i>signalMaximum</i>	Signal	Signal	Emits a change everytime the input has a bigger value than it had anytime before.
<i>eventMaximum</i>	Events	Signal	Emits a change everytime the input has a bigger value than it had anytime before or a default value if it is the biggest value occurred yet.
<i>signalMinimum</i>	Signal	Signal	Emits a change everytime the input has a smaller value than it had anytime before.
<i>eventMinimum</i>	Events	Signal	Emits a change everytime the input has a bigger value than it had anytime before or a default value if it is the biggest value occurred yet.
<i>sum</i>	Events	Signal	Emits the summed up value of all events that happened on the input upto that point.
<i>mrw</i>	Events	Signal	Emits a change everytime the input takes a new value. Not output complete because no new change is emitted if the last value of the input was the same as the current.

Tab. 3.3: List of input complete functions

not output complete, because changes of a signal are only produced if the value actually changes. For example, if the values of the inputs of an *add* are switched at a timestamp it would not produce a new change for that timestamp but consume changes from both inputs.

Name	Domain	Range	Explanation
<i>add</i>	Signal \times Signal	Signal	Adds both inputs.
<i>and</i>	Signal \times Signal	Signal	Performs a Boolean and over both inputs.
<i>div</i>	Signal \times Signal	Signal	Divides the first input by the second input.
<i>eq</i>	Signal \times Signal	Signal	Emits if both inputs are equal.
<i>geq</i>	Signal \times Signal	Signal	Emits if the first input is greater or equal to the second input.
<i>gt</i>	Signal \times Signal	Signal	Emits if the first input is greater than the second.
<i>implies</i>	Signal \times Signal	Signal	Emits the Boolean implies relationship between both inputs.
<i>leq</i>	Signal \times Signal	Signal	Emits if the first input is smaller or equal to the second.
<i>lt</i>	Signal \times Signal	Signal	Emits if the first input is smaller than the second.
<i>max</i>	Signal \times Signal	Signal	Emits the bigger value of both inputs.
<i>min</i>	Signal \times Signal	Signal	Emits the smaller value of both inputs.
<i>mul</i>	Signal \times Signal	Signal	Multiplies the first input by the second.
<i>or</i>	Signal \times Signal	Signal	Performs a Boolean or over both inputs.
<i>sub</i>	Signal \times Signal	Signal	Subtracts the second input from the first.
<i>filter</i>	Events \times Signal	Events	Emits events whenever an event occurs on the first input with the value of that event if the second input has the value true. It is not output complete because it doesn't emit events when the second input is false.
<i>ifThen</i>	Events \times Signal	Events	Emits an event with the value of the second input everytime an event occurs on the first input. It is not output complete because it only emits outputs when an event occurred on the first input.

Name	Domain	Range	Explanation
<i>ifThenElse</i>	Signal \times Signal \times Signal	Signal	Emits the value of the second input if the first is true, else of the third input.
<i>sample</i>	Signal \times Events	Events	Same as <i>ifThen</i> with switched arguments.
<i>occurAll</i>	Events \times Events	Events	Emits an event whenever events occur on both inputs. Not output complete because it only emits events whenever events happen on both inputs.

Tab. 3.4: List of incomplete functions

3.9.5 Timing Functions

Timing functions introduce a new challenge for TesslaServer. Timing functions are able to manipulate the timestamp of events. Without timing functions all events produced in an evaluation engine would have a timestamp that is equal to the timestamp of some input event. Furthermore, all functions that are not timing functions can at most produce one new event per input event, else two of the events would have to have the same timestamp which is forbidden in our definition of streams.

Timing functions can produce events at every timestamp and produce multiple events per input event. Table 3.5 lists all implemented timing functions.

Name	Domain	Range	Explanation
<i>timestamps</i>	Events	Events	Emits the timestamp of each input event as the value of an output event. Can also be categorized as a complete function since it only reads the timestamp.
<i>delayEventByCount</i>	Events	Events	Delays each event of an eventstream by a specified number of events. For example it can be used to change the timestamp of each event to the timestamp of the next event received.

Name	Domain	Range	Explanation
<i>delayEventByTime</i>	Events	Events	Delays each event of an eventstream by a specified time.
<i>delaySignalByTime</i>	Events	Events	Delays each change of a signal by a specified time.
<i>inPast</i>	Events	Signal	Emits a boolean signal which holds the value <i>true</i> whenever an event happened on the input before in a time specified.

Tab. 3.5: List of timing functions

3.10 State and History

All TESSLA evaluation engines have to hold a state, which encodes information necessary to continue the evaluation, and a history, which encodes what happened on all streams in the evaluation engine. The state of a whole evaluation engine is made up of the states of its nodes.

Each node has a state, which contains arbitray information, for example a counter for a *CountNode*, its input queues holding the non-processed events and, if they have signals as inputs, the last changes of them.

To distinguish between the two types of states, the state of the whole engine is called the *global state* and the state of a single node the *node state*. The set of all valid node states is called \tilde{N} .

The global state of an evaluation engine at a certain step is a map from its nodes to their node state. We will denote the set of all global states as S . A global state can be queried like $s(n) = \tilde{n}$ to yield the state of the node n .

Nodes, and therefore the whole evaluation engine, change their state when they are scheduled. The transition between states is described in Section 3.11

The history of an evaluation engine is defined at every step (read: after every computation of a node) as all events that were produced by any node upto that step.

3.11 Transitions

A transition describes what happens when the evaluation engine schedules a node: Events from the the inputs of a node are removed (at least one), output events can be generated (but don't have to) and distributed and the internal state of nodes are updated. Because the function of a node is evaluated at the evaluation timestamp, which is the minimal timestamp of all events on the heads of inputs, the events which are removed are exactly the ones that have the evaluation timestamp. To look at it in another way: a transition models the computation of a node and the progressing of the stream it produces towards the evaluation timestamp, which has to be bigger than the previous progress, since input queues are strictly ordered by their timestamp and the events with the minimal timestamps are removed after the computation. Therefore when we say 'Node a is scheduled' we mean that a transition is taken which models the computation of that node.

The set of all transitions is written as T . The function $node : T \rightarrow N$ returns the node of which the transitions models the computation.

One part of a transition is a relation between two sets of events, why sometimes we write $\tau = (\{e_1, e_2\}, \{e_3\})$ to visualize a transition, but remember that there is more to a transition than that. The relation $\tau = (\{e_1, e_2\}, \{e_3\})$ means that two events were consumed by a node and one event was produced based on them.

The empty transition, meaning no input was consumed and no output produced, is labeled with λ . Note that all transitions, which produces events have to consume at least one event (therefore no events can be created from nowhere) and that it's possible that no event was produced based on the consumed events (see Section 3.9). Furthermore with timing functions it's possible to create multiple events in one transition. For example think of an *EchoNode*, which duplicates an input after a specified amount of time.

Definition 12: Application of a Transition on a State.

Given global state s_0 and transition $\tau = (\tilde{E}, \tilde{E}') = (\{e_1, e_2, \dots, e_i\}, \{e'_1, e'_2, \dots, e'_i\})$

with $n = \text{node}(\tau_1)$ and N_c the set of all nodes that are children of n . When we apply τ to s_0 , written $\text{apply}(s_0, \tau) = s_1$, we get a new global state s_1 with

$\forall \tilde{n}_i = s_0(i)$
 if $\text{node}(\tilde{n}_i) \notin N_c \cup \{n\}$
 then $s_1(i) = \tilde{n}_i$
 (nothing changes for independent nodes)
 else if $\text{node}(\tilde{n}_i) \in N_c$
 then append all events in \tilde{E}' to the input representing the stream from n
 else
 remove all events in \tilde{E} from the inputs representing their streams
 and update the internal information based on the function n is modelling

This means that the new global state is built with the old global state by altering only the node states of the node identified by the transition and its children. The node states are altered by removing all events that were consumed by the function from the inputs of the scheduled node, updating its internal state and adding the produced events to the queues representing it in the node states of its children.

3.12 Run

A run of an evaluation engine is a sequence of transitions and states. The first element of the sequence is the empty transition and the initial state of the evaluation engine. It is a representation of the steps the engine takes to evaluate a specification over input streams. The length of a run can be retrieved with $\text{length}(r) = d \in \mathbb{N}$. A run can be queried by its index to return the element at that index: $r(i) = (\tau_i, s_i)$, $i \in [0, \text{length}(r)]$.

The run $\langle (\lambda, s_0), (\tau_1, s_1) \rangle$ means, that the engine was in its initial state, took the transition τ_1 and thereby reached the state s_1 .

Definition 13: Closeness of Runs.

The closeness δ of a run r_1 to a run r_2 is a pair $\delta(r_1, r_2) = (x, y)$, where x is the index before the first position where the two runs differ and y is the number of steps between the index of the first difference and the position where r_2 takes the transition that r_1 took after step x . The closeness of runs is ordered element-wise: $(x, y) > (x', y') \leftrightarrow ((x > x') \vee (x = x' \wedge y < y'))$. Therefore two runs with length d are equal, if their closeness is $d, 0$, which is the maximal closeness two runs of length d can have at all. Note that two runs have no closeness if they don't contain the same transitions.

Example 8: Closeness of Runs.

Let

$$r_1 = \langle (\lambda, s_0), (\tau_1, s_1), (\tau_2, s_2), (\tau_3, s_3), (\tau_4, s_4), (\tau_5, s_5), (\tau_6, s_6) \rangle$$

$$r_2 = \langle (\lambda, s_0), (\tau_1, s_1), (\tau_2, s_2), (\tau_5, s'_3), (\tau_4, s'_4), (\tau_6, s'_5), (\tau_3, s'_6) \rangle$$

$$r_3 = \langle (\lambda, s_0), (\tau_1, s_1), (\tau_2, s_2), (\tau_3, s_3), (\tau_5, s''_4), (\tau_4, s''_5), (\tau_6, s''_6) \rangle$$

Then is

- $\delta_{1,2} = \delta(r_1, r_2) = (3, 3)$ because r_1 takes τ_3 at step 3 while r_2 takes τ_5 and r_2 takes τ_3 at step $3 + 3 = 6$.
- $\delta_{1,3} = \delta(r_1, r_3) = (4, 1)$ because r_1 takes τ_4 at step 4 while r_3 takes τ_5 and r_2 takes τ_4 at step $4 + 1 = 5$.
- The explanations for the remaining cases is analogous and therefore not stated here.
- $\delta_{2,1} = \delta(r_2, r_1) = (3, 2)$, $\delta_{2,3} = \delta(r_2, r_3) = (3, 1)$
- $\delta_{3,1} = \delta(r_3, r_1) = (4, 1)$, $\delta_{3,2} = \delta(r_3, r_2) = (3, 2)$

The ordering of the distances is straightforward: $\delta_{1,3} < \delta_{2,3} < \delta_{2,1} < \delta_{3,1}$.

To reason about runs we have to restrict runs to the ones that are reasonable in the context of an evaluation engine. This means that only transitions are taken that are possible based on the global state. To do so we have to define when a node can compute based on its state. At first we will give a definition that only works for output complete TESSLA functions. Based on that definition we will see the problems that output incomplete functions have and modify the transition model to fix the problem.

Definition 14: Enabledness of a Node.

A node n with the node state \tilde{n} containing the input queues $\tilde{\sigma}$ in an evaluation engine is called enabled at a step i of a run r of that engine, if at least one input is buffered on each input queue:

$$\forall \sigma_x \in \tilde{\sigma} : \neg \text{empty}(\sigma_x)$$

Now it is possible to restrict runs to a subset where each run models a rational evaluation of a specification.

Definition 15: Valid Run.

A run r is called valid, if

- $\forall i \in [0, \text{length}(r)] : r(i) = (\tau_i, s_i) \wedge \text{node}(\tau_i)$ is enabled at step i
- and $s_i = \text{apply}(s_{i-1}, \tau_i)$

As stated, the definition for enabledness works for all output complete functions, while output incomplete functions have a problem. It is possible that an output incomplete functions will never produce a new output: Think of a *FilterNode* where the second input is always *false*. Even if new input events are added to the first input and the second input is known to be *false* upto any timestamp, the children of the node will never receive new events for that input queue. Therefore no children can ever again compute, because they don't have an event buffered on all inputs.

There are two ways to fix this:

All input queues could be extended with a progress timestamp, which denotes how far the parent node has progressed. Now everytime a node n with children N_c evaluates its function at a new evaluation timestamp the inputs of all nodes in N_c representing the stream of n would be updated to have that evaluation timestamp as their new progress, even if no new event was produced at that timestamp.

The other way, which is used in this thesis, doesn't alter the input queues: First we need a new type of events: *progress events*. A progress event holds no value and exists only to notify a node that a specific input has progressed upto the timestamp of the progress event. We will write a progress event as e_π^p where π is the timestamp of it.

Now whenever a node performs its computation at an evaluation timestamp and no new event was produced a progress event with the evaluation timestamp is added to the corresponding input queue of all children. When a node is scheduled and one of the consumed events is a progress event, what happens depends on the function of the node. Some examples are:

- An *AddNode* has at an evaluation timestamp a progress event as the first input and no event as the second input, which means there is an event buffered on input queue two, else the node wouldn't have been scheduled, and that event has a bigger timestamp than the evaluation timestamp, else its timestamp would have been the evaluation timestamp. The node can't produce a new change for its output, since all new information is, that input one hasn't changed upto the new timestamp. Therefore the node will distribute a progress event with the evaluation timestamp to its children.
- If the *AddNode* would have received a progress event on input one and a new change on input two at the same timestamp, it could emit a new change. The new event would have the value of the last change on the first input, which the node has to hold in its state, added with the value of the new change on the second input.
- A *MergeNode* has a progress event as the first input at an evaluation timestamp and no event on the second input. By the same reasoning as in the first

example, the node can't produce a new output and therefore will produce a progress event.

These examples show that progress events alters the categorization of TESSLA functions: output complete functions could produce no new normal output, if inputs are progress events. This is no problem, since the main reason for the categorization was to explain the necessity of progress events. With progress events all functions are output complete, because they always emit either normal events or progress events.

For the comparison of runs some more definitions are needed.

Definition 16: Independence of Nodes.

A node a is called independent of node b in an evaluation engine, if a is no descendant of b .

Definition 17: Independence of Transitions.

A transition τ_1 is called independent of another transitions τ_2 , if $node(\tau_1)$ is independent of $node(\tau_2)$.

Lemma 3: Exchange of Independent Transitions.

If a transition τ_2 is independent of a transition τ_1 , then for all runs of the evaluation engine that produces the runs the following holds:

*If $r_1 = \langle (\lambda, s_0), \dots, (\tau_1, s_i), (\tau_2, s_j), \dots, (\tau_l, s_l) \rangle$ is a valid run
then $r_2 = \langle (\lambda, s_0), \dots, (\tau_2, s'_i), (\tau_1, s_j), \dots, (\tau_l, s_l) \rangle$ is a valid run*

Proof. As a first step we will show, that the transitions τ_1 and τ_2 can be exchanged because their enabledness doesn't depend on each other.

Because $b = node(\tau_2)$ is no descendant of $a = node(\tau_1)$, the stream σ with $node(\sigma) = node(\tau_1)$ can be no input of b . Therefore the enabledness of b can't be changed by τ_1 , taken directly from the definition of enabledness. So b has to be enabled before τ_1 was taken in r_1 or else it couldn't be enabled afterwards and τ_2 couldn't be taken in the next step. Therefore r_2 also fulfills the requirements of a valid run up to and including the steps $(\tau_2, s'_i)(\tau_1, s_j)$.

As a second step we have to show, that the state s_j will stay the same, no matter the order of the two transitions and only the state s_i may be changed to s'_i in r_2 . If this holds, the subsequent states can't change either, since they are deterministically built by applying the same transitions. Now if all subsequent states stay the same, all subsequent transitions will stay enabled, since enabledness of a node only depends on the state. Figure 3.1 visualizes the exchange of the transitions and the argument why the state stays the same.




Fig. 3.1: Influence of the order of independent transitions on the global state of an evaluation engine

Remember that $s_i = \text{apply}(s_{i-1}, \tau_1)$ and $s'_i = \text{apply}(s_{i-1}, \tau_2)$. We have to show that $\text{apply}(\text{apply}(s_{i-1}, \tau_1), \tau_2) = \text{apply}(\text{apply}(s_{i-1}, \tau_2), \tau_1)$. This is straightforward when recalling what a transition does when it's being applied. $\text{apply}(\text{apply}(s_{i-1}, \tau_1), \tau_2)$ appends all events produced by τ_1 to the queue representing $n_1 = \text{node}(\tau_1)$ in the node states of the children of n_1 and updates the node state of n_1 . Then it appends all events produced by τ_2 to the queue representing $n_2 = \text{node}(\tau_2)$ in the node state of the children of n_2 and updates the node state of n_2 .

Since n_1 and n_2 are independent of each other, neither of them have input queues representing the other, therefore no event can be added to a queue of n_2 with τ_1 and the same with n_1 and τ_2 . Therefore all events that were appended by applying τ_1 will still be in the queues after τ_2 was applied, since only events from the queues of n_2 can be consumed, and the same if they were applied in reversed order.

To conclude: All events appended to any node states will still be present after both transitions were applied and the node states of n_1 and n_2 will be updates based on the same events, no matter in which order the transitions are run. Therefore the state after both transitions are applied have to be the same, no matter the order in which they are applied.

□

Lemma 4: Duration of Enabledness.

A node which is enabled stays enabled at least until it is scheduled. Formally: If a node n is enabled at step i in a run r it will stay enabled at least until the first step $j > i$ with $r(j) = (\tau, s)$, $\text{node}(\tau) = n$. Note that it doesn't have to be disabled after step j , because there could have been multiple events buffered on it's inputs.

Proof. Let n be a node enabled at step i in a run r . Let $(\tau_x, s_x) = r(x)$, $x = i + 1$. If $\text{node}(\tau_x) \neq n$, then the only influence that τ_x can have on the node state \tilde{n} of n is by appending produced events to one of the input queues, as per the definition of application of a transition. Since τ_x can't remove events from the input queues in \tilde{n} and on all input queues was at least one event buffered before the transition was

applied, else n wouldn't have been enabled at step i , \tilde{n} will have at least as many input events buffered on every input as in the step before. This means that n is still enabled at step x , and by induction at every later step until a transition with n as its node is taken. \square

Lemma 5: Finiteness of Enabledness.

Whenever a node is enabled in a run, it can only be scheduled continuously a limited number of times until becoming disabled.

Proof. Let n be an enabled node at step i with the node state \tilde{n} . First let's assume that no parent of n are scheduled after step i and before n becomes disabled. Since the input queues of n are filled by previous transitions and only finite number of events are produced at every step, all of the queues can only have a finite number of events buffered at step i . Since no parents of n are scheduled, the queues can't get fuller. Because all nodes represent functions, and all functions have to consume at least one input event (compare Section 3.9), everytime n is scheduled at least one input queue of n will have one event less buffered after it performed its computation. Therefore the worst case is when n only consumes one event per computation. The maximum number of time n can be scheduled is therefore bounded by the sum of the number of events on all input queues at step i .

Now let's assume that also input nodes of n can be scheduled after step i . Everytime an input node is scheduled it will add a finite amount of events to one input queue of n . This leads to a cyclic behaviour: If an input queue can be scheduled infinitely often, infinite many events will be added to the input queue and n can possibly be scheduled infinitely often. If input queues can not compute infinitely often, only a finite amount of events are added to the inputs of n and therefore n can only be scheduled a limited number of times. Closer inspection of the nature of evaluation engines show why this is no problem: Only if the external trace fed to an evaluation engine is infinite the source nodes can compute infinitely often. Hence for finite traces no node can compute infinitely often. Again the worst case is when n only consumes one event per computation. The number of times it can be scheduled is limited by the sum of the number of events buffered on the inputs at step i and the number of events that are produced by inputs of n after step i . \square

With the notion of runs and especially enabledness we can now define which schedules are seen as fair.

Definition 18: Fair Schedules.

A schedule of an evaluation engine is called fair, if for all runs r it produces the following holds:

$$\begin{array}{ll} \forall i < \text{length}(r) : & \text{if } n \text{ is enabled at step } i \\ & \text{then } \exists j \geq i \text{ such that } n \text{ is scheduled at step } j \end{array}$$

In other words, every enabled node is scheduled after a finite number of steps.

Building on this we will investigate different fair schedules in the next chapter.

Behaviours of Evaluation Engines

Based on the definitions in Chapter 3 we will now look at different schedules of evaluation engines and compare them. This is done in multiple steps: starting with a small subset of allowed schedules and functions and iteratively adding more complex cases.

For the comparison we will use timed transducers from Section 3.3. To do this the notion of a run of an evaluation engine is not sufficient: transducers describe a relationship between inputs and outputs, runs describe stepwise generation of internal and output events. Therefore the *behaviour* of a run is defined, which maps a run to relationship between inputs and outputs.

Definition 19: Behaviour of a Run.

Let r be a run of an evaluation engine. The behaviour β_r of it is a timed transducer: A set of tuples of timed sequences. It is calculated as follows:

1. Let $\beta_r = \emptyset$ and r_p an empty prefix of r .
2. Remove the prefix from r , where the first transition consumes an input upto but not including the next transition where an input is consumed, and append it to r_p .
3. Select the sequence of all output events O_p (which is possible empty) that are produced at any step in r_p .
4. Select the sequence of all input events E_p that are consumed at any step in r_p .
5. Add the tuple (E_p, O_p) to β_r .
6. Goto step 2 if r is not empty, else terminate.

Stated simple the run is chopped into pieces, where each piece begins with the consumption of an input events and ends before the next input is consumed. The pieces are then merged from left to right: First take all inputs and outputs consumed and produced upto the current piece and add them to the behaviour, then merge the current piece with the next and repeat.

Example 9: Construction of a Behaviour.

In Table 4.1 it is shown how a behaviour is built from a run. The run is denoted only by its transitions, events are labeled based on their type: e_i are external (read: input)

Run $(\{e_1\}, \{i_1\}) (\{i_1\}, \{i_2\}) (\{i_1, i_2\}, \{o_1\}) (\{e_2\}, \{i_3\}) (\{i_2, i_3\}, \{o_2\}) \leftarrow$
 $(\{i_3\}, \{o_3\}) (\{e_3\}, \{i_4\}) (\{e_4\}, \{i_5\}) (\{i_4, i_5\}, \{o_4\})$
 Tuples $(e_1, o_1) (e_1 e_2, o_1 o_2 o_3) (e_1 e_2 e_3, o_1 o_2 o_3) (e_1 e_2 e_3 e_4, o_1 o_2 o_3 o_4)$

Tab. 4.1: Example how the behaviour of a run is constructed

events, i_i are internal events and o_i are output events. Parts of the run in the same color are the transitions that end up in the same piece when applying the construction algorithm. The tuples show how the sequence of input and output events from the pieces are extracted. The behaviour of the run is the set of the tuples.

The behaviour of a run is a timed transducer since all events have timestamps and all consumed events are strictly ordered by their timestamp, since inputs to an evaluation engine are required to be ordered by their timestamp.

Since the behaviour encodes the relationship between inputs and outputs a run produces it provides the foundation to reason about equivalence between different runs and whole evaluation engines.

Definition 20: Equivalence of Runs.

Two runs are called equivalent if their behaviour is observational equivalent.

Now we can define when two evaluation engines are called equivalent based on their runs

Definition 21: Equivalence of Evaluation Engines.

Two evaluation engines are called equivalent, if for every run that one can produce there is an observational equivalent run in the other.

4.1 Schedules Without Timing Functions

For a first step we specify and compare behaviours of different approaches to evaluate TESSLA specifications without timing functions. Without timing functions all nodes work only on values or the presence of events and will emit exactly one event at every computation, either a normal or a progress event. This leads to behaviours that can be easily reason about, as seen in the next sections.

All systems to evaluate TESSLA specifications we will look at are based on the described structure in Section 3.8. While there are other approaches to evaluation, a DAG based approach seems to fit most naturally and focusing on one structure makes comparing systems easier.

Let's recap and summarize how an evaluation engine performs its computation. Each evaluation engine will work in steps, where each step is synonymous with

an index in the run of the system. Therefore at each step one enabled node is scheduled to perform its operation, represented as the transition in the run. The transition will encode one of the following three things that can happen:

- The next external event (external events have to be totally ordered by their timestamp) can be consumed by a source in the DAG, which generates internal events, that are propagated to its children.
- An internal node, which has at least one new input buffered on all of its input queues, can perform its computation and generate a new internal event, which is propagated to the children of that node.
- An output node, which has at least one new input buffered on all of its input queues, can produce a new output.

Evaluation engines are free in the way they are scheduling their nodes, only limited by causality (no event can be consumed before it's produced), which is guaranteed by the enabledness criteria. In the following evaluation engines are classified by their scheduling approaches.

4.1.1 Greedy Evaluation Engines

The first class of evaluation engines are called greedy.

Definition 22: Greedy schedule.

A schedule of an evaluation engine built by the following steps is called greedy.

1. *Select all nodes that are no sources, let their count be i*
2. *Label them with unique natural numbers from $[1, i]$ in reverse topological order*
3. *Label the remaining nodes with unique natural numbers bigger than i*
4. *Schedule the enabled node with the lowest label first*

We also call an evaluation engine greedy if we mean it's run with a greedy schedule.

Obviously for many DAGs there is no unique reverse topological order, therefore one can be chosen by the evaluation engine. We will show in Section 4.2.1 that all topological orders will produce observational equivalent behaviours.

The greedy schedule ensures that no node is scheduled which has a successor that can be scheduled, therefore events are *pushed* through the DAG towards an output node as fast as possible. As shown in Section 4.1.1 any schedule built like this is fair.




Fig. 4.1: Visualization of a simple evaluation engine with a greedy schedule. The left side shows the initial state of the evaluation engine with one input event that can be consumed. The right side shows the state after the input event was consumed and now two internal events are sent to the children of the input node.

Greedy evaluation engines offer a good start to reason about behaviours and will be used as a comparison for all other evaluation engines.

Definition 23: Valid Evaluation Engines.

An evaluation engine is called valid if it is equivalent to a greedy evaluation engine.

Figure 4.1 visualizes a greedy evaluation engine. It shows two DAGs representations of an evaluation engine where the nodes a to e are labeled in a reversed topological order and o and u represents two output streams. The left system is in its initial state and an input event e_1 is present and can be consumed by the input node a . When a node is chosen to compute by the scheduler, only node a is enabled, therefore it is scheduled. The right system is the representation of the next step: node a has consumed the external event and produced an internal event n_{a1} , which is propagated to all its children: nodes b and d . In the next step node, b would be scheduled, because it has the lowest number of any node that can compute (actually it's the only node that can compute at all, because d has to wait for the event from b). After b is scheduled, it would produce the internal event n_{b1} which would then be distributed to nodes c and d .

The complete run of the greedy engine for one input is the following, where the states are omitted:

$$\langle (\lambda, s_0), ((\{n_{a1}\}, \{n_{b1}\}), s_1), ((\{n_{b1}\}, \{o_1\}), s_2), \\ ((\{n_{a1}, n_{b1}\}, \{n_{d1}\}), s_3), ((\{n_{d1}\}, \{u_1\}), s_4) \rangle$$

If there were more than one input event, at this point node a would be scheduled again. It would consume the next external event and the following nodes would be scheduled in the same order as before, extending the run in an obvious way.

Fairness of Greedy Schedules

It remains to show that greedy schedules are fair.

Lemma 6: Greedy Schedules are Fair.

Any greedy schedule is fair.

Proof. Let a be a node with the label n , which is enabled at step i and is no source. Because evaluation engines can only contain a finite number of nodes there can only be a finite number of enabled nodes with a smaller label than n . Because of Lemma 5, all nodes with a smaller label than n will become disabled after a finite number of steps. Let that number be j . The only way new events could enter the system are through sources, but they have bigger labels than n , as by the definition of the schedule, and therefore can't be scheduled before n . Because of Lemma 4, a will still be enabled after these steps. So a is the enabled node with the lowest label at step $i + j$ and therefore will be scheduled.

Now let a be a source. Sources are only scheduled when no internal node is enabled since they are labeled with higher numbers than all internal nodes. Based on the same reasoning as in the first case at some point all internal nodes will become disabled, therefore a source node has to be scheduled. This source can either be a or another source, recall that only one source is enabled at any time because of the environment of an evaluation engine. If another source was scheduled, after a finite amount of steps all internal nodes will have to become disabled again. Since finite traces are evaluated at some point either the trace will end without ever feeding an input to a , then a will never be enabled, or at some point a will receive an external event. When a receives an external event, it will be the only enabled source, else no input would be fed to an input at that step. Therefore a will be scheduled the next time no internal nodes are enabled. \square

4.1.2 Fair Evaluation Engines

Obviously greedy schedules are only a small subset of all fair schedules. As the next step we will look at the rest of them.

Definition 24: Fair Evaluation Engines.

A fair evaluation engine is one with a fair schedule.

In contrast to a greedy evaluation engine a fair one has no fixed schedule, meaning that at each step any enabled node can be scheduled. Therefore predecessors of enabled nodes can perform multiple computations before their children are scheduled and events are not *pushed* through the DAG as fast as possible.

The difference between greedy and fair schedules are similar to the ones of synchronous and asynchronous transducers: A greedy schedule will ensure that outputs are produced as fast as possible while a fair can *delay* the outputs by consuming multiple inputs first and scheduling internal nodes multiple times before scheduling an output node. But note that there is an important difference between synchronous transducers and the behaviour of a greedy evaluation engine: A greedy evaluation engine can produce multiple events at every step.

4.2 Equivalence of Different Schedules Without Timing Functions

The behaviour of a run of an evaluation engine with a given schedule allows us to reason about equivalence.

As by Definition 23 any evaluation engine has to be equivalent to a greedy one to be valid.

The equivalence is shown in two steps: first in Section 4.1.1 it is shown, that all possible greedy engines for a specification are equivalent, so there is only one valid evaluation for a specification over a fixed input. Afterwards in Section 4.2.2 it is shown that any fair evaluation engine is equivalent to a greedy one.

4.2.1 Equivalence of Greedy Systems

When given a series of input events, two greedy evaluation engines for a specification with different schedules will have different runs. But both will produce all outputs that can be produced after consuming one specific input before the next input is consumed as reasoned in Section 4.1.1. Also both runs will obviously have the same length (both engines are the same DAG, so they have the same number of nodes), let that length be l .

To proof the equivalence of both engines we can prove the equivalence of their runs. To show the equivalence it is shown that two runs r_1 and r_2 of two evaluation engines based on the same graph but with a different greedy schedule can always be reordered to become closer while preserving observational equivalence. If such

a closer run always exist, we will show that the run with closeness $(l, 0)$ to r_1 , which has to be r_1 itself, is also observational equivalent to r_2 .

Theorem 1: Equivalence of Different Greedy Evaluation Engines.

Two greedy evaluation engines for a specification with different schedules are equivalent.

Proof. Let r_1, r_2 be the runs of two greedy evaluation engines for the same specification which received the same inputs. Because each TESSLA specification contains only a finite amount of functions and works on finite traces, the runs also have to be finite.

If the two runs aren't equal, they must have a closeness which is smaller than $(l, 0)$. Let $[r_2]$ be the set of all runs that are observational equivalent to r_2 . If r_1 is in this set, we would be done. Let's assume that r_1 is not in the set. Therefore all runs in the set have to have a smaller closeness than $(l, 0)$ to r_1 , since the only run with a closeness of $(l, 0)$ to another run is the run itself. Select one run $r'_2 \in [r_2]$ which has the biggest closeness to r_1 . Let $(d, k) = \delta(r_1, r'_2)$.

This means that at step d the run r'_2 has taken a different transition than run r_1 . Let the transitions the runs have taken be τ_1 for r_1 and τ_2 for r'_2 . Run r'_2 will take transition τ_1 at step $d + k$ (as per the definition of the closeness). Obviously the two transitions have to be independent of each other, else they couldn't have been taken in different order by the two runs.

If $k > 1$ there will be a transition $\tau'_2 \neq \tau_1$ which is taken by the run r'_2 at step $d + (k - 1)$. While this transition τ'_2 must also be taken in the first run as per Lemma 4, it's not possible, that it was taken before τ_1 , because then the two runs wouldn't have been the same upto the point where τ_1 was taken. Therefore τ_1 has to be independent of τ'_2 , and because τ'_2 was scheduled by the second run before τ_1 both transitions are independent of each other.

As of Lemma 3 which one of them is taken first won't change the rest of the run at all after both transitions were applied and the run stays valid. Therefore there is a valid run r''_2 , which is equal to r'_2 , except that the transitions τ_1, τ'_2 are scheduled the other way around. See Figure 4.2 for a visualization of the runs. Dotted edges represent multiple transitions. The visualization shows how the three runs *branch* and *merge* after certain steps. The run r_1 is marked in red, r''_2 in green and r'_2 in blue. As you can see the run r''_2 has a bigger closeness to r_1 since it takes τ_1 before τ'_2 . Note that r_1 is only known upto step d , especially the transitions τ_2, τ'_2 will be taken eventually, but it isn't known when or in which order.

When two adjacent transitions are exchanged basically two things can happen:

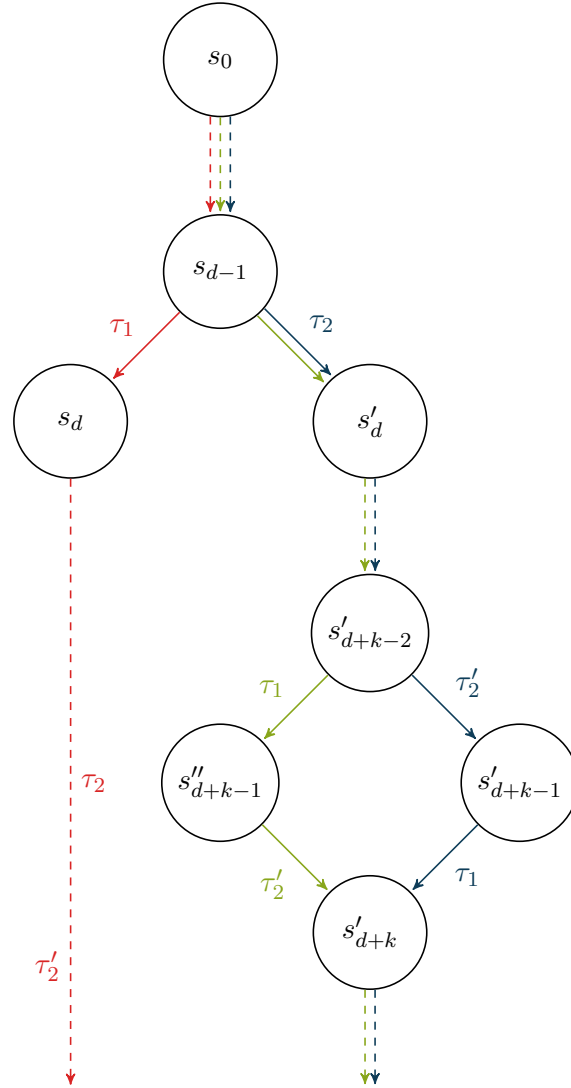


Fig. 4.2: Visualization of three runs of an evaluation engine. The left run differs from the two right runs after state s_{d-1} while the two right runs only differ by the order of the transitions τ_1 and τ'_2 .

- Outputs can be produced later, maybe moving them to the next piece in the construction of a behaviour (this happens when a transition producing outputs is exchanged with the next transition which consumes an external event)
- Outputs can be produced earlier, which can move them to an earlier piece in the construction of the behaviour (this happens in the opposite case, where a transition consuming an external event is pushed before one that produces outputs)

While the order of outputs in the tuples of the behaviour will change when two transitions producing outputs are exchanged, observational equivalence isn't influenced, since it is defined over reordered outputs.

Let's test if r_2'' is observational equivalent to r_2' : We will compare how the behaviour of the run r_2' changes when the transition τ_1 is exchanged with τ_2' . The comparison is based on the different cases the transitions can encode. Some of the cases are not feasible, they are listed for the sake of completeness and to explain why they can't happen. The cases are:

1. No inputs are consumed and no outputs produced by either transition. This obviously won't change the behaviour at all.
2. τ_2' consumes an input and doesn't produce an output and
 - a) τ_1 doesn't consume an input and doesn't produce an output. This won't change the behaviour, since τ_1 didn't add anything to it in the first place.
 - b) τ_1 doesn't consume an input, but produces one or more outputs. This changes the behaviour, since τ_1 is now part of the piece starting before τ_2' . Therefore the produced outputs are now part of one more tuple of the behaviour. But note that it still is part of all tuples built from the later pieces of the run. Therefore the two behaviours are still observational equivalent: All tuples in the behaviour are still the same, except that one or more output events are produced one step earlier, which doesn't hurt observational equivalence.
 - c) τ_1 consumes an input but doesn't produce an output. This case can't happen, since inputs to evaluation engines are totally ordered by their timestamp and therefore τ_2' couldn't have been scheduled after τ_1 in r_1 and before τ_1 in r_2 .
 - d) τ_1 consumes an input and produces outputs. This case can't happen for the same reason.
3. τ_2' produces one or more outputs and doesn't consume an input and

- a) τ_1 doesn't consume an input and doesn't produce an output. This won't change the behaviour, since τ_1 didn't added anything to it in the first place.
 - b) τ_1 doesn't consume an input, but produces one or more outputs. This only changes the order of the output events in the behaviour, but since they are reordered by their timestamp and the order of events with the same timestamp isn't important for observational equivalence, the new run is still observational equivalent.
 - c) τ_1 consumes an input but doesn't produce an output. This will *delay* the production of the outputs from τ_2' by one piece of the chopped run. While this changes the behaviour it preserves observational equivalence.
 - d) τ_1 consumes an input and produces outputs. This is kind of a combination of the previous two cases. The outputs from τ_2' are *delayed* by one piece and the outputs of τ_1 are now produced before them. But still all outputs are produced, only in different order and maybe one step later. Therefore observational equivalence holds. Also note that such a transition isn't very useful as argued in the next case.
4. τ_2' produces one or more outputs and consumes an input. First of all note that this is a rather made up combination that can only happen when a source is an output node at the same time and therefore doesn't have much of a purpose. But for the sake of completeness let's look at the cases following from this:
- a) τ_1 doesn't consume an input and doesn't produce an output. Again won't change the behaviour, as in earlier cases.
 - b) τ_1 doesn't consume an input, but produces one or more outputs. Preserves observational equivalence since the outputs of τ_1 are only produced one step earlier than before.
 - c) τ_1 consumes an input. This case can't happen as reasoned in Case 2c.

So for all cases that can happen, the run which is obtained by exchanging the two adjacent transitions is observational equivalent to the run without the change. The exchange of the two transitions brings the closeness of the new run by construction one step closer to r_1 . This means, since observational equivalence is transitive, that there is an observational equivalent run to r_2 , the run r_2'' , which has at least the closeness $(d, 1)$.

If $k = 1$ the transition τ_2' from the previous case is equal to τ_2 . The reasoning for all cases stays exactly the same, in the end we will obtain a run r_2'' which is observational equivalent to r_2' but has the closeness $(d, 0)$ to r_1 . This obviously doesn't make sense: The first element of the closeness is the last step where both runs

are equal, the second element describes how many steps afterwards the differing transition was taken. But if it was taken right in the step after the last equal step, there is no difference at that position, so the closeness of r_1 and r_2'' can be at least $(d + 1, x), x \in \mathbb{N}_{>0}$. This also contradicts our initial statement that r_2' was the run with the biggest closeness to r_1 which is observational equivalent to r_2 .

Combined we can now say, that there is no upper bound on the closeness of observational equivalent runs of r_2 to r_1 , therefore the run with the closeness $(l, 0)$ also has to be equivalent to r_2 . And as already stated, only the run r_1 can have the closeness of $(l, 0)$ to r_1 . Therefore r_1 has to be observational equivalent to r_2 . \square

This characteristic of greedy schedules gives us a baseline to compare other schedules to. Since all greedy schedules of an evaluation engine produce observational equivalent behaviours we can choose any run produced by such a schedule and compare any other run to it. In the next section we will do this for all fair schedules.

4.2.2 Equivalence of Greedy and Fair Evaluation Engines

Let's recap what fairness of a schedule means: Whenever a node becomes enabled in a run it has to be scheduled eventually. What makes fair schedules harder to reason about than greedy ones is that for one they don't have to be deterministic and furthermore that it's possible that an enabled node is not scheduled for a very long time.

Before reasoning about equivalence of greedy and fair schedules let's look at a kind of fair schedule that can be seen as worst case. Basically it is the reverse of a greedy schedule: Always schedule the enabled node that is closest to a source. Note that this schedule is not fair for infinite input traces. Stated simple this schedule will consume all input events and produce all output events per *level* of the DAG, starting at the sources and moving towards the outputs. The behaviours of runs with such a schedule are pretty special: Since no output is produced before all external events are consumed (except if a source is also an output) only one tuple of the behaviour will contain any outputs, to be specific the one which contains the sequence of all inputs as the first element. An abbreviated example of such a run is $(e_1, ()) (e_1e_2, ()) (e_1e_2e_3, ()) (e_1e_2e_3e_4, o_1o_2o_3o_4o_5)$. Such a run needs obviously more reordering than a greedy one to become observational equivalent to another greedy one since greedy schedules try to produce outputs as early as possible.

This example shows what the difference is when reordering a fair run in contrast to a greedy run: basically more transitions have to be reordered since outputs can be produced later.

Let's revisit the cases from Section 4.2.1: Actually the Cases 2b and 4b can't happen for greedy runs. If in a run r_1 a transition τ_1 , which produces an output, has to be exchanged with an earlier transition τ_2 , that consumed an external event, to become closer to a greedy run r_2 , r_1 couldn't have been greedy in the first place: Greedy schedules ensure by construction that all outputs that can be produced based on consumed external events are produced before the next external event is consumed. If τ_1 happened before τ_2 in a greedy run, the outputs from τ_1 can be produced without consuming the external event from τ_2 before, hence a run in which τ_2 happens before τ_1 couldn't be produced by a greedy schedule.

It's noteworthy that actually the whole proof of Theorem 1 doesn't depend on the fact, that the runs are produced by greedy schedules. The only real requirement is fairness, meaning that all transitions that can happen will eventually happen. Therefore the proof does hold without change for Theorem 2.

Theorem 2: Equivalence of Fair and Greedy Evaluation Engines.

Any fair evaluation engine is equivalent to a greedy evaluation engine for the same specification.

4.3 Timing functions

As stated before timing functions add complexity. But, as we will see, the reasoning from previous sections still hold for evaluation engines which include timing functions. The two characteristics of timing functions are, that they can produce events with timestamps that is not equal to the timestamp of any input event and that they can produce multiple output events for every input event.

The proof for equivalence of asynchronous evaluation engines does not depend on the fact, that events only carry timestamps that are present on input events. Therefore this characteristic of timing functions doesn't interfere with our theory.

That a timing function can generate multiple outputs for input is more critical. The proof requires that evaluation engines terminate eventually when the input trace is finite. If each node of an evaluation engine can only generate at most one event per input, the number of total events that can be produced is bound by the size of the input trace, the number of nodes and the number of edges of the evaluation engine. If a timing function is allowed to generate infinitely many outputs for an input the evaluation engine would never terminate. For now we choose to not allow such a function.

For future work we like to note that the fact that we only consider fair schedules which will schedule each enabled node eventually. Therefore even timing functions

which generate infinite many events cannot produce a lifelock in an evaluation engine.

Implementation Details

Besides the theoretical basics presented in Chapter 2 the TESSLA runtime of this thesis is built upon a number of technologies. To better understand the decisions made during the implementation this chapter will give an overview of these decisions and show why they were chosen.

As already mentioned, the implemented runtime itself is independent on the way traces are generated. Therefore we will not only look at the building blocks for the runtime itself but also examine related projects which can be used to obtain traces, which then can be monitored by the runtime. Because the format of the traces can differ heavily, depending on how and why they were collected, they are not only used to test the runtime but also to determine how the runtime can consume them.

5.1 TesslaServer

The runtime that evaluates specifications against traces is implemented in the programming language Elixir, which itself is built on top of Erlang, the Bogdan/Björn's Erlang Abstract Machine (BEAM) Virtual Machine (VM) and Open Telecom Platform (OTP). To understand why this platform was chosen we will look at the Erlang ecosystem in the next section. Note that in the following sections terminology from Chapters 2 and 3 as well as from Erlang and Elixir is extensively used.

5.1.1 Erlang and Elixir

Erlang was originally developed 1987 as a language to program systems with limited resources which had to be highly fault tolerant. The primary purpose of the language were phone switches, which have to handle large amounts of connections at the same time. Since the switches weren't deployed at a central location but wherever they are needed crashes would entail long downtimes of the system. Also, since customers expect permanent service, the platform had to provide a way to update the software without downtimes. Erlang and the OTP platform are built on top of these requirements.

While the requirements of TeslaServer are quite different we will see that the Erlang platform is a great fit for the implementation.

The rather new programming language Elixir¹ can be seen as a dialect of Erlang. Elixir code is compiled into bytecode for BEAM and can therefore interoperate with Erlang code. The rationale to use Elixir instead of Erlang directly is twofold: On the one hand Erlangs syntax is pretty different from that of most modern general purpose programming languages while Elixirs syntax was developed based on modern language design principles. Also, Elixir supports metaprogramming, meaning you can write code that generates code at compile time, a feature we use heavily, as described in Section 5.1.2.

One of the core strengths of the Erlang platform is its support to exploit multiple processor cores, even if these cores are deployed over multiple machines in a network. The platform offers tools to develop code that can be distributed over multiple processors. This distribution is transparent to the developer. The underlying concept of the distribution is the actor model, first introduced in [HBS73].

An actor is basically a self contained entity, that holds a state and can receive and send messages to other actors. Since an actor manages its own state and is the only one that can manipulate it, an actor can be scheduled on any core as long as the runtime guarantees transparent message delivery. When an actor receives a message from another actor the BEAM VM will eventually schedule the code responsible for handling the message on an available core. This code can then access the state, alter it and send a response to the sender of the message or send messages to other actors. In this sense an actor can be seen as a state machine, which alters its state everytime it receives a message. Since actors are independent of each other they can be scheduled in parallel on multiple processors. Only when two actors synchronously communicate one actor has to wait for the other.

Another reason to choose the Erlang platform was its support for multiple platforms, including resource constrained ones. While this thesis only considers offline monitoring it may be a future goal to perform online monitoring with TeslaServer. To enable this feature the runtime has to be able to run on the same hardware architecture as the monitored program sharing resources. An example of the variety of the supported platforms of Erlang and Elixir is the Nerves project² which allows developers to build embedded software.

¹<http://elixir-lang.org>

²<http://nerves-project.org>

5.1.2 Architecture

As described in Section 3.8 TESSLA specifications form a DAG of nodes, which perform transformations on streams and send their computed streams to children nodes. Streams can be seen as a sequence of events or changes that can be represented as messages between the nodes. This form of specification can be easily implemented as an actor based system, where each vertex in the DAG is implemented as an actor and the communication between adjacent vertices is realised with message passing between the corresponding actors. Interestingly, as we will see in the following, this architecture enables the enforcement of greedy schedules, which are described in Definition 22, while the Erlang runtime itself only guarantees a fair scheduling. To understand how a greedy schedule can be enforced we have to look at the internals of message passing and handling in Erlang and Elixir.

Basically there are two ways two Erlang processes can communicate with each other through messages: synchronous and asynchronous. The *call* API is used to send a message synchronously. A *call* will send a message to another process and block until a response is received. The *cast* API is the asynchronous counterpart, which will send a message and immediately passes back control to the process that used it.

As can be easily seen the usage of the synchronous API will lead to a valid greedy schedule, since the sources in the DAG will have to wait for their children to finish and the children will transitively have to wait for their children. This means that each new message gets distributed through the whole graph as fast as possible. While this behaviour is an interesting observation the actual implementation uses the asynchronous API to take better advantage of parallel evaluation.

An important characteristic of TESSLA specifications is that they can specify properties targeting realtime characteristics. On one hand this enables specifications that are not feasible with more classical specification approaches, like LTL or LOLA, which work on synchronous streams. On the other hand it adds complexity to monitoring approaches, since it adds asynchronicity to all parts of the system. One point where this is important is in the way systems have to be monitored, or more precisely how their generated events are encoded. For synchronous monitoring approaches, encoding the fact that an event happened is sufficient. However for asynchronous ones it is important to know at which exact time each event happened. For an implementation of an asynchronous monitoring approach this simply means, that the representation of events has to include information about the time at which the event happened. Another consequence of the asynchronous nature is discussed in Section 5.1.4 with the notion of the *front* of events.

```

{
  "type": "java.util.Collections$UnmodifiableSet",
3  "items": [{
    "id": 1,
    "nodetype": "TesslaServer.Node.Lifted.Add",
6    "operands": [3, 2]
  }, {
    "id": 2,
9    "nodetype": "TesslaServer.Node.Literal",
    "options": {
      "value": 5
12  }
  }, {
    "id": 3,
15    "nodetype": "TesslaServer.Node.Literal",
    "options": {
      "value": 3
18  }
  }
}]
}

```

Listing 5.1: Minimal example of the JSON based specification format. The specification describes a DAG with three nodes: two literals as the sources and an adder as their child.

5.1.3 Synthesis of the Evaluation Engine

The first step that TesslaServer has to perform to evaluate a specification is to synthesize the evaluation engine that will consume event traces and perform the corresponding computation. For this step a specification is compiled into a JavaScript Object Notation (JSON) based format that describes the nodes and their relationship. Listing 5.1 shows a minimal example of the format, which includes three nodes: two literals and a node adding their value.

The compiler performs multiple checks, including type checks and ensures that no loops are present in the specification, and transformations, namely resolve macros and other syntactic elements of the specification language. Since the compiler acts as a safety guard, TesslaServer assumes that a given specification is error free and performs no redundant safety checks. Invalid specifications can therefore lead to all kinds of wrong behaviour if fed to TesslaServer.

The JSON based specification is then translated into actors as follows. For each node described in the *items* object an actor is started with the Elixir module specified by the *nodetype* key as the message handling code. When a node is started as an actor it will receive the additional information present in the JSON object, namely the values for the *id*, *operands* and *option* keys, as arguments to its initialization

handler. The node will use this information to build up the initial state and to register itself with a central process registry provided by erlang under its *id*.

After all actors are started, TesslerServer will send each actor a message asking them to subscribe to their operands. When a node subscribes to an operand it will send the operand a message containing the nodes *id* with the request to add this node as a child. The node representing the operand will add the *id* to the list of children in its internal state. Later, during the actual evaluation of input traces, each node will use the list of children to send messages of new generated events using the central process registry.

The evaluation engine is considered to be fully synthesized when all nodes subscribe to their operands and can start to evaluate the specification over traces. The next section will explain the implementation of nodes to understand how the evaluation works.

5.1.4 Node Implementation

Nodes (or *computations* due to namespace errors with the Erlang standard library) are responsible for the actual evaluation of a specification over traces. TESSLA defines a standard set of nodes (called functions in the original specification) but leaves it open to the runtime to support only part of them or extend it. To support extensions of TESSLA and the runtime one of the main focuses was to make it easy to implement new node types.

This is achieved by building upon an abstraction from OTP called *GenServer* and providing a new abstraction which is tailored towards implementing a node for stream transformation, called *GenComputation*. Figure 5.1 shows the control flow of a single node with two inputs and one child node.

The *GenServer* abstraction is provided by the Erlang and Elixir platforms and is basically an implementation of the actor pattern mentioned above. *GenServer* provides an API to register actors, to send messages to other actors and to handle incoming messages as well as maintaining the actor state. Furthermore *GenServer* enables monitoring, crash recovery and hot code upgrades, mechanism that are not used in TesslerServer as of now.

The next layer is provided by the *GenComputation* abstraction. Before examining its responsibilities we will look at its implementation. *GenComputation*, similar to *GenServer*, heavily relies on Elixir metaprogramming, which is achieved with macros. Elixir uses macros, a mechanism mainly known from the LISP programming language. Macros enables the programmer to write code that generates further code. Since all nodes perform a similar task, performing computations on one




Fig. 5.1: Schema of the control flow of a node. A node interacts with other nodes using the *GenServer* abstraction, buffers events using the *GenComputation* abstraction and performs its calculation in the *logic* part. Computed events are then send to children using the same abstraction mechanisms.

or more input streams, it makes sense to generate code for shared responsibilities instead of duplicating this kind of code for every node. This use of macros achieves a similar goal as inheritance in object oriented programming languages, namely code reuse.

Let's now focus on the responsibilities of *GenComputation*. When a node receives a message from another node, the message will be handled by the *GenServer* abstraction. *GenServer* itself requires the user to implement a callback method which is responsible for actually handling a received message by updating the state or sending new messages as required. Since TesslerServer uses exactly one format for messages *GenComputation* is able to handle them for all actual node implementations in a general way and will only invoke actual node logic when needed. The concrete handling of messages differs between the two versions of TesslerServer whose details are described in Sections 5.1.5 and 5.1.6.

On a high level *GenComputation* stores messages in a buffer until it can be sure that all inputs have progressed since the last time outputs were generated by the node, which means that either new events can be generated or at least the output stream can be progressed. A similar concept is presented in [Hal16] which we discussed in Section 2.3.2 called the *front* of the input streams. Because BeepBeep is not using timed specifications, the computation of the front is easier: the head of each input when every input has at least one event buffered. For TESSLA the computation of the *front* is not so easy, because the timestamps of the events in the front can differ. As a result, *GenComputation* has to perform multiple steps to determine the appropriate actions to take based on the events of the front. The exact steps are described in the respective sections of the two versions of TesslerServer. When *GenComputation* has determined that at least one new transformation can happen it invokes the actual node logic by invoking a callback method that each node has to implement.

5.1.5 TesslerServer V1: Stream passing

The first version of TesslerServer was built with two goals in mind: safety and hiding complexity. This led to implementation decisions that had had big performance impacts (see Section 6.1) and made it hard to implement complex node types.

One of the main ideas of this implementation is to make streams a central data structure that is able to guarantee some safety aspects like the ordering of events. Each node maintains a number of streams: one for each parent node and one output stream. Whenever a node computes new events, its output stream is updated and the whole updated stream is send to all of its children. When a node receives a message from a parent node containing an updated stream, the node will update

its own state by replacing the last saved stream from the parent with the newly received. Then the node will determine if with this updated stream new events can be computed. To do so the node looks at the input stream with the minimal progress and compare it with the progress of its own output. If these two timestamps are equal, the node cannot produce a new front, since at least one input has not progressed since the last computation. Only if the minimal progress of all inputs is bigger than the progress of the output the node has to perform the appropriate computations.

This is implemented by generating a sequence of partial fronts for specific timestamps as follows:

1. Look at all input streams and find all events that happened after the progress of the output upto the minimal progress.
2. Take the timestamps of these events in chronological order. We call these timestamps the *change timestamps* since they denote that at least on one input stream something has changed.
3. Iterate over the timestamps in order and build partial fronts by getting all events that happened on any input stream at that timestamp.
4. Invoke the actual logic of the node for each partial front to perform the corresponding transformation.
5. Add the generated events to the output stream and send the updated output stream to all children.

It is important to understand that all steps except Step 4 are performed by the *GenComputation* abstraction. Hence, a node implementation - at least in theory - only has to implement the logic to combine a partial front to a new output.

The problems of this approach are twofold. First, to implement more complex node types it was necessary to overwrite a lot of the provided abstractions, for example to manipulate timestamps. But more important were scalability issues. Since every node stores a copy of all input streams in its state and streams contained all events produced, the random access memory (RAM) usage grows exponentially with the number of nodes and input events used to evaluate a specification. This can be seen in Section 6.1.2. Another problem is that the messages between nodes also grow with the number of events, since the whole stream is sent every time.

To alliviate these issues the TesslaServer v2 architecture was designed.

5.1.6 TesslaServer V2: Event passing

The second version used the insights of the first version to provide better abstractions. Scalability and handling complex node types were the main goals of the new architecture.

To achieve these goals some changes has to be made. Streams are no longer an explicit data structure in the system but mere an attribute of events to denote on which stream they happened. The new architecture of *GenComputation* achieves a simpler and clearer control flow of nodes and a very small API at the cost of some safety guarantees, which explicit streams provided.

In the new architecture simple node types implement more logic, since they have to decide how to handle progress events. This is necessary in the new architecture to propagate that a stream has progressed to a new timestamp without an event happening on it. In the old architecture the *GenComputation* abstraction handled these cases for all nodes, which was not appropriate for some node types. To avoid too much code duplication the new architecture provides a new abstraction *GenLifted*. *GenLifted* can be used as the building block for node types that *lift* a function - which normally would run on two values - to run on two signals. This approach avoids the problems of the old architecture by moving concerns out of the base *GenComputation* abstraction and making it optional to use the new *GenLifted* abstraction.

The new approach to sending messages between nodes is to send each generated event as one message. This will lead to an overall increase in messages but simplifies the handling of each individual message. In the new architecture, nodes contain a buffer for each parent node as part of their state. A node saves all events received from that parent, that weren't part of a front upto that point in time, in this buffer. The process of handling new messages and computing the partial fronts implemented in the new version of *GenComputation* is the following:

1. Add the newly received event to the end of the buffer that stores events of that parent node.
2. Test if on each buffer at least one event is stored.
 - End if at least one buffer has no events.
3. Else determine the minimal timestamp over the first events of all buffers.
4. Remove all events from the head of the buffers with that exact timestamp, as this set of events form a partial front.
5. Invoke the actual logic of the node for the partial front to perform its transformation.

6. This will generate at least a new progress event or one or more normal events: send these events to all children of the node.
7. Go to Step 2.

Note again how only at Step 5 the actual node logic is invoked, meaning only that part has to be implemented for each new node type. Nonetheless this procedure adds more responsibilities to the programmer of new node types. In the old approach the actual node logic did not have to handle progress events and caching of events that are important for future computations. This can be described as the concept of making complexity explicit. Because the implementation has to actually handle complex edge cases in contrast to the old approach which tried to hide this complexity.

One side effect of the new implementation is that one limitation on input traces is no longer needed. In the first implementation traces had to be totally ordered over all streams, but the new implementation works as long as traces are ordered per stream. This is especially useful when using traces that were generated by multithreaded systems. In the systems where one can assume that each stream in the trace is exclusive for one thread the generated trace file can be directly fed to `TesslaServer`. This can easily be achieved by including a unique identifier per thread in the stream identifier.

5.2 Instrumentation Pass

While the implementation of the actual runtime was the main goal of this thesis another project was also developed, mainly to provide test traces to the runtime. We describe the used technologies in this chapter, since the project uses some interesting technologies and can be extended to support a wide variety of trace data generation.

When no suitable test data for the runtime could be found the need for a tool to generate traces arised. As reasoned in Section 2.5 all trace data that was available was not suitable for `TESSLA` for a number of different reasons. Therefore a tool was implemented to generate traces tailored towards evaluating `TESSLA`. This did open up the opportunity for the runtime as a trace collection tool. One of the central ideas of the runtime is, that it does not make assumptions about the platform of the monitored program. While code instrumentation with the goal to emit a trace obviously relies on the language in which the code is written, the LLVM project provides abstractions that can be used to implement an instrumentation mechanism that does not rely on the language the instrumented code was written in.

```
variable_values:write_ptr 843071489 1463991050 176761
variable_values:write_ptr 843071490 1463991050 176832
function_calls:process nil 1463991050 176901
```

Listing 5.2: Trace data generated by the instrumentation pass. Each line represents an event and each event consists of four pieces separated by spaces: a stream name, an optional value and the timestamp in unix format followed by the amount of microseconds since that timestamp.

To understand how this is possible recall Section 2.5.6 and the way LLVM works. A frontend compiles code from a source language into an IR, performs compiler passes on the IR and then finally compiles into native machine code. If the instrumentation pass works at the level of the IR it would work for all languages that have a frontend for LLVM.

It is easy to implement such an instrumentation with the provided C++ API from LLVM to implement compiler passes that can analyse and transform IR code. The implementation uses the *ModulePass*³ base class to analyse whole modules. Modules represent constructs like classes from C and C++ in IR. A module consists of *Functions*⁴. On invocation implemented *ModulePass* iterates over the *functions* and checks if the current function is one that should be instrumented. If so it builds a *CallInst*⁵, which is the IR equivalent of calling a function. The *CallInst* will call a specified log method provided by a logging library. The instruction is then inserted as the first instruction of the *Function* that is instrumented.

At runtime the instrumented program will then log an event everytime the instrumented function is called. Events generated by the instrumented program will have the format shown in Listing 5.2.

Our implementation of the compiler pass works and was used to generate authentic test data, but it is limited and has some remaining challenges. As we will see in Section 6.2 the pass can add a lot of overhead and interfere with compiler optimizations. As a first measure to reduce the overhead the logging mechanism is implemented on top of a library called *zlog*⁶ to buffer generated events. Without this buffer the overhead would be much higher since constant writes of the generated events to an output device would have to happen.

A second measure is that the pass will only instrument functions that are specified by the user when invoking the pass. This enables the generation of multiple different instrumented versions of the same code, of which each will only generate traces that are interesting for a certain specification. For example one could have

³http://llvm.org/docs/doxygen/html/classllvm_1_1ModulePass.html

⁴http://llvm.org/docs/doxygen/html/classllvm_1_1Function.html

⁵http://llvm.org/docs/doxygen/html/classllvm_1_1CallInst.html

⁶<https://github.com/HardySimpson/zlog>

two instrumented versions of the same code, one which generates events used by a specification to monitor buffer size limits and another which generates events used by a specification describing performance constraints.

As of now the instrumentation will also produce erroneous traces when the instrumented program is multithreaded and will generate events on more than one stream. Due to race conditions the logged events can then be in the wrong order with respect to their timestamps. For offline monitoring this can be solved by simply reordering the events, but for online monitoring this is not possible. For the second version of TessaServer the problem only occurs when different threads can generate events with the same stream, since this version only requires the inputs to be totally ordered over the same input stream. This can be solved by adding a thread unique identifier to the stream name and adapting the specifications to take into account that events might happen on different threads out of order.

As a final limitation the current instrumentation pass only supports function calls to be instrumented. To do so it would have been enough to implement a *FunctionPass*⁷. A *ModulePass* was chosen because the pass could be easily expanded to generate other events as well, for example function returns, variable definitions, variable overwrites and assignments of null values.

⁷http://llvm.org/docs/doxygen/html/classllvm_1_1FunctionPass.html

Evaluation

In this chapter performance characteristics of the implemented runtime and also of the instrumentation pass are presented. To do so both systems are tested under a number of different circumstances to get an overview of how they might perform in real world usage. As the complete benchmark data is too big in this chapter only summaries are presented. The complete datasets can be found in Chapter 8.

6.1 Runtime Benchmarks

Over the course of this thesis two different approaches to handling the progress of streams in the runtime are explored, which were also briefly mentioned in Section 3.12: progress timestamps and progress events. Progress timestamps were used first, which worked, but led to big overhead when a great number of computations had to happen, since the whole stream had to be sent to all children of the node that performed the computation. Therefore, later the runtime was refactored to use progress events, so that only events had to be sent between nodes. The following benchmarks will show the performance characteristics of both approaches under different circumstances.

6.1.1 Number of Processors

Since one of the main motivations to choose Erlang as the platform for the runtime was its support for parallel execution on many processor cores, we explore the performance characteristic in regard to available cores in this section. To do so a simple TESSLA specification is given, that includes as many nodes as the maximal number of processors available. Such a spec for eight processors is given in Listing 6.1.

The specification is evaluated over appropriate traces with a specified amount of processor cores available. The needed time of the evaluation engine from its start until it emits the conclusion, that the stream *done* is *true*, is measured for different amounts of processors. All benchmarks were performed on the same machine with up to 16 cores and 48GB of RAM. Figure 6.1 shows the results of the benchmarks. The complete data can be found in Tables 8.1 to 8.1.

```

define num_events: Signal<Int> := literal(10000)

3 define add_calls: Events<Unit> := function_calls("add")
  define add_call_sum: Signal<Int> := eventCount(add_calls)

6 define overhead_0: Signal<Int> := signalAbs(add_call_sum)
  define overhead_1: Signal<Int> := signalAbs(overhead_0)
  define overhead_2: Signal<Int> := signalAbs(overhead_1)
9 define overhead_3: Signal<Int> := signalAbs(overhead_2)
  define overhead_4: Signal<Int> := signalAbs(overhead_3)

12 define done Signal<Boolean> := eq(overhead_0, num_events)

```

Listing 6.1: TESSLA specification with eight nodes on the critical path




Fig. 6.1: Performance of the runtime with different number of used cores. V1 denotes the old implementation, V1.5 the adaption with limited amount of saved events and V2 the new approach.

One thing that was recognized during benchmarking was that the older implementation approach used big amounts of RAM and the usage grows exponentially with the amount of cores used. For example, with two used cores the runtime would use around 2GB RAM, with 4 Cores already over 8GB. This can be explained easily: in the old implementation all data has to be sent between all cores while in the new implementation only the generated events have to be sent around. The excessive usage of RAM even lead to timeout under some circumstances and crashes the whole runtime. The amount of data send between processes can even lead to negative performance impacts when more cores are added.

This behaviour was one of the main reasons to switch to a new approach. To test if reduced RAM usage would lead to better performance and eliminate crashes, the old architecture was changed in a small way: Each stream was limited to only hold the last 20 events, all older events were thrown away. Since the specification used for benchmarking only works on the most recent event of each stream, this would not alter the conclusion of the engine. For example a specification computing the average of the last 21 events wouldn't work with this adaption. See Section 6.1.2 for a benchmark of RAM usage. The obvious reduction of RAM usage of the adapted first approach motivated the decision to refactor the runtime to use the new approach.

Figure 6.1 shows how the adapted first approach and second approach are both performing much better than the original one. The graph of the adapted first version shows, that it suffers from a too high number of cores. It is suspected that this comes from the fact that the messages between the processes are still quite big. The new version does not only has the best performance with a single core, it also scales with more cores and doesn't suffer under a high number of cores. While not tested we expect the new approach to scale very well with even more cores, especially for more complex specifications and more input events.

6.1.2 Number of Events

Another metric looked at, is how the runtime behaves in regard to the number of events that are fed to the engine, and therefore how many events are generated during its execution. To obtain these metrics the specification from Listing 6.1 is evaluated with different values for the comparison node are run. The specification is extended to use 16 nodes, which means that each input event will generate 16 internal events. All benchmarks in this section were performed on a processor with 4 cores at 2.4GHz speed and 8GB of RAM.

The first benchmark compares the execution time of the implementation approaches when fed with different number of input events. Figure 6.2 gives an overview of the data, the whole dataset can be found in Tables 8.4 to 8.6.




Fig. 6.2: Average execution time of the different implementations when fed with different amounts of input events. V1 refers to the old approach, V1.5 to the adapted old approach and V2 to the new approach. V1 has no data for 10 000 events since the enormous RAM usage constantly lead to timeouts.

The numbers clearly show that the first version has big scaling problems when the number of input events grows. The adapted approach shows, that the reduction of size of the messages is an important factor to improve the scalability. Version 2 does perform even better than the adapted approach and also scales very well with big numbers of events.

The number of events also increase memory usage, espacially in the first implementation approach. The second benchmark therefore tests the RAM usage of the different approaches under different number of input events. Figure 6.3 illustrates the RAM usage of the different approaches. The complete data can be found in Tables 8.7 to 8.9.

Again it can be seen how the old version doesn't scale well and uses an exponential amount of memory with a growing number of events. The adapted and the new version both require a very small memory overhead and are nearly constant in their memory consumption. The new version uses a bit more memory, we suspect this comes from the fact that more messages have to be sent between the processes and each message carries a bit of overhead.

6.1.3 Number of Nodes

As a last metric the performance of the runtime is evaluated over specifications with different amount of nodes. Therefore the specification from Listing 6.1 was modified to include higher number of nodes by appending more nodes computing the absolute. All benchmarks in this section were performed on a processor with 4




Fig. 6.3: RAM usage of the different versions of the runtime in regard to input events. V1 refers to the old approach, V1.5 to the adapted old approach and V2 to the new approach. V1 has no data for 10 000 events since the enormous RAM usage constantly lead to timeouts.

cores at 2.4GHz speed and 8GB of RAM. Figure 6.4 visualizes the performance, the complete data can be found in Tables 8.10 to 8.12.

While there is no real world data of the size of typical TESSLA specifications, 128 nodes seems to be a rational upper bound. As the data shows that all version of the runtime scale nicely with complex specification, the data even suggests a somewhat logarithmic increase of the execution time in regard to node count. It can also be seen that each version of the runtime has better performance than the previous one.




Fig. 6.4: Average execution time of the runtime over specifications with different amount of nodes. All specifications did count the number of function calls and stopped when it reached 1000. V1 refers to the old approach, V1.5 to the adapted old approach and V2 to the new approach.

6.2 Instrumentation Benchmarks

After the evaluation of the runtime itself it remains to evaluate the C instrumentation program. Note that the instrumentation is mostly a proof of concept and its main purpose for now was to generate test data for the runtime. But it seems feasible that it can be optimized and extended to become a general purpose trace collection tool, therefore some benchmarks were performed.

For reliable trace collection of software the performance impact of the instrumentation is important. To measure this a trivial C program was exercised with and without instrumentation. The program increments each integer from 0 to 100 000 000 by one and, if the incremented number is divisible by a given parameter c , adds it to an intermediate result. Note that the program can and will perform some integer overflows. The code is shown in Listing 6.2.

The code is then instrumented to log each call of the add method. For each benchmark the compiled program was run 50 times. All benchmarks in this section were performed on an Intel Core i5 with four cores at a clock speed of 2.4GHz and 8GB of Ram.

```
#include <stdio.h>

3 int intermediate = 0;

    int inc(int input) {
6     return input + 1;
    }

9 void add(int input) {
    // Instrumentation Point
    intermediate += input;
12 }

    int main() {
15     for(int i = 0; i < 100000000; i++) {
        int result = inc(i);
        if(result % 100 == 0) {
18             add(result);
        }
    }
21     printf("%i", intermediate);
    return 0;
}
```

Listing 6.2: Example C program for benchmark purposes




Fig. 6.5: Performance Comparison of an example C program with and without instrumentation.

6.2.1 Performance Comparison with non Instrumented Code and Compiler Optimizations

One interesting metric is the intrusiveness of the instrumentation, describing how an instrumented program performs in contrast to the same program without the added instrumentation. For this the parameter c of the example program is set to 100, so that around 1% of all function calls are instrumented.

One thing that can be recognized is that the impact of instrumentation is not predictable when using compiler optimizations. Aggressive compiler optimizations can remove function calls or inline values if the optimization doesn't change the program behaviour. When instrumentation calls are added no such optimization can happen, since otherwise the behaviour of the program would be altered, namely the logging would be removed when the function calls are removed.

Therefore we compare the instrumented and non-instrumented code for all optimization levels. The comparison can be seen in Figure 6.5. The complete data can be found in Tables 8.13 and 8.14.

The choice to instrument around 1% of the function calls was made arbitrary after some experimentation showed that a higher amount leads to huge performance impacts as shown in Section 6.2.2. Intuitively it sounds reasonable that in a real life example only a small subset of the function calls are interesting for trace generation, since TESSLA specifications should be used to monitor only critical parts of a system.




Fig. 6.6: Execution time of a program with different amount of calls to an instrumented function

6.2.2 Performance Impact in Regard to Instrumentation Percentage

To further investigate the impact of instrumentation we will look at the performance impact with respect to the percentage of function calls that are instrumented. Therefore the variable c of the program is changed, which leads to different amount of calls to the instrumented function. For all values the program was compiled with the maximal optimization setting. In Figure 6.6 the results of changing the parameter can be seen. The complete dataset can be found in Table 8.15.

Note that the x -axis is logarithmic. As expected the performance impact scales directly with the amount of calls to the instrumented function. At some point no more speed up will happen, since no calls to the instrumented functions will be made.

6.3 Practical Examples

Additionally to the theoretic benchmarks it was also important to evaluate the run-time against some more practical examples and test a bigger amount of node types. While the specifications for the benchmarks from the previous section did use some

node types and showed that these types perform the correct computations (else the benchmarks wouldn't have worked), they only used a subset of all available types. Thus the following examples concentrate on a wider range of different nodes.

To do so traces from real world examples were collected, either with the developed instrumentation pass or by adjusting existing logs to fit the needs of TESSLA. The programs or traces were then modified to deliberately include errors to test if appropriate TESSLA specifications will find them when evaluated in the runtime.

As the main source for test trace data a ringbuffer, implemented in C, was modified by the instrumentation to log specific events. The code of the ringbuffer can be found in Listing 8.1. Events that are logged are of two types: calls to the process function and changes of the `write_ptr` variable. The generated traces can be used for a variety of specifications, including performance tests, race conditions or wrong initializations.

The Following specifications were tested over the traces:

- Data is only processed when it is available, meaning that at each point of the trace the amount of changes of `write_ptr` is bigger than the number of calls to process.
- The given buffer size of five is never exceeded, meaning the difference of the amount of changes of both event types is never bigger than five.
- Everytime new data is added, there is a call to process within the next milisecond.

The specifications were intentionally chosen and written in such a way, that they use a big amount of different nodes. The actual TESSLA specifications used to describe the requirements can be found in Section 8.4. For each of these specification the traces were deliberately altered to violate them. Evaluation of the specifications over the traces with TesslerServer always led to the correct conclusions.

Conclusion

In this chapter we will summarize the rest of the thesis and describe areas of future work. In the final section of this chapter we highlight some challenges that can be attached based on the work that is already done.

7.1 TesslaServer

As the evaluation in Sections 6.1 and 6.3 shows the implemented runtime is able to evaluate a variety of specifications in a performant way. The new architecture for the runtime improved the performance in many ways, especially enabling the evaluation of complex specifications with many input events. Nonetheless, the current version of the runtime is a proof of concept developed for this thesis and has not been tested under real life circumstances or with actual systems.

At the moment there are two drawbacks of the runtime, that are based on limitations of the current version of the TESSLA compiler. These limitations are missing type information and missing output information.

Let us first describe the lack of type information and its impact on the runtime. The TESSLA compiler is able to infer a lot of type information, therefore making it possible to omit type information from TESSLA specifications. Many functions in TESSLA are generic functions over the type they consume or produce. Recall that each function in TESSLA is generic, since functions work either on signals or eventstreams which carry an underlying type themselves, so a function could have for example the output type `Events<Int>` or `Signal<Boolean>`. Furthermore a function can be generic not only over the type of the values of its streams but also on what kind of stream (signal or eventstream) it consumes or produces. Additionally the same function can be defined with different arities, for example an *add* function could be defined for 2, 3 or any number of arguments (note that TESSLA does not support yet methods with a varying number of arguments, in some languages known as *varargs*).

This genericity leads to a situation where for example the *delay* function, which delays stream by an amount, has actually three generic types:

- `Events<A>, time -> Events<A>`
- `Signal<A>, time, A -> Events<a>`
- `Events<A>, Int -> Events<A>`.

The first and second functions are delaying a stream by a certain time, the second one includes a default value as the third argument, and the last function delays the events on an eventstream by a given amount of steps, meaning the first event will be delayed to the timestamp of the second event.

The compiler is now able to match the signatures of a function to the correct instance and will check if the whole specification is correct with respect to the type system. The problem is in the output of the compiler which does not contain type informations anymore, meaning that the representation of the first and third *delay* functions are exactly the same and the runtime would have to figure out which concrete type of the function to use. We propose two workarounds for the runtime and the compiler until the type informations are included in the output.

The first workaround is for functions that work with signals and eventstreams and which have the same arity. To make these functions work without type information the functions are split into multiple functions, one for each argument, that can take either a signal or an eventstream. For example the two *delay* functions with the same arity are split into a *delayEventsByTime* and a *delayEventsByCount* function. This moves some responsibility to the user writing a spec since the user has to use the right function and cannot rely on the compiler to figure out the correct version. This workaround leads to an increased implementation effort for the runtime, since more node types have to be implemented. The second workaround is much simpler: functions which work on streams with generic type of values (e.g. `Signal<A>`) will use `Int` as the default type and will not support any other types.

The best solution to the problem is, that the compiler includes the type signature of functions in its output. The signature could then be used at runtime to dynamically convert the values of events to the corresponding matching types.

The problem of the lack of output information is easier to understand and solve. TESSLA specifications denote that a stream is an output stream using the `out` keyword. At the moment the compiler is not handling this keyword at all and therefore the information, which streams are outputs, is missing in the output. The workaround for now is to specify that output nodes when the runtime is started. A

typical invocation of `TesslaServer` looks like `./tessla_server example.spec -o 4:error` where the `-o` flag denotes an output, 4 is the id of the node that generates the output stream and `error` is a name that can be chosen by the user.

A last missing functionality we want to highlight is based on the asynchronous nature of the runtime. Since nodes are independent of each other and share no knowledge, except the one that is explicitly exchanged by messages, there is no global state that can be used to track if all nodes have finished their computation. Furthermore there is no shared progress between independent nodes in the new architecture, or in other words each source will be progressed to a different timestamp.

This leads to a problem of detecting when the trace actually ends: Consider a specification with two independent sources and a trace which contains events on the stream of the first source upto timestamp t_2 but on the stream of the second source only to the timestamp t_1 , with $t_1 < t_2$. When the trace ends the second source will have only progressed to t_1 while the trace implicitly contains the information that there is no event happening on its input stream at any time later than t_1 (and upto t_2). This information could lead to new generated events of a child node, for example a *timeout* node which could then lead to new output events.

A simple workaround would be to implement a *flush* mechanism, which will send each source node a message that there will be no more input events and therefore they can progress to the maximal timestamp. With this mechanism all children of sources will then receive the information that their inputs have progressed without generating new events, enabling these nodes to progress to the maximal time. The information that all inputs have finished will transitively flow through the whole system and eventually all nodes will have performed their final computation and the system will have generated all possible outputs.

7.2 Instrumentation Pass

The benchmarks reveal that the instrumentation adds significant overhead, especially when compiler optimizations are turned on or the percentage of instrumented function calls is large. While the overhead is large, it is stable (see the standard deviation in Section 6.2.1).

The current instrumentation should be seen as a proof of concept. The generated traces should be mostly used to analyse test settings where compiler optimizations are turned off. When used for timing specifications, the instrumented code can be benchmarked against a non instrumented version of the code and the results can be used to transform the actual timing requirements to the corresponding ones

for the instrumented code. This will be an approximation of the actual results. Also it is recommended to write small and specific TESSLA specifications, so that the instrumentation only has to be applied to a small subset of functions. Obviously it is strongly discouraged to use the instrumented code in any production setting.

Two enhancements could be easily integrated into the instrumentation pass:

1. Enabling the logging of more event types, like variable reads or memory allocation, while also including data in the events.
2. Introducing a better configuration format to specify what events should be logged.

As explained in Section 5.2 the compiler pass is instrumented as a *ModulePass*, enabling the logging of all kinds of events. While the logging of function calls was sufficient for testing and evaluation purposes this restriction together with the total lack of logging data is an obvious obstacle for more serious use cases. The trace format already includes a field for data values and the instrumentation pass has the ability to extract data, for example function arguments or return values, from the source code it is transforming.

Also the expansion to more event types is possible harnessing the *InstVisitor*¹ mechanism. An *InstVisitor* can be used to abstract the actual iteration over the building blocks of the IR of a program, so that the user only has to specify the transformation that should be happening when a specific instruction type is seen. For example the instrumentation of return statements could use the *visitReturnInst* API to handle all return statements and perform a log event.

In the context of adding more event types to the instrumentation pass the configuration format becomes a problem. For now a user specifies each function that should be instrumented when calling the pass with an option to the compiler pass like `-instrument FUNCTIONNAME`. This approach obviously does not scale when multiple event types are implemented. Therefore we advocate for the design of a common configuration format that is processed by the instrumentation pass. This approach would also enable the sharing of configuration files between different developers and the versioning of them using version control systems.

7.3 Further Work

In addition to the adjustments mentioned in the two previous sections there are also some ideas to extend the works reported in this thesis in ways that require more theoretic work and larger changes of the architecture. This section will give

¹http://llvm.org/docs/doxygen/html/classllvm_1_1InstVisitor.html

an overview of these ideas and will present sketches of approaches to implement them.

7.3.1 Implementation for the JVM

While the Erlang platform has been tested in production for many years and provides a great performance for our implementation, there is a certain interest to provide an implementation on top of the JVM. The main reason for this is that the TESSLA compiler itself is implemented in Scala and therefore already uses the JVM. An implementation using the same platform would enable the distribution of a single executable that can be used to compile and evaluate TESSLA specifications altogether. Also a runtime on the JVM could interact with monitored programs that are written in a JVM language and therefore eliminate the need for an instrumentation pass.

The Akka project² provides a Scala implementation of the actor model. Since the runtime relies heavily on the actor model, Akka might enable to port the runtime to scala in a very straightforward way. As a side note there is also an implementation of Akka on top of the .NET platform³ which might be used to additionally port the runtime to yet another platform.

7.3.2 Composition of Evaluation Engines

Another idea, motivated by the CEP field, is the composition of multiple evaluation engines. This is somewhat contrary to the RV approach, since RV is aimed to answer the question of whether a run of a system satisfies a given property. Therefore a specification in TESSLA that is used for RV purposes would have exactly one output: a boolean stream denoting if the property is satisfied at a given time.

The composition of evaluation engines would only make sense if the output of specifications would be streams that are interesting as the input to a new evaluation engine. As it turns out, there are no restrictions applied to output streams, neither by the TESSLA specification language nor the runtime.

This means that it should be possible to combine multiple evaluation engines with only small adjustments to the output format of the runtime.

²<http://akka.io>

³<http://getakka.net>

7.3.3 Parameterized Streams

A quite common model in RV are parameterized streams, where one parameterized stream represents a dynamic number of actual streams. A very common example to highlight why parameterization is required is the specification that a user cannot have more than three failed login attempts in a row. When the number of users is not known it is not possible to write such a specification without parameterization. A parameterized version of the specification would generate a signal for each user holding the value of failed log in attempts in a row.

The common approach to handle parameterized specifications, called *slicing*, is discussed in [CR09]. The basic idea of *slicing* is to partition a stream of events, where each event holds a data value, into multiple streams with respect to that data. In the context of the example specification restricting failed consecutive login attempts the following would happen: An input stream `login_attempt(int, bool)` denotes that a user with a given id (the first value of the stream) tries to login, where the second value denotes if the login was succesful. This input stream would be partitioned, or *sliced*, by the first argument, generating a variable number of streams `login_attempt_id(bool)`. Each of these streams could then be evaluated by the TESSLA runtime.

While streams in TESSLA are already carrying data, only a restricted form of generating new streams based on this data is possible. This limited slicing can be performed for example using the *ifThen* function, which can be used to get only that events of a stream that cary a specified value. But since the real expressiveness of parameterization is based on the dynamic nature of the input data that approach is not sufficient. Therefore, to allow parameterized streams, the TESSLA specification language would have to be extended.

The changes required to allow slicing in the runtime are not very complex. Basically, a *slice* node would have to be implemented with one input stream and a variable number of output streams. Whenever the node receives an event which carries a value that has not been seen before the node would have to start new children nodes that are responsible to handle the new stream. Since nodes are already generated at runtime, when the evaluation engine is synthesized, it should be no problem to start new nodes during the evaluation process.

7.3.4 Easier Definition of New Node Types

A problem that is still present in the current architecture is how complex it is to implement new node types. First of all the TESSLA compiler would have to be changed to include the signature of the new function and to support the mapping

of the function to its JSON representation. Then, a corresponding node has to be implemented for the runtime. While the *GenLifted* abstraction provides a starting point for a generic generation of new node types it still requires manual work to implement new node types.

To improve this situation TESSLA could be extended with a mechanism to specify new functions without the requirement to explicitly implement these functions in the compiler and the runtime. A somewhat related technique that is already present in TESSLA are macros which can be used in a specification to compose existing functions into new ones. While this feature enables the user to write more expressive specifications by leveraging common subexpressions, this feature does not support the generation of new functions.

One approach that looks promising is presented in [Hal16] in the context of the BeepBeep tool. In BeepBeep computations can be generated by combining two base parts, *processors* and *functions*. A processor is responsible to apply a transformation on one or more streams while functions specify how the exact transformations look like. As an example consider the following two processors and their combination with one function to perform different transformations.

Let the example processors be called *Combiner* and *Aggregator*. The *Combiner* processor is responsible for combining two input streams by applying a function for each front and the *Aggregator*, which computes an aggregation over one stream by applying a function to each new event combined with the last computed value. Each function that takes two input values and generates one output can be combined with these processors. The $+$ function when combined with a *Combiner* would add the values of two streams at each position while it would compute the sum of the values of all events on one stream when combined with an *Aggregator*.

This approach could lead to a much smaller implementation effort for the runtime, but how exactly it could be integrated into TESSLA remains to be discussed.

7.3.5 Runtime Optimizations

Section 6.1 shows that the implemented runtime is able to perform evaluation of TESSLA specifications in an efficient and scalable way. Nonetheless some additional optimizations can be performed.

As explained in Section 5.1.6 the new architecture generates more events and messages than the old approach. While this leads to a more intuitive evaluation model and better abstraction it carries some overhead. A possible optimization is to bundle multiple events into a single message when possible. For example, a node could generate only one message containing all produced events whenever it is scheduled.

Currently a node generates a message for each partial front. This approach allows also to remove unnecessary progress events.

The following optimization is more complex and is specific to some node types. Think of an *ifThenElse* node that has three input signals and one output signal. The output signal will carry the value of the second input when the first input has a *true* value and the value of the third input otherwise. When the node has events buffered for the first and second input upto a timestamp t_1 and none for the third input the node will perform no computation until it receives an event for the third input because of the way that the *GenComputation* abstraction works. But the node could actually perform a computation if the buffered event for the first input has a *true* value and generate events upto the timestamp t_1 .

More generally stated the problem is that the progress of the output of a node is always equal or smaller to the minimal progress of all of the nodes inputs. While this is indeed necessary for most node types, there are important exceptions as we just illustrated. The added complexity necessary to handle such cases did not justify its implementation in the current context.

7.3.6 Error prevention

RV can be used not only to detect errors but also to react when an error is detected. Interweaved monitors can influence program behaviour when errors are detected, for example it could prevent cascading errors or notify a third party that can then take appropriate actions in response to the error. As the TESSLA runtime is implemented as a standalone monitor that shares nothing with the monitored program such a behaviour is harder to achieve.

While the reimplementing of the runtime on another platform as described in Section 7.3.1 is a possible solution to that problem we also want to show another approach. To enable communication between the monitored program and TessaServer the monitored code has to be changed, but we solved this problem with the instrumentation pass. The instrumentation pass changes a program to generate traces, which can be interpreted as a way to communicate with the monitor.

The instrumentation pass could be reworked so that it not only changes a program to emit traces but also to include a way to receive feedback messages from an external program and react to these messages.

Appendix

8.1 Runtime Benchmark Data

This section includes all data from the benchmarks of TesslaServer.

8.1.1 Execution Time in Regard to Used Processor Cores

Tab. 8.1: Execution time of multiple runs of TesslaServer V1 with 10 000 input events with different number of used processor cores.

Run #	1	2	4	8	16
1	38.2684 s	95.8390 s	33.2289 s	18.3235 s	21.9672 s
2	38.9833 s	98.9827 s	32.3660 s	17.9681 s	19.4063 s
3	38.1630 s	77.4360 s	36.0462 s	17.0076 s	27.6495 s
5	38.4424 s	222.9723 s	29.4773 s	17.2925 s	17.3606 s
6	39.2239 s	75.5560 s	32.1869 s	15.3906 s	19.1662 s
7	38.7815 s	81.0928 s	44.7217 s	18.3882 s	21.9974 s
8	38.7514 s	119.1012 s	38.0911 s	18.0104 s	20.2454 s
9	37.8491 s	60.2407 s	35.0671 s	17.9517 s	17.2825 s
10	38.2721 s	83.5956 s	37.5839 s	17.1108 s	15.5887 s
11	38.5657 s	105.2514 s	43.8580 s	15.8287 s	22.0196 s
12	38.1699 s	75.6985 s	32.4470 s	15.7641 s	18.3230 s
13	37.9775 s	104.0994 s	37.0281 s	24.5554 s	17.0857 s
14	38.8215 s	75.0534 s	34.8001 s	14.3733 s	19.5598 s
15	38.3913 s	74.5584 s	32.7344 s	17.9163 s	17.3179 s
16	40.4536 s	111.3214 s	35.7434 s	16.6162 s	20.2855 s
17	39.0735 s	87.0386 s	35.6367 s	15.5962 s	21.7681 s
18	38.8830 s	114.2697 s	33.4537 s	28.1556 s	17.2382 s
19	39.0799 s	103.1956 s	32.7134 s	17.7034 s	23.5276 s
20	39.3866 s	117.2423 s	39.1457 s	15.6895 s	18.4146 s

Tab. 8.2: Execution time of multiple runs of TesslaServer V1.5 with 10 000 input events with different number of used processor cores

Run #	1	2	4	8	16
1	6.3722 s	0.1256 s	0.1309 s	0.1378 s	2.6756 s
2	6.2417 s	0.1300 s	0.1333 s	0.1444 s	2.6292 s
3	6.2953 s	0.1301 s	0.1325 s	0.1379 s	2.6453 s
5	6.5148 s	0.1433 s	0.1315 s	0.1451 s	2.5836 s
6	6.6828 s	0.1347 s	0.1480 s	0.1468 s	2.6670 s
7	6.8090 s	0.1386 s	0.1294 s	0.1414 s	2.6671 s
8	6.3252 s	0.1307 s	0.1330 s	0.1385 s	2.7755 s
9	6.0975 s	0.1304 s	0.1339 s	0.1362 s	2.6097 s
10	6.1274 s	0.1328 s	0.1325 s	0.1371 s	2.6335 s
11	5.7481 s	0.1479 s	0.1299 s	0.1398 s	2.6477 s
12	6.0317 s	0.1318 s	0.1435 s	0.1311 s	2.6343 s
13	6.1565 s	0.1293 s	0.1260 s	0.1361 s	2.6603 s
14	6.3439 s	0.1223 s	0.1309 s	0.1313 s	2.5560 s
15	6.3749 s	0.1318 s	0.1335 s	0.1338 s	2.6739 s
16	6.3289 s	0.1321 s	0.1323 s	0.1503 s	2.5466 s
17	6.3894 s	0.1253 s	0.1323 s	0.1397 s	2.6240 s
18	6.3943 s	0.1288 s	0.1320 s	0.1316 s	2.6939 s
19	6.4089 s	0.1319 s	0.1347 s	0.1302 s	2.5940 s
20	6.3865 s	0.1698 s	0.1336 s	0.1312 s	2.6347 s

Tab. 8.3: Execution time of multiple runs of TesslaServer V2 with 10 000 input events with different number of used processor cores

Run #	1	2	4	8	16
1	3.620s	1.598s	1.024s	0.948s	1.010s
2	3.483s	1.650s	0.964s	0.974s	0.971s
3	3.488s	1.753s	0.989s	0.959s	0.989s
5	3.439s	1.597s	1.009s	1.033s	1.082s
6	3.521s	1.651s	1.013s	0.931s	1.151s
7	3.448s	1.725s	0.976s	0.987s	0.982s
8	3.550s	1.648s	0.994s	0.979s	1.098s
9	3.496s	1.624s	1.098s	0.967s	1.137s
10	3.586s	1.615s	1.039s	1.011s	1.042s
11	3.504s	1.634s	1.092s	1.075s	1.101s
12	3.443s	1.580s	0.944s	1.003s	1.060s
13	3.840s	1.609s	1.017s	1.009s	1.071s
14	3.605s	1.579s	0.964s	1.083s	1.189s
15	3.499s	1.535s	0.981s	0.997s	1.023s
16	3.535s	1.608s	0.989s	1.287s	1.121s
17	3.446s	1.739s	1.015s	1.007s	1.129s
18	3.505s	1.544s	1.047s	1.175s	1.161s
19	3.586s	1.680s	1.013s	1.035s	1.067s
20	3.377s	1.611s	1.008s	1.096s	1.100s

8.1.2 Execution Time in Regard to Number of Input Events

Tab. 8.4: Execution time of multiple runs of TeslaServer V1 with different number of input events

Run #	500	1000	5000	10000
1	0.5815 s	1.1913 s	39.0046 s	timeout
2	0.5958 s	1.1991 s	33.5771 s	timeout
3	0.5963 s	1.2122 s	26.8252 s	timeout
5	0.580 s	1.2569 s	23.8944 s	timeout
6	0.5948 s	1.2280 s	23.8837 s	timeout
7	0.5825 s	1.2519 s	23.4110 s	timeout
8	0.6271 s	1.2636 s	23.1929 s	timeout
9	0.5705 s	1.2324 s	23.3227 s	timeout
10	0.5765 s	1.1783 s	21.9043 s	timeout
11	0.5893 s	1.1928 s	22.9516 s	timeout
12	0.6101 s	1.2759 s	22.8377 s	timeout
13	0.5684 s	1.1955 s	26.1974 s	timeout
14	0.5746 s	1.2047 s	25.7536 s	timeout
15	0.5588 s	1.2044 s	24.6289 s	timeout
16	0.5698 s	1.1958 s	22.3709 s	timeout
17	0.5911 s	1.2271 s	23.9404 s	timeout
18	0.5955 s	1.1981 s	21.4453 s	timeout
19	0.5702 s	1.1718 s	23.3482 s	timeout
20	0.5863 s	1.2571 s	21.9753 s	timeout

Tab. 8.5: Execution time of multiple runs of TeslaServer V1.5 with different number of input events

Run #	500	1000	5000	10000
1	0.4512 s	0.5901 s	2.1607 s	2.8476 s
2	0.4349 s	0.5703 s	2.0242 s	2.8715 s
3	0.4424 s	0.5511 s	2.1535 s	2.8250 s
5	0.4373 s	0.5487 s	2.1715 s	2.7025 s
6	0.4239 s	0.5564 s	2.2763 s	2.7671 s
7	0.4670 s	0.5466 s	4.2085 s	2.6719 s
8	0.4492 s	0.5401 s	2.6771 s	2.6956 s
9	0.4649 s	0.6345 s	1.9132 s	2.5634 s
10	0.440 s	0.5618 s	1.5309 s	2.6090 s
11	0.4641 s	0.5577 s	2.5762 s	2.6738 s
12	0.4306 s	0.5415 s	1.6873 s	2.7178 s
13	0.4255 s	0.5829 s	2.0342 s	2.7294 s
14	0.460 s	0.5442 s	3.0314 s	2.6789 s
15	0.4401 s	0.5443 s	3.2066 s	2.6045 s
16	0.4516 s	0.5848 s	2.5746 s	2.5837 s
17	0.4927 s	0.5406 s	2.4933 s	2.7866 s
18	0.4635 s	0.5669 s	2.6895 s	2.6719 s
19	0.4506 s	0.5576 s	2.2802 s	2.7056 s
20	0.5080 s	0.5898 s	2.4456 s	2.6697 s

Tab. 8.6: Execution time of multiple runs of TeslaServer V2 with different number of input events

Run #	500	1000	5000	10000
1	0.4415 s	0.5358 s	1.4555 s	1.9343 s
2	0.4569 s	0.5471 s	1.4538 s	1.9907 s
3	0.4549 s	0.4996 s	2.8283 s	1.9546 s
5	0.4569 s	0.4985 s	1.8618 s	1.9731 s
6	0.4603 s	0.5482 s	1.3788 s	1.9725 s
7	0.4610 s	0.5489 s	1.1817 s	2.0493 s
8	0.4585 s	0.5302 s	1.1973 s	1.9670 s
9	0.4407 s	0.5581 s	1.3172 s	2.2681 s
10	0.4414 s	0.5407 s	1.3002 s	2.0725 s
11	0.4392 s	0.5349 s	1.2226 s	1.9567 s
12	0.4652 s	0.5557 s	1.2435 s	1.9850 s
13	0.4436 s	0.5462 s	1.2375 s	2.0038 s
14	0.4458 s	0.5240 s	1.2271 s	2.0334 s
15	0.4480 s	0.5093 s	1.2287 s	1.9577 s
16	0.4508 s	0.5071 s	1.8341 s	2.0005 s
17	0.4498 s	0.5390 s	1.5943 s	2.1850 s
18	0.4820 s	0.5252 s	1.4755 s	2.4828 s
19	0.4365 s	0.5576 s	1.2696 s	2.5152 s
20	0.4379 s	0.5334 s	1.2343 s	2.3986 s

8.1.3 Ram Usage with Respect to Number of Input Events

Tab. 8.7: RAM usage of multiple runs of TeslaServer V1 with different number of input events

Run #	500	1000	5000	
1	98.50 MB	182.93 MB	1.55 GB	timeout
2	95.47 MB	226.23 MB	1.69 GB	timeout
3	84.59 MB	223.75 MB	1.62 GB	timeout
5	102.98 MB	185.86 MB	1.78 GB	timeout
6	90.24 MB	222.20 MB	1.77 GB	timeout
7	84.54 MB	181.79 MB	1.75 GB	timeout
8	89.31 MB	171.32 MB	1.60 GB	timeout
9	89.84 MB	228.12 MB	1.75 GB	timeout
10	106.07 MB	184.48 MB	1.67 GB	timeout
11	89.41 MB	194.26 MB	1.47 GB	timeout
12	97.62 MB	163.42 MB	1.80 GB	timeout
13	97.54 MB	160.26 MB	1.51 GB	timeout
14	88.77 MB	221.16 MB	1.96 GB	timeout
15	99.99 MB	223.81 MB	1.87 GB	timeout
16	97.40 MB	216.00 MB	1.87 GB	timeout
17	87.62 MB	250.93 MB	1.68 GB	timeout
18	120.70 MB	206.97 MB	1.71 GB	timeout
19	91.89 MB	246.55 MB	1.64 GB	timeout
20	88.91 MB	252.26 MB	1.87 GB	timeout

Tab. 8.8: RAM usage of multiple runs of TeslaServer V1.5 with different number of input events

Run #	500	1000	5000	1000
1	52.09 MB	53.71 MB	61.49 MB	58.37 MB
2	48.77 MB	56.98 MB	59.74 MB	60.16 MB
3	52.06 MB	52.60 MB	57.21 MB	75.88 MB
5	52.42 MB	51.23 MB	57.45 MB	61.79 MB
6	54.13 MB	52.44 MB	56.55 MB	74.53 MB
7	50.30 MB	50.41 MB	76.10 MB	62.92 MB
8	47.91 MB	52.48 MB	53.78 MB	63.47 MB
9	50.97 MB	51.99 MB	58.85 MB	60.24 MB
10	52.26 MB	51.00 MB	57.04 MB	59.36 MB
11	51.25 MB	54.08 MB	57.74 MB	62.32 MB
12	47.60 MB	52.06 MB	53.92 MB	65.00 MB
13	50.15 MB	52.54 MB	57.23 MB	64.99 MB
14	51.03 MB	52.09 MB	59.37 MB	64.33 MB
15	52.56 MB	62.94 MB	53.98 MB	65.32 MB
16	50.79 MB	50.54 MB	58.80 MB	67.25 MB
17	61.64 MB	49.85 MB	56.02 MB	58.18 MB
18	50.69 MB	51.54 MB	60.00 MB	59.93 MB
19	49.59 MB	52.21 MB	61.75 MB	82.53 MB
20	50.04 MB	51.91 MB	67.09 MB	72.43 MB

Tab. 8.9: RAM usage of multiple runs of TeslaServer V2 with different number of input events

Run #	500	1000	5000	1000
1	50.95 MB	51.71 MB	65.99 MB	71.82 MB
2	50.55 MB	59.71 MB	63.25 MB	71.65 MB
3	52.09 MB	50.17 MB	73.34 MB	79.18 MB
5	50.45 MB	51.35 MB	66.12 MB	65.73 MB
6	62.37 MB	52.33 MB	75.56 MB	72.34 MB
7	50.24 MB	49.91 MB	70.64 MB	81.99 MB
8	49.84 MB	51.51 MB	61.73 MB	70.57 MB
9	53.22 MB	52.20 MB	61.24 MB	70.05 MB
10	55.05 MB	52.32 MB	77.86 MB	74.08 MB
11	51.11 MB	50.86 MB	65.75 MB	66.15 MB
12	52.04 MB	51.57 MB	63.21 MB	68.69 MB
13	49.97 MB	50.91 MB	93.26 MB	66.23 MB
14	53.91 MB	51.40 MB	66.14 MB	61.75 MB
15	49.71 MB	68.42 MB	71.99 MB	67.96 MB
16	49.84 MB	64.05 MB	63.66 MB	68.89 MB
17	50.39 MB	51.54 MB	71.29 MB	64.43 MB
18	49.29 MB	50.30 MB	68.03 MB	69.73 MB
19	52.94 MB	61.20 MB	64.49 MB	67.56 MB
20	53.41 MB	60.22 MB	66.41 MB	83.87 MB

8.1.4 Execution Time in Regard to Number of Nodes

Tab. 8.10: Execution time of multiple runs of TesslerServer V1 in regard to different amount of nodes in a specification

Run #	8	16	32	64	128
1	0.997 s	1.295 s	2.893 s	3.780 s	7.166 s
2	0.875 s	1.280 s	2.903 s	4.139 s	7.071 s
3	1.269 s	1.277 s	2.837 s	4.055 s	7.083 s
5	0.888 s	1.262 s	3.693 s	3.811 s	6.937 s
6	0.917 s	1.330 s	4.930 s	4.150 s	7.075 s
7	0.881 s	1.419 s	3.761 s	3.826 s	7.327 s
8	0.959 s	1.398 s	3.115 s	3.835 s	7.158 s
9	0.973 s	1.294 s	2.944 s	4.093 s	7.102 s
10	0.980 s	1.257 s	3.143 s	4.208 s	6.981 s
11	0.918 s	1.285 s	2.788 s	4.092 s	7.081 s
12	0.939 s	1.329 s	2.914 s	4.033 s	7.198 s
13	0.965 s	1.225 s	2.850 s	3.815 s	7.135 s
14	0.871 s	1.320 s	3.003 s	3.817 s	6.943 s
15	0.898 s	1.336 s	2.922 s	3.905 s	6.901 s
16	0.897 s	1.357 s	3.212 s	3.968 s	7.153 s
17	1.073 s	1.401 s	3.057 s	4.081 s	7.10 s
18	1.182 s	1.342 s	2.759 s	3.826 s	7.311 s
19	1.019 s	1.288 s	2.945 s	3.991 s	7.064 s
20	0.928 s	1.314 s	3.304 s	5.343 s	7.161 s

Tab. 8.11: Execution time of multiple runs of TesslerServer V1.5 in regard to different amount of nodes in a specification

Run #	8	16	32	64	128
1	0.474 s	0.602 s	0.811 s	1.353 s	2.134 s
2	0.476 s	0.605 s	0.813 s	1.410 s	2.431 s
3	0.487 s	0.577 s	0.884 s	1.331 s	2.015 s
5	0.502 s	0.606 s	0.933 s	1.273 s	2.138 s
6	0.511 s	0.615 s	0.930 s	1.231 s	2.287 s
7	0.518 s	0.577 s	0.813 s	1.269 s	2.638 s
8	0.527 s	0.603 s	0.898 s	1.306 s	2.90 s
9	0.540 s	0.596 s	0.946 s	1.195 s	2.874 s
10	0.511 s	0.618 s	0.834 s	1.319 s	3.80 s
11	0.494 s	0.624 s	0.796 s	1.260 s	2.842 s
12	0.561 s	0.601 s	0.893 s	1.497 s	3.275 s
13	0.487 s	0.619 s	0.843 s	1.256 s	2.898 s
14	0.541 s	0.582 s	0.868 s	1.203 s	2.951 s
15	0.482 s	0.627 s	0.80 s	1.257 s	3.118 s
16	0.487 s	0.604 s	0.861 s	1.347 s	4.078 s
17	0.476 s	0.591 s	0.862 s	1.213 s	7.687 s
18	0.574 s	0.596 s	0.846 s	1.284 s	5.113 s
19	0.555 s	0.590 s	0.824 s	1.277 s	3.892 s
20	0.507 s	0.619 s	0.879 s	1.275 s	3.991 s

Tab. 8.12: Execution time of multiple runs of TesslaServer V2 in regard to different amount of nodes in a specification

Run #	8	16	32	64	128
1	0.534 s	0.548 s	0.611 s	0.784 s	1.027 s
2	0.510 s	0.533 s	0.598 s	0.741 s	1.039 s
3	0.501 s	0.546 s	0.631 s	0.718 s	1.057 s
5	0.494 s	0.552 s	0.594 s	0.731 s	1.293 s
6	0.489 s	0.563 s	0.598 s	0.747 s	1.231 s
7	0.521 s	0.539 s	1.041 s	0.737 s	1.236 s
8	0.496 s	0.548 s	0.676 s	0.721 s	1.221 s
9	0.507 s	0.541 s	0.618 s	0.745 s	1.262 s
10	0.483 s	0.559 s	0.579 s	0.748 s	1.229 s
11	0.507 s	0.554 s	0.579 s	1.271 s	1.342 s
12	0.487 s	0.539 s	0.836 s	0.739 s	1.060 s
13	0.507 s	0.526 s	0.574 s	0.720 s	1.174 s
14	0.545 s	0.530 s	0.586 s	0.782 s	1.274 s
15	0.518 s	0.563 s	0.597 s	0.888 s	1.218 s
16	0.561 s	0.528 s	0.606 s	1.214 s	1.278 s
17	0.605 s	0.528 s	0.602 s	1.006 s	1.247 s
18	0.550 s	0.557 s	0.651 s	0.738 s	1.043 s
19	0.550 s	0.523 s	0.645 s	0.750 s	1.027 s
20	0.541 s	0.534 s	1.051 s	0.735 s	0.987 s

8.2 Ringbuffer Code

```
#include <stdio.h>
#include <stdatomic.h>
3 #include <pthread.h>
#include <stdlib.h>
#include <sys/time.h>
6
int BUF_SIZE = 5;

9 char* ring_buffer;
char* _Atomic read_ptr;
char* write_ptr;
12
void init_ring_buffer() {
    ring_buffer = malloc(BUF_SIZE);
15    read_ptr = ring_buffer;
    write_ptr = ring_buffer;

18 }

char* buffer_next(char* ptr) {
21    return (char*) ring_buffer + (((ptr - ring_buffer) + 1)
        % BUF_SIZE);
}

24 void process(char* data){
    struct timespec ts;
    ts.tv_sec = 0;
27    ts.tv_nsec = 1000000;
    nanosleep(&ts, NULL);
}
30
void producer_main() {
    char* next_write_ptr;
33    while (!feof(stdin)){
        next_write_ptr = buffer_next(write_ptr);
        if (next_write_ptr != read_ptr) {
36            *write_ptr = getc(stdin);
            write_ptr = next_write_ptr;
```

```

39     struct timespec ts;
        ts.tv_sec = 0;
        ts.tv_nsec = 100000000ULL * rand() / RAND_MAX ;
42     nanosleep(&ts, NULL);
        }
    }
45 }

void* consumer_main(void* thread_id) {
48     int tid = (int) thread_id;
        char* current_rptr;
        char* current_wptr;
51     char data = 0;
        char* next_read_ptr;
        char* local_read_ptr;
54
        while (1) {
            local_read_ptr = read_ptr;
57
            if (local_read_ptr != write_ptr) {
                data = *local_read_ptr;
60                next_read_ptr = buffer_next(local_read_ptr);
                if (atomic_compare_exchange_weak(&read_ptr, &
                    local_read_ptr, next_read_ptr)) {
                    process(&data);
63                }
                struct timespec ts;
                ts.tv_sec = 0;
66                ts.tv_nsec = 100000000ULL * rand() / RAND_MAX ;
                nanosleep(&ts, NULL);
            }
69     }
}

72 void create_consumers(int num_consumers) {
    pthread_t threads[num_consumers];
    int thread_ids[num_consumers];
75     int rc;
    for(int t = 0; t < num_consumers; t++){
        thread_ids[t] = t;
78     rc = pthread_create(&threads[t], NULL, consumer_main,
        (void* ) t);

```

```

        if (rc){
            pthread_exit(NULL);
81    }
    }
}
84
int main() {
    init_ring_buffer();
87
    create_consumers(2);

90    producer_main();

    return 0;
93 }

```

Listing 8.1: Code of the ringbuffer example programm

8.3 Instrumentation Benchmark Data

Tab. 8.13: Execution time of multiple runs of an uninstrumented example C program compiled with different optimization levels

Run #	O3	O2	O1	O0
1	0.1161 s	0.1191 s	0.2948 s	0.6530 s
2	0.1157 s	0.1156 s	0.2738 s	0.6632 s
3	0.1167 s	0.1148 s	0.2762 s	0.6660 s
5	0.1141 s	0.1238 s	0.2678 s	0.6449 s
6	0.1164 s	0.1141 s	0.2763 s	0.6779 s
7	0.1149 s	0.1187 s	0.2759 s	0.6569 s
8	0.1200 s	0.1117 s	0.2793 s	0.6452 s
9	0.1100 s	0.1159 s	0.2894 s	0.6577 s
10	0.1103 s	0.1151 s	0.2816 s	0.6565 s
11	0.1138 s	0.1219 s	0.2807 s	0.6649 s
12	0.1163 s	0.1176 s	0.2723 s	0.6588 s
13	0.1170 s	0.1168 s	0.2781 s	0.6618 s
14	0.1127 s	0.1200 s	0.2799 s	0.6571 s
15	0.1124 s	0.1124 s	0.2763 s	0.6544 s
16	0.1231 s	0.1198 s	0.2752 s	0.6740 s
17	0.1203 s	0.1191 s	0.2750 s	0.6417 s
18	0.1163 s	0.1229 s	0.2662 s	0.6602 s
19	0.1186 s	0.1143 s	0.2685 s	0.6608 s
20	0.1182 s	0.1160 s	0.2935 s	0.6540 s
21	0.1166 s	0.1128 s	0.2743 s	0.6557 s
22	0.1196 s	0.1137 s	0.2758 s	0.6767 s
23	0.1156 s	0.1164 s	0.2886 s	0.6637 s
25	0.1203 s	0.1166 s	0.2829 s	0.6816 s
26	0.1173 s	0.1203 s	0.2849 s	0.6397 s
27	0.1160 s	0.1158 s	0.2814 s	0.6493 s
28	0.1161 s	0.1233 s	0.2841 s	0.6558 s
29	0.1086 s	0.1245 s	0.2797 s	0.6555 s
30	0.1138 s	0.1158 s	0.2746 s	0.6605 s
31	0.1159 s	0.1194 s	0.2767 s	0.6677 s
32	0.1231 s	0.1197 s	0.2818 s	0.6489 s
33	0.1154 s	0.1125 s	0.2842 s	0.6527 s
34	0.1147 s	0.1216 s	0.2762 s	0.6488 s
35	0.1205 s	0.1136 s	0.2732 s	0.6657 s
36	0.1184 s	0.1135 s	0.2743 s	0.6611 s
37	0.1149 s	0.1160 s	0.2773 s	0.6637 s
38	0.1183 s	0.1157 s	0.2722 s	0.6638 s
39	0.1138 s	0.1105 s	0.2787 s	0.6738 s
40	0.1113 s	0.1144 s	0.2707 s	0.6705 s

Tab. 8.14: Execution time of multiple runs of an instrumented example C program compiled with different optimization levels

Run #	O3	O2	O1	O0
1	3.3911 s	3.2436 s	3.4216 s	3.8694 s
2	3.5288 s	3.6410 s	3.3192 s	3.6850 s
3	3.3494 s	3.4974 s	3.3641 s	3.6085 s
5	3.5640 s	3.4832 s	3.5960 s	3.7213 s
6	3.3368 s	3.4711 s	3.4086 s	3.6712 s
7	3.5097 s	3.5463 s	3.4823 s	3.6867 s
8	3.3817 s	3.3919 s	3.4265 s	3.8107 s
9	3.6170 s	3.4889 s	3.5831 s	3.6996 s
10	3.3764 s	3.3402 s	3.4814 s	3.7145 s
11	3.4867 s	3.7752 s	3.5054 s	3.7900 s
12	3.3393 s	3.3320 s	3.6760 s	3.7082 s
13	3.2769 s	3.3553 s	3.7271 s	3.6436 s
14	3.5930 s	3.3282 s	3.4532 s	3.6124 s
15	3.5464 s	3.5488 s	3.5590 s	3.7060 s
16	3.6068 s	3.5292 s	3.4408 s	3.6440 s
17	3.3833 s	3.3312 s	3.6440 s	3.6952 s
18	4.6447 s	3.2691 s	3.6565 s	3.7749 s
19	3.9256 s	3.3252 s	3.5309 s	3.6850 s
20	3.3639 s	3.3914 s	3.6782 s	3.6187 s
21	3.6016 s	3.3473 s	4.5602 s	3.6672 s
22	3.2825 s	3.3359 s	4.2268 s	3.5982 s
23	3.1969 s	3.3892 s	4.3686 s	3.6702 s
25	3.3095 s	3.3870 s	3.5713 s	3.7751 s
26	3.4130 s	3.4193 s	3.6797 s	3.6604 s
27	3.5319 s	3.3532 s	3.5412 s	3.6823 s
28	3.3935 s	3.3253 s	3.5808 s	3.6570 s
29	3.3294 s	3.3051 s	3.4674 s	3.6164 s
30	3.4536 s	3.3509 s	3.5215 s	3.6644 s
31	3.4952 s	3.4388 s	3.8695 s	3.7006 s
32	3.4600 s	3.3684 s	4.1289 s	3.6317 s
33	3.5761 s	3.4953 s	4.5513 s	3.6137 s
34	3.5816 s	3.4220 s	4.4331 s	3.7160 s
35	3.6319 s	4.3748 s	3.6797 s	3.6431 s
36	3.4037 s	3.7254 s	3.5498 s	3.6337 s
37	3.5317 s	3.3875 s	3.9773 s	4.0382 s
38	3.6149 s	3.2542 s	3.6003 s	3.6492 s
39	3.3683 s	3.1844 s	3.6274 s	3.6503 s
40	3.3448 s	3.3463 s	3.6728 s	3.6539 s

Tab. 8.15: Execution time of multiple runs of an instrumented program with different percentages of instrumented function calls

Run #	10%	5%	2%	1%	0.5%	0.1%	0.01%	0.001%
1	32.2744 s	16.5746 s	6.8182 s	3.4653 s	1.9216 s	0.5291 s	0.3014 s	0.1262 s
2	32.0180 s	17.2402 s	6.7273 s	4.7988 s	1.8523 s	0.5143 s	0.2784 s	0.1243 s
3	32.5687 s	16.5926 s	6.9581 s	4.0173 s	1.7571 s	0.4944 s	0.2864 s	0.1245 s
5	31.7277 s	16.4390 s	6.5665 s	3.9327 s	1.7359 s	0.4897 s	0.2750 s	0.1209 s
6	33.2172 s	16.2865 s	6.6697 s	4.3147 s	1.7712 s	0.4748 s	0.2754 s	0.1266 s
7	32.0802 s	16.4187 s	6.6856 s	3.9003 s	1.7668 s	0.4553 s	0.2718 s	0.1314 s
8	36.0954 s	16.7155 s	6.6875 s	3.8426 s	1.8202 s	0.4432 s	0.2712 s	0.1301 s
9	36.9455 s	16.7840 s	6.5936 s	4.2499 s	1.7469 s	0.4662 s	0.2723 s	0.1200 s
10	33.0339 s	20.1691 s	7.4481 s	3.4560 s	1.8394 s	0.4528 s	0.3151 s	0.1175 s
11	35.2251 s	19.8842 s	7.2807 s	3.8017 s	1.7728 s	0.4460 s	0.3352 s	0.1189 s
12	31.8794 s	20.2660 s	7.0673 s	5.2281 s	1.6975 s	0.4340 s	0.2739 s	0.1209 s
13	32.9260 s	19.2307 s	6.7634 s	4.3974 s	1.7026 s	0.4441 s	0.2708 s	0.1195 s
14	31.9317 s	17.8926 s	6.7118 s	3.7419 s	1.7544 s	0.4361 s	0.2752 s	0.1216 s
15	32.6378 s	17.5177 s	7.7485 s	3.3878 s	1.8489 s	0.5027 s	0.2767 s	0.1240 s
16	32.1711 s	16.3863 s	6.7146 s	3.3716 s	2.0388 s	0.4380 s	0.2893 s	0.1266 s
17	32.4985 s	16.4408 s	6.6772 s	4.1358 s	1.8717 s	0.4349 s	0.2669 s	0.1210 s
18	32.1387 s	16.4152 s	7.1151 s	3.2681 s	1.8253 s	0.4491 s	0.2719 s	0.1217 s
19	32.0004 s	16.1371 s	6.7711 s	3.6237 s	1.7483 s	0.4274 s	0.2678 s	0.1215 s
20	32.5547 s	16.4341 s	7.2219 s	3.7627 s	1.7674 s	0.4555 s	0.2658 s	0.1222 s
21	32.1219 s	16.1507 s	6.7387 s	3.4525 s	1.7188 s	0.4372 s	0.2741 s	0.1280 s
22	32.7006 s	16.2854 s	6.5879 s	3.4259 s	1.7273 s	0.4447 s	0.2731 s	0.1207 s
23	32.0705 s	15.9990 s	6.6927 s	3.6898 s	1.7741 s	0.4446 s	0.2728 s	0.1271 s

Tab. 8.15: Execution time of multiple runs of an instrumented program with different percentages of instrumented function calls

Run #	10%	5%	2%	1%	0.5%	0.1%	0.01%	0.001%
25	32.5558 s	16.2882 s	6.7336 s	3.2830 s	1.8821 s	0.4399 s	0.2744 s	0.1245 s
26	35.2615 s	17.6253 s	6.5753 s	3.3661 s	1.8344 s	0.4657 s	0.2745 s	0.1279 s
27	46.2991 s	16.6608 s	6.6603 s	3.3115 s	1.7554 s	0.4431 s	0.2678 s	0.1266 s
28	34.3145 s	16.7270 s	6.7223 s	3.5270 s	1.7145 s	0.4346 s	0.2728 s	0.1251 s
29	33.3674 s	15.9586 s	6.7363 s	3.3917 s	1.7318 s	0.4426 s	0.2735 s	0.1257 s
30	35.3249 s	15.8946 s	7.1817 s	3.2812 s	1.7364 s	0.4255 s	0.2741 s	0.1274 s
31	37.5802 s	17.0438 s	8.0833 s	3.9514 s	1.7507 s	0.4614 s	0.2701 s	0.1250 s
32	35.6964 s	17.9818 s	6.6254 s	3.4312 s	1.8008 s	0.4516 s	0.2693 s	0.1183 s
33	32.3928 s	16.2117 s	6.8224 s	3.4366 s	1.7850 s	0.4451 s	0.2717 s	0.1215 s
34	34.6837 s	16.1150 s	8.0794 s	3.4316 s	1.7732 s	0.4375 s	0.2807 s	0.1304 s
35	32.9116 s	16.6887 s	7.7993 s	3.6139 s	1.7912 s	0.4468 s	0.2682 s	0.1248 s
36	32.2931 s	17.4771 s	6.6512 s	3.4691 s	1.9558 s	0.4335 s	0.2717 s	0.1269 s
37	32.5576 s	16.2437 s	6.6823 s	3.3764 s	1.8431 s	0.4469 s	0.2719 s	0.1273 s
38	32.4156 s	16.1199 s	6.6837 s	3.3485 s	1.8451 s	0.4372 s	0.2792 s	0.1234 s
39	32.7322 s	16.1902 s	6.6598 s	3.5290 s	1.8274 s	0.4583 s	0.2744 s	0.1252 s
40	32.7395 s	16.4886 s	6.7993 s	3.3594 s	1.8212 s	0.4294 s	0.2709 s	0.1233 s

8.4 Example TESSLA Specifications

```
define values: Signal<Int> :=  
    variable_values("buffer.c:write_ptr")  
3 define writes: Events<Int> := changeOf(values)  
    define processed: Events<Unit> :=  
        function_calls("buffer.c:process")  
6  
    define bufLevel: Signal<Int> :=  
        sub(eventCount(writes), eventCount(processed))  
9  
    define error: Signal<Boolean> :=  
        signalNot(and(  
12     leq(literal(0), bufLevel), leq(bufLevel, literal(5))  
        )  
    )
```

Listing 8.2: TESSLA specification that checks if the size of a buffer is always between zero and five. It counts the number of events happened on two input streams. The first stream denotes changes of the variable *write_ptr*, the second calls of the function *process*. An error is encountered when more data is processed than is present or if more than five items are written onto the buffer.

```

1 define values: Signal<Int> :=
    variable_values("buffer.c:write_ptr")
    define new_data: Events<Int> := changeOf(values)
4 define data_processed: Events<Unit> :=
    function_calls("buffer.c:process")

7 define delayed_new_data: Events<Unit> :=
    delayEventByTime(new_data, 1000000)

10 define consumed_in_past: Signal<Boolean> :=
    inPast(1000000, data_processed)

13 define error: Events<Boolean> :=
    eventNot(sample(consumed_in_past, delayed_new_data))

```

Listing 8.3: TESSLA specification describing a performance constraint. The specification describes that after each change of the *write_ptr* variable there has to be a call to the *process* function after at most 1 millisecond. Since TesslaServer only implements the past time functions of TESSLA the specification is a bit more complex than it would be when using funvtions that can depend on future values.

Bibliography

- [AD94] Rajeev Alur and D.L. Dill. „A theory of timed automata“. In: *Theoretical computer science* 126.2 (1994), pp. 183–235 (cit. on p. 22).
- [AF16] Duncan Paul Attard and Adrian Francalanza. „A Monitoring Tool for a Branching-Time Logic“. In: *Runtime Verification - 16th International Conference, Madrid*. 163406. 2016, pp. 473–481 (cit. on pp. 8, 12).
- [Ber+16] Dean Michael Berris, Alistair Veitch, Nevin Heintze, Eric Anderson, and Ning Wang. *XRay: A Function Call Tracing System*. Tech. rep. 2016 (cit. on p. 14).
- [Bla+06] Stephen M Blackburn, Robin Garner, Chris Hoffmann, et al. „The DaCapo Benchmarks : Java Benchmarking Development and Analysis“. In: *Proceedings of the 21st annual ACM SIGPLAN conference on Object-oriented programming systems, languages, and applications* (2006), pp. 169–190 (cit. on p. 13).
- [BLS06] Andreas Bauer, Martin Leucker, and Jonathan Streit. „SALT - Structured assertion language for temporal logic“. In: *Lecture Notes in Computer Science (including subseries Lecture Notes in Artificial Intelligence and Lecture Notes in Bioinformatics)* 4260 LNCS (2006), pp. 757–775 (cit. on pp. 11, 12).
- [CE82] Edmund M Clarke and E Allen Emerson. „Design and synthesis of synchronization skeletons using branching time temporal logic“. In: *Logics of Programs* 131 (1982), pp. 52–71 (cit. on p. 11).
- [Che07] Feng Chen. „MOP : An Efficient and Generic Runtime Verification Framework “. In: *ACM SIGPLAN Notices* 42.10 (2007), pp. 569–588 (cit. on p. 13).
- [CR09] Feng Chen and Grigore Rou. „Parametric trace slicing and monitoring“. In: *Lecture Notes in Computer Science (including subseries Lecture Notes in Artificial Intelligence and Lecture Notes in Bioinformatics)* 5505 LNCS (2009), pp. 246–261 (cit. on p. 86).
- [CSL04] Bryan M Cantrill, Michael W Shapiro, and Adam H Leventhal. „DTrace: Dynamic Instrumentation of Production Systems“. In: *Development July 2002* (2004), pp. 15–28 (cit. on p. 15).
- [DAC99] M.B. Dwyer, G.S. Avrunin, and J.C. Corbett. „Patterns in property specifications for finite-state verification“. In: *Proceedings of the 1999 International Conference on Software Engineering (IEEE Cat. No.99CB37002)* (1999), pp. 411–420 (cit. on pp. 12, 13).
- [DAn+05] Ben D’Angelo, Sriram Sankaranarayanan, César Sánchez, et al. „LOLA: Runtime monitoring of synchronous systems“. In: *Proceedings of the International Workshop on Temporal Representation and Reasoning* (2005), pp. 166–175 (cit. on pp. 1, 2, 8, 9).
- [DTH16] Normann Decker, Daniel Thoma, and Jannis Harder. „TESSLA A Temporal Stream-based Specification Language“. 2016 (cit. on pp. 1, 10, 11).

- [EC00] Hans-Dieter Ehrich and Carlos Caleiro. „Specifying communication in distributed information systems“. In: *Acta Informatica* 616 (2000), pp. 1–24 (cit. on p. 7).
- [Fay+16] Peter Faymonville, Bernd Finkbeiner, Sebastian Schirmer, and Hazem Torfah. „A Stream-Based Specification Language for Network Monitoring“. In: *Runtime Verification: 16th International Conference, RV 2016, Madrid, Spain, September 23–30, 2016, Proceedings*. Ed. by Yliès Falcone and César Sánchez. Cham: Springer International Publishing, 2016, pp. 152–168 (cit. on p. 9).
- [Fid88] Colin J Fidge. „Timestamps in Message-Passing Systems That Preserve the Partial Ordering“. In: *Proc of the 11th Australian Computer Science Conference ACSC88* 10.1 (1988), pp. 56–66. arXiv: 10614036 (cit. on p. 7).
- [Hal16] Sylvain Hallé. „When RV Meets CEP“. In: *Runtime Verification - 16th International Conference, Madrid*. Ed. by Yliès Falcone and César Sánchez. Vol. 10012. Lecture Notes in Computer Science April. Cham: Springer International Publishing, 2016, pp. 68–91 (cit. on pp. 10, 65, 87).
- [Hav08] Klaus Havelund. „Runtime Verification of C Programs“. In: (2008) (cit. on pp. 1, 5, 7, 12).
- [HBS73] Carl Hewitt, Peter Bishop, and Richard Steiger. „A Universal Modular ACTOR Formalism for Artificial Intelligence“. In: *Ijcai* (1973), pp. 235–245 (cit. on p. 60).
- [Kic+01] Gregor Kiczales, Erik Hilsdale, Jim Hugunin, et al. „An Overview of AspectJ“. In: *ECOOP 2001 Object-Oriented Programming: 15th European Conference Budapest, Hungary, June 1822, 2001 Proceedings* 2072.4 (2001), pp. 327–354 (cit. on p. 12).
- [Koy90] Ron Koymans. „Specifying real-time properties with metric temporal logic“. In: *Real-Time Systems* 2.4 (1990), pp. 255–299 (cit. on p. 11).
- [LA04] Chris Lattner and Vikram Adve. „LLVM: A compilation framework for lifelong program analysis & transformation“. In: *International Symposium on Code Generation and Optimization, CGO c* (2004), pp. 75–86 (cit. on p. 16).
- [Lam78] Leslie Lamport. „Time, Clocks, and the Ordering of Events in a Distributed System“. In: *Communications of the ACM* 21.7 (1978), pp. 558–565. arXiv: 10614036 (cit. on p. 7).
- [LS07] Martin Leucker and César Sánchez. „Regular Linear Temporal Logic“. In: *Theoretical Aspects of Computing ICTAC 2007*. Vol. 4711 LNCS. Berlin, Heidelberg: Springer Berlin Heidelberg, 2007, pp. 291–305 (cit. on p. 11).
- [MB15] Menna Mostafa and Borzoo Bonakdarpour. „Decentralized Runtime Verification of LTL Specifications in Distributed Systems“. In: *2015 IEEE International Parallel and Distributed Processing Symposium* (2015), pp. 494–503 (cit. on pp. 1, 7).
- [MN04] Oded Maler and Dejan Nickovic. „Monitoring Temporal Properties of Continuous Signals“. In: *Formal Techniques, Modelling and Analysis of Timed and Fault-Tolerant Systems*. 2004, pp. 152–166 (cit. on p. 11).

- [MZA12] L Marek, Y Zheng, and D Ansaloni. „DiSL: a domain-specific language for byte-code instrumentation“. In: *Proceedings of the 11th annual international conference on Aspect-oriented Software Development*. 2012, pp. 239–250 (cit. on pp. 12, 13).
- [Nav+13] Samaneh Navabpour, Yogi Joshi, Wallace Wu, et al. „RiTHM: A Tool for Enabling Time-triggered Runtime Verification for C Programs“. In: *Proceedings of the 2013 9th Joint Meeting on Foundations of Software Engineering* (2013), pp. 603–606 (cit. on p. 12).
- [Nec+02] George C. Necula, Scott McPeak, Shree P. Rahul, and Westley Weimer. „CIL: Intermediate language and tools for analysis and transformation of C programs“. In: *Lecture Notes in Computer Science (including subseries Lecture Notes in Artificial Intelligence and Lecture Notes in Bioinformatics)* 2304 (2002), pp. 213–228 (cit. on pp. 7, 14).
- [Pik+10] Lee Pike, Alwyn Goodloe, Robin Morisset, and Sebastian Niller. „Copilot: A hard real-time runtime monitor“. In: *Lecture Notes in Computer Science (including subseries Lecture Notes in Artificial Intelligence and Lecture Notes in Bioinformatics)* 6418 LNCS.Rv (2010), pp. 345–359 (cit. on pp. 1, 6).
- [Pnu77] Amir Pnueli. *The Temporal Logic of Programs*. 1977 (cit. on p. 11).
- [RHF16] Giles Reger, Sylvain Hallé, and Yliès Falcone. „Third International Competition on Runtime Verification“. In: 2016, pp. 21–37 (cit. on p. 14).
- [RS97] Jean-François Raskin and Pierre-yves Schobbens. „State clock logic: A decidable real-time logic“. In: 1997, pp. 33–47 (cit. on p. 11).
- [RSS16] Carl Martin Rosenberg, Martin Steffen, and Volker Stolz. „Leveraging DTrace for Runtime Verification“. In: *Runtime Verification - 16th International Conference, Madrid*. 2016, pp. 318–332 (cit. on p. 16).
- [Sen+04] Koushik Sen, Abhay Vardhan, Gul Agha, and Grigore Rosu. „Efficient Decentralized Monitoring of Safety in Distributed Systems“. In: *26th International Conference on Software Engineering, ICSE 2004 V.October* (2004), pp. 418–427 (cit. on p. 7).
- [SH08] Margaret H. Smith and Klaus Havelund. „Requirements capture with RCAT“. In: *Proceedings of the 16th IEEE International Requirements Engineering Conference, RE'08* (2008), pp. 183–192 (cit. on p. 7).
- [Wu+16] Rongxin Wu, Xiao Xiao, Shing-Chi Cheung, Hongyu Zhang, and Charles Zhang. „Casper: An Efficient Approach to Call Trace Collection“. In: (2016), pp. 678–690 (cit. on p. 13).
- [Zhe+15] Xi Zheng, Christine Julien, Rodion Podorozhny, and Franck Cassez. „BraceAssertion: Runtime verification of cyber-physical systems“. In: *Proceedings - 2015 IEEE 12th International Conference on Mobile Ad Hoc and Sensor Systems, MASS 2015* (2015), pp. 298–306 (cit. on p. 1).
- [Zho+14] Jingwen Zhou, Zhenbang Chen, Ji Wang, Zibin Zheng, and Michael R. Lyu. „Trace Bench: An Open Data Set for Trace-Oriented Monitoring“. In: *2014 IEEE 6th International Conference on Cloud Computing Technology and Science* (2014), pp. 519–526 (cit. on p. 13).

List of Figures

2.1	Visualization of TESSLA stream model, taken from [DTH16]	11
3.1	Influence of the order of independent transitions on the global state of an evaluation engine	42
4.1	Visualization of a simple evaluation engine with a greedy schedule. . .	48
4.2	Visualization of three runs of an evaluation engine	52
5.1	Control flow of a node	64
6.1	Performance of the runtime with different number of used cores	72
6.2	Average execution time of the different implementations when fed with different amounts of input events.	74
6.3	RAM usage of the different versions of the runtime	75
6.4	Average execution time of the runtime over specifications with different amount of nodes.	76
6.5	Performance Comparision of an example C program with and without instrumentation.	78
6.6	Execution time of a program with different amount of calls to an instrumented function	79

List of Tables

3.1	List of complete functions	32
3.2	List of output complete functions	33
3.3	List of input complete functions	33
3.4	List of incomplete functions	35
3.5	List of timing functions	36
4.1	Example how the behaviour of a run is constructed	46
8.1	Execution time of multiple runs of TesslaServer V1 with 10 000 input events with different number of used processor cores.	89
8.2	Execution time of multiple runs of TesslaServer V1.5 with 10 000 input events with different number of used processor cores	90
8.3	Execution time of multiple runs of TesslaServer V2 with 10 000 input events with different number of used processor cores	90
8.4	Execution time of multiple runs of TesslaServer V1 with different number of input events	91
8.5	Execution time of multiple runs of TesslaServer V1.5 with different number of input events	91
8.6	Execution time of multiple runs of TesslaServer V2 with different number of input events	92
8.7	RAM usage of multiple runs of TesslaServer V1 with different number of input events	93
8.8	RAM usage of multiple runs of TesslaServer V1.5 with different number of input events	93
8.9	RAM usage of multiple runs of TesslaServer V2 with different number of input events	94
8.10	Execution time of multiple runs of TesslaServer V1 in regard to different amount of nodes in a specification	95
8.11	Execution time of multiple runs of TesslaServer V1.5 in regard to different amount of nodes in a specification	95
8.12	Execution time of multiple runs of TesslaServer V2 in regard to different amount of nodes in a specification	96
8.13	Execution time of multiple runs of an uninstrumented example C program compiled with different optimization levels	101

8.14	Execution time of multiple runs of an instrumented example C program compiled with different optimization levels	102
8.15	Execution time of multiple runs of an instrumented program with dif- ferent percentages of instrumented function calls	103

Glossary

LOLA A specification language and algorithms for the online and offline monitoring of synchronous systems including circuits and embedded systems. 1, 8, 9, 12, 61

RCAT Requirements CApture Tool. 7

RMOR Requirement Monitoring and Recovery. 1, 7, 12

TESSLA A temporal, stream based specification Language. 1–3, 5–14, 19, 27–30, 36, 39, 41, 46, 51, 59, 61, 63, 65, 68, 71, 72, 75, 78, 80–82, 84–88, 105, 106, 111

mHML Monitorable Hennessy-Milner Logic. 8

API Application Programming Interface. 10, 61, 63, 67, 69, 84

BEAM Bogdan/Björn's Erlang Abstract Machine. 59, 60

CEP Complex Event Processign. 10, 12, 85

CIL C Intermediate Language. 7, 14

CRV Competition on Runtime Verification. 14

CTL Computational Tree Logic. 11

DAG Directed Acyclic Graph. 29, 46–48, 50, 55, 61, 62

eSQL Event Stream Query Language. 10

FIFO First In First Out. 29

FPGA Field Programmable Gate Array. 2

GCC GNU Compiler Collection. 15

HDFS Hadoop Distributed Filesystem. 13

IR Intermediate Representation. 16, 17, 69, 84

ISP Institute for Software Engineering and Programming Languages. 1

JSON JavaScript Object Notation. 62, 87

JVM Java Virtual Machine. 5, 85

LLVM Low Level Virtual Machine. 3, 15–17, 68, 69

LTL Linear Temporal Logic. 7, 11, 16, 61

MTL Metric Temporal Logic. 11

OTP Open Telecom Platform. 59, 63

RAM random access memory. 66, 73–75, 111

RISC Reduced Instruction Set Computing. 16

RiTHM Runtime Time-triggered Heterogeneous Monitoring. 12

RLTL Regular Linear Temporal Logic. 11

RV Runtime Verification. 1–3, 5–8, 10–13, 16, 85, 86, 88

SALT Structured Assertion Language for Temporal Logic. 11, 12

STL Signal Temporal Logic. 11

TL Temporal Logic. 1, 9

TLTL Timed Linear Temporal Logic. 11

VM Virtual Machine. 59, 60

List of Theorems

1	Definition (Transducer)	19
2	Definition (Deterministic Transducer)	19
3	Definition (Synchronous Transducer)	20
4	Definition (Asynchronous Transducer)	20
5	Definition (Causal and Clairvoyant Transducers)	20
6	Definition (Asynchronous equivalence of Transducers)	20
1	Lemma (Asynchronous equivalence is an equivalence Relation) . . .	21
7	Definition (Timed Sequence)	22
8	Definition (Monotonicity of Timed Sequences)	23
9	Definition (Timed Transducer)	23
10	Definition (Boundedness of Timed Transducers)	23
11	Definition (Observational Equivalence)	24
2	Lemma (Observational Equivalence is an Equivalence Relationship for Bounded Transducers)	24
12	Definition (Application of a Transition on a State)	37
13	Definition (Closeness of Runs)	38
14	Definition (Enabledness of a Node)	39
15	Definition (Valid Run)	39
16	Definition (Independence of Nodes)	41
17	Definition (Independence of Transitions)	41
3	Lemma (Exchange of Independent Transitions)	41
4	Lemma (Duration of Enabledness)	42
5	Lemma (Finiteness of Enabledness)	43
18	Definition (Fair Schedules)	44
19	Definition (Behaviour of a Run)	45
20	Definition (Equivalence of Runs)	46
21	Definition (Equivalence of Evaluation Engines)	46
22	Definition (Greedy schedule)	47
23	Definition (Valid Evaluation Engines)	48
6	Lemma (Greedy Schedules are Fair)	49
24	Definition (Fair Evaluation Engines)	49
1	Theorem (Equivalence of Different Greedy Evaluation Engines) . . .	51
2	Theorem (Equivalence of Fair and Greedy Evaluation Engines) . . .	56

