

An Asynchronous Evaluation Engine for Stream Based Specifications

Alexander Schramm

November 20, 2016
Version: My First Draft

University of Luebeck



UNIVERSITÄT ZU LÜBECK

Institute For Software Engineering and Programming Languages

isp

Master Thesis

An Asynchronous Evaluation Engine for Stream Based Specifications

Alexander Schramm

1. Reviewer

Prof. Dr. Martin Leucker

Institute For Software Engineering and Programming Languages
University of Luebeck

Supervisor

Cesar Sanchez

November 20, 2016

Alexander Schramm

An Asynchronous Evaluation Engine for Stream Based Specifications

Master Thesis, November 20, 2016

Reviewers: Prof. Dr. Martin Leucker

Supervisors: Cesar Sanchez

University of Luebeck

Institute For Software Engineering and Programming Languages

Ratzeburger Allee 160

23562 Luebeck

Abstract

Abstract (Deutsch)

Acknowledgement

Contents

1	Introduction	1
1.1	Motivation and Problem Statement	1
1.2	Results	2
1.3	Thesis Structure	2
2	Related Work	5
2.1	TESSLA	5
2.2	LOLA	6
2.3	Distributed Verification Techniques	7
2.4	Copilot	7
2.5	RMOR	8
3	System	11
3.1	TESSLA Runtime	11
3.1.1	Erlang and Elixir	11
3.1.2	Implementation	11
3.2	Trace Data	11
3.2.1	Debie	12
3.2.2	TraceBench	12
3.2.3	Aspect oriented programming	12
3.2.4	CIL	12
3.2.5	Google XRay	12
3.2.6	GCC instrument functions	12
3.2.7	Sampling	12
3.2.8	LLVM/clang AST matchers	12
4	Preliminaries	13
4.1	Time	13
4.2	Transducers	13
4.3	Timed Transducers	16
4.4	Labeled Timed Transducers	21
4.5	Events	21
4.6	Streams	21
4.7	Functions	22

4.8	Nodes	23
4.9	TESSLA Evaluation Engine	23
4.10	TESSLA Functions	24
4.10.1	Complete Functions	25
4.10.2	Output Complete Functions	25
4.10.3	Input Complete Functions	25
4.10.4	Incomplete Functions	26
4.10.5	Timing Functions	29
4.11	State and History	29
4.12	Transitions	30
4.13	Run	31
5	Behaviours of Evaluation Engines	39
5.1	Schedules Without Timing Functions	40
5.1.1	Greedy Evaluation Engines	41
5.1.2	Fair Evaluation Engines	43
5.2	Equivalence of Different Schedules Without Timing Functions	44
5.2.1	Equivalence of Greedy Systems	44
5.2.2	Equivalence of Greedy and Fair Evaluation Engines	49
5.3	Behaviour with Timing functions	50
5.4	Equalities with Timing functions	50
5.5	Parallel computation	50
6	Evaluation	51
6.1	Runtime Benchmarks	51
6.1.1	Number of Processors	51
6.1.2	Number of Events	52
6.1.3	Number of Nodes	53
6.2	Practical Examples	53
6.2.1	Ringbuffer Examples	53
6.3	Instrumentation Benchmarks	53
6.3.1	Performance Comparison with non Instrumented Code and Compiler Optimizations	55
6.3.2	Performance Impact in Regard to Instrumentation Percentage	55
6.3.3	Conclusion for the Instrumentation	55
7	Conclusion	59
7.1	Further Work	59
7.1.1	Error prevention	59
8	Appendix	61
8.1	Ringbuffer Code	61
8.2	Instrumentation Benchmark Data	63

Bibliography	69
Glossary	75

Introduction

1.1 Motivation and Problem Statement

Software Verification is an important tool to harden critical systems against faults and exploits. Due to the raising importance of computer based systems, verification has become a big field of research in computer science.

While pure verification approaches try to proof the correct behaviour of a system under all possible executions, Runtime Verification (RV) limits itself to single, finite runs of a system, trying to proof it conforms to a given specification under specific conditions, like input sequences or scheduling. These specifications can be given in various ways, e.g. as a Temporal Logic (TL) formula or in specification languages that are specifically developed for RV. Examples for this are RMOR [4], LOLA [2] and others [9, 7, 5], which we will look at more closely in Section 2.

The language TeSSLa aims to make it easy to specify behaviour of streams. To gain this it introduces a number of language features and syntax sugar to expressively describe the conditions a stream should fulfill. The evaluation of TeSSLa specifications is done in two steps: first the specification is compiled by a compiler written at the Institute for Software Engineering and Programming Languages (ISP) of the University of Lübeck. The output is a canonical representation of the operations on the streams in the specification. In the second step the compiled specification is connected with a system that produces some kind of traces, which are treated as the input streams of the specification.

The second step can be done in different ways: online or offline, interweaving the monitors into the monitored program (e.g. [4]) or having a standalone system. These different approaches lead to different manipulations of the original program that should be monitored. When the monitors are interweaved into the program, they can produce new errors or even suppress others. When the monitors are run in a different process or even on different hardware, the overhead and influence to the system can be much smaller, but there will be a bigger delay between the occurrence of events in the program and their evaluation in the monitor. Furthermore interweaved monitors can optionally react towards errors by changing the program execution, therefore eliminating cascading errors, while external executions of monitors can't directly modify the program but can still produce warnings to prevent such errors. While online monitoring can be used to actively react to

error conditions, either automatically or by notification of a third party, offline monitoring can be thought of as an extension to software testing ([2]).

At the beginning of this thesis there was one implementation of a runtime for TESSLA specifications that is based on FPGAs that have to be manually reconfigured for each new specification. While this is a very performant approach for actual monitoring it isn't feasible for testing and prototyping. Therefore it is wanted to implement a runtime for TeSSLa specifications that can be run independent of specific hardware.

Furthermore most RV approaches are specific to one programming language or environment and combine ways of generating the data, which is used for monitoring, and the monitoring itself. TESSLA specifications themselves are independent of any implementation details of the monitored system, working only on streams of data, which can be gathered in any way. This can be used to implement a runtime that is also independent of the monitored system and how traces of it are collected.

During the thesis it is proven, that the actual approach of this runtime, a functional, actor based, asynchronous system, will generate the same observations on input traces as a synchronous evaluation of the specification. While TESSLA specifications can work on all kinds of streams, especially on traces on all levels of a program, e.g. on instruction counters or on spawning processes, in this thesis we will mainly focus on the level of function calls and variable reads/writes. Other applications of the system can easily extend it to use traces of drastically different fields, e.g. health data, temperatures, battery levels, web services and more.

To test the software based runtime, different specifications will be tested on multiple traces, some of which are generated by actually running a program, which was instrumented by hand to generate traces, others which are generated or modified by hand to deliberately introduce bugs which should be detected by the system.

1.2 Results

1.3 Thesis Structure

As the whole evaluation engine is built on top of different technical and theoretical ideas, it is structured to show the reasoning behind the decisions that were made during the development. Furthermore it will prove equality of different kinds of systems in multiple steps that build on one another. In the following a quick overview of the different parts of the thesis is given.

Chapter 2

In this chapter the theoretical foundation for the system is explained. Furthermore multiple approaches solving similar problems are shown and it is highlighted which concepts of them were used in the new system and which were disregarded and why.

Chapter 3

Moving towards the implementation of the new evaluation engine, in this chapter practical concepts and systems that are used to implement the runtime are shown. It is explained how they are used and which alternatives exist.

Chapter 4

Building on the theoretical and practical findings of the previous chapters new definitions are presented, which are needed to reason about the implemented system.

Chapter 5

The work from the previous chapter is put to work to reason about the semantics of the implemented runtime and to show its correctness.

Chapter 7

To show the value of the implemented system it is thoroughly tested with real world examples and traces. The results of this testing is used to evaluate the implementation.

Related Work

As Runtime Monitoring and Verification is a widely researched field, multiple approaches to attain its goals were developed.

As stated in [4] most approaches are geared towards software written in Java, while many critical systems are written in C and there are countless other systems that could benefit from monitoring and verification written in all kinds of programming languages. With TESSLA as a specification language over streams, which has no assumptions on the environment of the system that produces the streams, as the base for our monitoring approach, we recognized the possibility to abstract the monitoring platform from the monitored program. This means that the developed runtime for TESSLA is not restricted to monitor programs written in a specific language but can monitor anything that can produce streams of data.

To show that the runtime is valuable in the context of existing approaches, we will show ways to generate traces from systems that were used to evaluate other monitoring techniques. Afterwards we will compare the expressiveness of TESSLA and the runtime with other approaches, based on the generated traces, to show what kinds of specifications can be monitored with TESSLA and where the language or the runtime can be extended.

The following chapter will highlight the systems against which TESSLA and the runtime is evaluated, furthermore it will also give insights into other work that TESSLA and this thesis is based on.

Just a collection of thoughts for now, needs to be polished a lot

2.1 TESSLA

The implemented runtime and the theoretic work of this thesis is built upon the TESSLA project from [3]. For that project a syntax and a formal semantic of a specification language was defined.

Specifications in TESSLA are based on streams of data. Streams are the representation of data over time, e.g. a variable value in a program or the temperature of a processor. To model streams TESSLA defines a timing model. That model is




Fig. 2.1: Visualization of TESSLA stream model, taken from [3]

based on timestamps that are isomorphic to real numbers \mathbb{R} . Figure 2.1 shows how streams behave over time.

The syntax of TESSLA is pretty small, but can be used to define complex functions and specifications:

```

spec ::= define name[: stype] := texpr
      out texprspec spec
texpr := expr[: type]
expr  := name | literal | name(texpr(, texpr)*)
stype := btype | stype
stype := Signal<btype> | Events<btype>

```

2.2 LOLA

The concepts of LOLA [2] are very similar to the ones of TESSLA. Both approaches built upon streams of events. The biggest difference in the modelation is, that while streams in LOLA are based on a discrete model of time TESSLA uses a continuous timing model.

The specification language of LOLA is very small (expressions are built upon three operators) but the expressiveness surpasses TLs and many other formalisms [2]. Expressions in LOLA are built by manipulating existing streams to form new ones. Therefore streams depend on other streams, so they can be arranged in a weighted dependency graph, where the weight describes the amount of steps a generated Stream is delayed compared to the parent.

Based on this graph a notion of efficiently monitorable properties is given and an algorithm to monitor them is presented.

TESSLA takes concepts of LOLA and applies them to a continuous model of time and introduces a language and a rich set of functions that can be applied to streams. The dependency graph is a core concept of TESSLA and is used to check if specifications are valid (e.g. cycle free) and is also the core concept to evaluate specifications over traces in this thesis.

2.3 Distributed Verification Techniques

While most implementations of RV systems don't consider or use modern ways of parallelism and distribution and focus on programs running locally, in [5] a way to monitor distributed programs is presented. To do this distributed monitors, which have to communicate with one another, are specified and implemented.

As stated earlier, the TESSLA runtime doesn't care about the environment of the monitored program, so it doesn't distinguish between traces from distributed and non distributed programs. But the runtime itself is highly concurrent and can be distributed easily to many processors or even different computers. Therefore many of the definitions for distributed monitors can be used to reason about the behaviour of the runtime.

2.4 Copilot

The realtime runtime monitor system Copilot was introduced in [7]. Copilot is designed to overcome the shortcomings of existing RV tools in regards to hard-realtime software written in C.

To do so they first define characteristics a monitoring approach has to fulfill to be considered valuable for this domain. The four principles are:

Functionality Monitors cannot change the functionality of the observed program unless a failure is observed.

Schedulability Monitors cannot alter the schedule of the observed program.

Certifiability Monitors must minimize the difficulty in re-validating the observed program; in particular, we make it our goal to avoid modifying the observed programs source code.

SWaP overhead Monitors must minimize the additional overhead required including size, weight, and power (SWaP).

The monitors follow a sampling based approach, where at specified steps the values of global variables are observed and the monitors are evaluated on that values. While sampling based approaches are widely disregarded in RV, because they can lead to both false positives and false negatives, they argue:

In a hard real-time context, sampling is a suitable strategy. Under the assumption that the monitor and the observed program share a global clock and a static periodic schedule, while false positives are possible, false negatives are not. [7]

A special detail of Copilot is that monitors aren't inlined into the program but can be scheduled as independent processes. The implementation of the TESSLA runtime in this thesis follows a similar approach: It is a totally independent program, and therefore also has some of the gains in regard to the specified four characteristics. Because the runtime works with all kinds of traces, it is insignificant how they are produced: It can work with traces based on sampling, working in a similar fashion as Copilot, or by actually instrumenting code to generate traces, which alters the semantics of the program.

2.5 RMOR

RMOR [4] is another approach on monitoring C programs. It does so by transforming C code into an *armored* version, which includes monitors to check conformance to a specification.

Specifications are given as a textual representation of state machines, which is strongly influenced by RCAT [8]. The specifications are then interweaved into the program using CIL [6]. Specifications work on the level of function calls and state properties like *write may never be called before open was called*. Because software developers are often working at the same abstraction level (in contrast to e.g. assembler or machine instructions), they can define specifications without having to learn new concepts. In the TESSLA runtime support for traces at the same abstraction level (function calls, variable reads and writes) is present and used in most of the tests in Section ??.

Because RMOR specifications are interweaved into the program, their observations can not only be reported but also used to recover the program or even to prevent errors by calling specified functions when a critical condition is encountered. The

TESSLA runtime doesn't support this out of the box, as it's primary purpose is testing and offline monitoring, but in Section 7.1.1 we will look at possible extensions to support this.

System

Besides the theoretical basics presented in Section 2 the TESSLA runtime of this thesis is built upon a number of technologies. To better understand decisions made during the implementation this chapter will give an overview of them and show why they were chosen.

As already mentioned, the implemented runtime itself is independent of the way traces are generated. Therefore we will not only look at building blocks for the runtime itself but also examine related projects which can be used to obtain traces, which then can be monitored by the runtime. Because the format of the traces can differ heavily, depending on how and why they were collected, they are not only used to test the runtime but also to determine how it can consume them.

3.1 TESSLA Runtime

The runtime to evaluate specifications is implemented in the programming language Elixir, which itself is built on top of Erlang. To understand why this platform was chosen we will look at the Erlang ecosystem in the next section.

3.1.1 Erlang and Elixir

3.1.2 Implementation

3.2 Trace Data

Problem: Many traces don't carry timestamp (see DeCapo, CRV 15)

<http://ltnng.org> <http://diamon.org/ctf/> <https://github.com/efficios/barectf>

BEAM,
Ac-
tors/Thread,
multi-
plat-
form
(nerves
project)

Timing
model:
reason
why
ctf
events
have to
carry

times-
tamps
in con-

3.2.1 [Debie](#)

3.2.2 [TraceBench](#)

3.2.3 [Aspect oriented programming](#)

3.2.4 [CIL](#)

3.2.5 [Google XRay](#)

3.2.6 [GCC instrument functions](#)

3.2.7 [Sampling](#)

3.2.8 [LLVM/clang AST matchers](#)

Preliminaries

In this chapter we will define concepts that are used in Chapter 5 to reason about the implemented runtime.

While the TESSLA specification itself defines a set of semantics, for this thesis we will slightly alter some of it and add some new definitions based on them. This is necessary to reason about the specifics how the runtime is built (Note that TESSLA doesn't define an operational semantic, therefore we will define our own) and how it behaves.

4.1 Time

TESSLA has a model of continuous time, where timestamps $\pi \in \mathbb{T}$ are used to represent a certain point in time and \mathbb{T} has to be isomorphic to \mathbb{R} .

4.2 Transducers

Fundamentally TESSLA is a special kind of a transducer. Therefore in this section we will define a model of transducers which can be used to reason about the evaluation of a TESSLA specification.

A transducer is a system, which consumes an input and produces an output. Let Φ, Γ be two alphabets and ϵ the empty word.

Definition 1: Transducer.

A transducer t is a relation $t \subseteq \Phi^ \times \Gamma^*$, Φ is called the input alphabet, Γ the output alphabet.*

TESSLA specifications are deterministic for any input, meaning they should produce the same output for the same input.

Definition 2: Deterministic Transducer.

A deterministic transducer relates each input to at most one output.

Example 1: Deterministic and Nondeterministic Transducers.

$t_d = \{(a, 1), (b, 2), (ab, 12), (ba, 21)\}$ is a deterministic transducer; $t_{nd} = \{(a, 1), (a, 2)\}$ is nondeterministic, because it relates input a to both outputs 1 and 2.

Transducers can furthermore be categorized as synchronous, asynchronous, causal and clairvoyant transducers: synchronosity is a property over the behaviour of a transducer when it's consuming input per element. If it is synchronous, it will produce an output element for each input element.

Definition 3: Synchronous Transducer.

Let $\vec{i} \in \Phi^*$, $i \in \Phi$, $\vec{o} \in \Gamma^*$, $o \in \Gamma$. A transducer t is called synchronous, when it satisfies, that: if $(\vec{i} \circ i, \vec{o} \circ o) \in t$ then $(\vec{i}, \vec{o}) \in t$

An asynchronous transducer can produce zero, one or many outputs for each input it consumes.

Definition 4: Asynchronous Transducer.

Let $\vec{i} \in \Phi^*$, $i \in \Phi$, $\vec{o} \in \Gamma^*$. A transducer t is called asynchronous when it satisfies the formula: if $(\vec{i} \circ i, \vec{o}) \in t$ then $\exists \vec{o}', \vec{o}'' \in \Gamma^*$ so that $\vec{o} = \vec{o}' \circ \vec{o}''$ and $(\vec{i}, \vec{o}') \in t$

Example 2: Synchronous and Asynchronous Transducers.

$t_s = \{(a, 1), (b, 2), (ab, 12), (ba, 21)\}$ is a synchronous transducer, $t_{as} = \{(a, \epsilon), (ab, 12)\}$ is asynchronous.

A causal transducer is one, where the output depends only on consumed inputs and not on future inputs:

Definition 5: Causal and Clairvoyant Transducers.

A transducer t is called causal, when it satisfies, that: if $(\vec{i}, \vec{o}) \in t$ then $\forall \vec{i}' \in \Phi^*$ with $(\vec{i} \circ \vec{i}', \vec{o}) \in t$ it holds, that $\vec{o} \sqsubseteq \vec{o}'$

A transducer that isn't causal is called clairvoyant.

Example 3: Causal and Clairvoyant Transducers.

$t_{cl} = \{(a, 1), (b, 2), (ab, 12), (ba, 21)\}$ is a causal transducer, because each output only depends on the inputs seen upto that point, $t_{cl} = \{(a, 1), (ab, 22), (aa, 11)\}$ is clairvoyant, because the output when the letter a is seen depends on the next input.

When talking about transducers, it is interesting to know if two transducers are equivalent. There are multiple possible definitions for equivalence of transducers, we will look at two, which are interesting for this thesis. In the following σ_i is used to get the element at position i and $\sigma_{[i,j]}$ to get the infix of σ which starts at position i and ends at position j (With 0 as the index of the first element).

Definition 6: Asynchronous equivalence of Transducers.

Let t_1, t_2 be two asynchronous transducers from Φ^* to Γ^* . They are called asynchronous equivalent, written $t_1 \equiv_a t_2$, if they satisfy:

$\forall \sigma \in \Phi^*$:

- $\forall (\sigma_{[0,k]}, \vec{o}) \in t_1: \exists k' \geq k$ with $(\sigma_{[0,k']}, \vec{o}) \in t_2$ and $\vec{o} \sqsubseteq \vec{o}'$
- and $\forall (\sigma_{[0,k]}, \vec{o}) \in t_2: \exists k' \geq k$ with $(\sigma_{[0,k']}, \vec{o}) \in t_1$ and $\vec{o} \sqsubseteq \vec{o}'$

Lemma 1: Asynchronous equivalence is an equivalence Relation.

Asynchronous equivalence is symmetric, reflexive and transitive.

Proof.

Symmetry is trivial, since the second part of the definition is requiring it.

Reflexivity is also trivial, for $(\sigma_{[0,k]}, \vec{o})$ select $k' = k$.

For transitivity:

Let $t_1 \equiv_a t_2, t_2 \equiv_a t_3$.

First case:

Since $t_1 \equiv_a t_2 : \forall (\sigma_{[0,k_1]}, \vec{o}_1) \in t_1 :$

$\exists k_2$ such, that $(\sigma_{[0,k_2]}, \vec{o}_2) \in t_2$ with $\vec{o}_1 \sqsubseteq \vec{o}_2$

and since $t_2 \equiv_a t_3$

$\exists k_3$ such, that $(\sigma_{[0,k_3]}, \vec{o}_3) \in t_3$ with $\vec{o}_2 \sqsubseteq \vec{o}_3$

With $\vec{o}_1 \sqsubseteq \vec{o}_2 \sqsubseteq \vec{o}_3$ it follows, that $t_1 \equiv_a t_3$

The second case works the same, just change t_1 and t_3 .

□

Example 4: Asynchronous equivalence of Transducers.

Let $\Phi = \{a\}, \Gamma = \{1\}$ and

$$\begin{array}{llll} t_1 = \{ & (a, \epsilon), & (aa, \epsilon), & (aaa, 111) & \} \\ t_2 = \{ & (a, 1), & (aa, 1), & (aaa, 111) & \} \\ t_3 = \{ & (a, \epsilon), & (aa, 1), & (aaa, 11) & \} \end{array}$$

All three transducers are asynchronous and causal. Let's see which ones are asynchronous equivalent:

$$t_1 \stackrel{?}{\equiv}_a t_2$$

(a, ϵ)	$\in t_1, k = 1$	$\rightarrow k' = 1, (a, 1) \in t_2,$	$\epsilon \sqsubseteq 1$
(aa, ϵ)	$\in t_1, k = 2$	$\rightarrow k' = 2, (aa, 1) \in t_2,$	$\epsilon \sqsubseteq 1$
$(aaa, 111)$	$\in t_1, k = 3$	$\rightarrow k' = 3, (aaa, 111) \in t_2,$	$111 \sqsubseteq 111$
$(a, 1)$	$\in t_2, k = 1$	$\rightarrow k' = 3, (aaa, 111) \in t_1,$	$1 \sqsubseteq 111$
$(aa, 1)$	$\in t_2, k = 2$	$\rightarrow k' = 3, (aaa, 111) \in t_1,$	$1 \sqsubseteq 111$
$(aaa, 111)$	$\in t_2, k = 3$	$\rightarrow k' = 3, (aaa, 111) \in t_1,$	$111 \sqsubseteq 111$

$$\Rightarrow t_1 \equiv_a t_2$$

$$t_1 \stackrel{?}{\equiv}_a t_3$$

$$(aaa, 111) \in t_1, k = 3 \rightarrow \nexists k'$$

$$\Rightarrow t_1 \not\equiv_a t_3$$

Because of Lemma 1 $\Rightarrow t_2 \not\equiv_a t_3$.

4.3 Timed Transducers

For the second kind of equivalence we need to introduce *timed sequences*, originally introduced as *timed words* in [1], and *timed transducers*. Note that timed sequences don't have to be monotonically increasing like in the original definition. Quite on the contrary the unorderedness of outputs is an important key principle to much of the later work as you will see.

Let \mathbb{T} be a timing model that is isomorphic to \mathbb{R} . For the examples we will use \mathbb{R} for \mathbb{T} .

Definition 7: Timed Sequence.

A sequence is called *timed*, if every element of it is associated with a timestamp: $\sigma \in (\Gamma \times \mathbb{T})^*$. For brevity a timed sequence can be written with the timestamps as the index of the elements: $\sigma = e_0 e_{0.5} e_1$.

The function

$$timed : (\Gamma \times \mathbb{T})^* \rightarrow (\Gamma \times \mathbb{T})^*$$

reorders a timed sequence σ by its timestamps, such that:

$$\forall i, j \in \mathbb{N} : \text{ if } i < j \text{ then } \pi_i < \pi_j \text{ with } (o_i, \pi_i) = \sigma_i \text{ and } (o_j, \pi_j) = \sigma_j$$

The function

$$upto : \mathbb{T} \times (\Gamma \times \mathbb{T})^* \rightarrow (\Gamma \times \mathbb{T})^*$$

removes all elements from a timed sequence, that have a timestamp bigger than the first argument.

The function

$$maxTime : (\Gamma \times \mathbb{T})^* \rightarrow \mathbb{T}$$

returns the biggest Timestamp in a timed sequence.

Example 5: Functions on Timed Sequences.

Let $\sigma = a_1 a_{0.5} a_{1.5} a_0$.

Then is

$$timed(\sigma) = a_0 a_{0.5} a_1 a_{1.5}$$

$$upto(1.3, \sigma) = a_1 a_{0.5} a_0$$

$$maxTime(\sigma) = 1.5$$

Definition 8: Monotonicity of Timed Sequences.

A timed sequence σ with alphabet Φ is called monotonic, if $timed(\sigma) = \sigma$

Definition 9: Timed Transducer.

A timed transducer t with input alphabet Φ and output alphabet Γ works on monotonic, timed sequences as inputs and has timed sequences as outputs:

$$t \subset (\Phi \times \mathbb{T})^* \times (\Gamma \times \mathbb{T})^*$$

Example 6: Timed Transducers.

Let $\Phi = \{a\}, \Gamma = \{b\}$.

$t_{tsc} = \{(a_0, b_0), (a_0 a_1, b_0 b_1)\}$ is a timed, causal and synchronous transducer.

$t_{tac} = \{(a_0, \epsilon), (a_0 a_1, b_0 b_1)\}$ is a timed, causal and asynchronous transducer.

For later theoretic work we have to restrict timed transducers.

Definition 10: Boundedness of Timed Transducers.

A timed transducer t with input alphabet Φ and output alphabet Γ is called bounded, if it satisfies:

$$\begin{aligned}
& \forall \sigma \in (\Phi \times \mathbb{T})^* : \\
& \quad \text{if } (\sigma_{[0,k]}, \vec{o}) \in t \\
& \quad \text{then } \exists k' > k \text{ with} \\
& \quad \quad (\sigma_{[0,k']}, \vec{o} \circ \vec{o}') \in t \\
& \quad \text{and } \forall k'' > k' \text{ with } (\sigma_{[0,k'']}, \vec{o} \circ \vec{o}' \circ \vec{o}'') \in t \text{ it holds, that} \\
& \quad \quad \text{upto}(\maxTime(\vec{o}), \text{timed}(\vec{o} \circ \vec{o}')) = \text{upto}(\maxTime(\vec{o}), \text{timed}(\vec{o} \circ \vec{o}' \circ \vec{o}''))
\end{aligned}$$

Based on the definitions we can define an equivalence relationship on bounded timed transducers.

Definition 11: Observational Equivalence.

Let t_1, t_2 be two bounded timed transducers with input alphabet Φ and output alphabet Γ . They are called observational equivalent, written $t_1 \equiv_o t_2$, if they satisfy:

$$\begin{aligned}
& \forall \sigma \in (\Phi \times \mathbb{T})^* : \\
& \quad \forall (\sigma_{[0,k]}, \vec{o}) \in t_1 : \exists k', k'' \geq k \text{ such that} \\
& \quad \quad (\sigma_{[0,k']}, \vec{o} \circ \vec{o}') \in t_1 \\
& \quad \quad \text{and } (\sigma_{[0,k'']}, \vec{o}_2) \in t_2 \\
& \quad \quad \text{and } \text{timed}(\text{upto}(\maxTime(\vec{o}), \vec{o} \circ \vec{o}')) = \text{timed}(\text{upto}(\maxTime(\vec{o}), \vec{o}_2))
\end{aligned}$$

and the same for switched t_1, t_2 .

What does observational equivalence between two transducers intuitively mean? It means that two transducers eventually produce the same output values for the same timed inputs, maybe in a different order, but with the same timestamps, which is very important. Since the values are associated with timestamps the outputs can be reordered by them and therefore be exactly equal.

Lemma 2: Observational Equivalence is an Equivalence Relationship for Bounded Transducers.

\equiv_o is symmetric, reflexive and transitive for bounded timed transducers.

Proof.

Let t_1, t_2, t_3 be bounded timed transducers. Symmetry follows directly from the definition.

Reflexivity: For $(\sigma_{[0,k]}, \vec{o})$ select $k' = k''$ as the k , for which the transducer is bounded for that input.

Transitivity:

- Let $t_1 \equiv_o t_2, t_2 \equiv_o t_3$.

- First case:

Since $t_1 \equiv_o t_2 : \forall (\sigma_{[0,k_1]}, \vec{o}_1) \in t_1 :$

$\exists k'_1, k_2 > k_1$ with $(\sigma_{[0,k'_1]}, \vec{o}_1 \circ \vec{o}_1') \in t_1$ and $(\sigma_{[0,k_2]}, \vec{o}_2) \in t_2$

with $\text{timed}(\text{upto}(\text{maxTime}(\vec{o}_1), \vec{o}_1 \circ \vec{o}_1'))$

$= \text{timed}(\text{upto}(\text{maxTime}(\vec{o}_1), \vec{o}_2))$

(★)

and since $t_2 \equiv_o t_3 : \exists k'_2, k_3 > k_2$ with $(\sigma_{[0,k'_2]}, \vec{o}_2 \circ \vec{o}_2') \in t_2$

and $(\sigma_{[0,k_3]}, \vec{o}_3) \in t_3$

with $\text{timed}(\text{upto}(\text{maxTime}(\vec{o}_2), \vec{o}_2 \circ \vec{o}_2'))$

$= \text{timed}(\text{upto}(\text{maxTime}(\vec{o}_2), \vec{o}_3))$

(★★)

$\text{maxTime}(\vec{o}_1)$ has to be smaller than $\text{maxTime}(\vec{o}_2)$

else (★) couldn't hold, therefore, combined with boundedness and (★★) :

$\text{timed}(\text{upto}(\text{maxTime}(\vec{o}_1), \vec{o}_2))$

$= \text{timed}(\text{upto}(\text{maxTime}(\vec{o}_1), \vec{o}_3))$

which concludes $\text{timed}(\text{upto}(\text{maxTime}(\vec{o}_1), \vec{o}_1 \circ \vec{o}_1'))$

$= \text{timed}(\text{upto}(\text{maxTime}(\vec{o}_1), \vec{o}_3))$

- The second case works the same, just switch t_1 and t_3 .

□

Example 7: Observational Equivalence.

Let

$$\begin{aligned} t_1 &= \{ (a_0, \epsilon), (a_0a_1, b_1), (a_0a_1a_2, b_1b_2b_0) \} \\ t_2 &= \{ (a_0, \epsilon), (a_0a_1, \epsilon), (a_0a_1a_2, b_2b_1b_0) \} \\ t_3 &= \{ (a_0, b_0), (a_0a_1, b_0), (a_0a_1a_2, b_2b_1) \} \end{aligned}$$

All three are causal, asynchronous timed transducers.

Let's see which ones are observational equivalent:

$$t_1 \stackrel{?}{\equiv}_o t_2$$

$$(a_0, \epsilon) \in t_1, k = 1, \text{maxTime}(\epsilon) = 0$$

$$\rightarrow k' = 1, (a_0, \epsilon) \in t_1$$

$$\rightarrow k'' = 1, (a_0, \epsilon) \in t_2$$

$$(a_0a_1, b_1) \in t_1, k = 2, \text{maxTime}(b_1) = 1$$

$$\rightarrow k' = 3, (a_0 a_1 a_2, b_1 b_2 b_0) \in t_1$$

$$\rightarrow k'' = 3, (a_0 a_1 a_2, b_2 b_1 b_0) \in t_2$$

$$\text{timed}(\text{upto}(1, b_1 b_2 b_0)) = b_0 b_1 = \text{timed}(\text{upto}(1, b_2 b_1 b_0))$$

$$(a_0 a_1 a_2, b_1 b_2 b_0) \in t_1, k = 3, \text{maxTime}(b_1 b_2 b_0) = 2$$

$$\rightarrow k' = 3, (a_0 a_1 a_2, b_1 b_2 b_0) \in t_1$$

$$\rightarrow k'' = 3, (a_0 a_1 a_2, b_2 b_1 b_0) \in t_2$$

$$\text{timed}(\text{upto}(2, b_1 b_2 b_0)) = b_0 b_1 b_2 = \text{timed}(\text{upto}(2, b_2 b_1 b_0))$$

$$(a_0, \epsilon) \in t_2, k = 1, \text{maxTime}(\epsilon) = 0$$

$$\rightarrow k' = 1, (a_0, \epsilon) \in t_2$$

$$\rightarrow k'' = 1, (a_0, \epsilon) \in t_1$$

$$(a_0 a_1, \epsilon) \in t_2, k = 2, \text{maxTime}(\epsilon) = 0$$

$$\rightarrow k' = 2, (a_0 a_1, \epsilon) \in t_2$$

$$\rightarrow k'' = 2, (a_0 a_1, b_1) \in t_1$$

$$\text{timed}(\text{upto}(0, \epsilon)) = \epsilon = \text{timed}(\text{upto}(0, b_1))$$

$$(a_0 a_1 a_2, b_2 b_1 b_0) \in t_2, k = 3, \text{maxTime}(b_2 b_1 b_0) = 2$$

$$\rightarrow k' = 3, (a_0 a_1 a_2, b_2 b_1 b_0) \in t_2$$

$$\rightarrow k'' = 3, (a_0 a_1 a_2, b_1 b_2 b_0) \in t_1$$

$$\text{timed}(\text{upto}(2, b_2 b_1 b_0)) = b_0 b_1 b_2 = \text{timed}(\text{upto}(2, b_1 b_2 b_0))$$

$$\Rightarrow t_1 \equiv_a t_2$$

$$t_1 \stackrel{?}{\equiv}_a t_3$$

$$(a_0 a_1 a_2, b_1 b_2 b_0) \in t_1, k = 3, \text{maxTime}(b_1 b_2 b_0) = 2$$

$$\rightarrow k' = 3, (a_0 a_1 a_2, b_1 b_2 b_0) \in t_1$$

$$\rightarrow \vec{A}(\vec{i}, \vec{o}) \in t_3 \text{ with } \exists n \in \mathbb{N} : \vec{o}_n = b_0$$

$$\rightarrow \vec{A}(\vec{i}, \vec{o}) \in t_3 \text{ with } \text{timed}(\text{upto}(2, b_1 b_2 b_0)) = b_0 b_1 b_2 = \text{timed}(\text{upto}(2, \vec{o}))$$

$$\Rightarrow t_1 \not\equiv_a t_3$$

If t_3 weren't bounded (and therefore not finite) there would be no way to know, if it was equivalent to t_1 , because it could always produce a missing event at a later time.

Because of Lemma 2 $\Rightarrow t_2 \not\equiv_a t_3$.

4.4 Labeled Timed Transducers

Maybe necessary, maybe not

4.5 Events

Events are the atomic unit of information that all computations are based on. There are three types of events: external, output and internal events.

The set of all events is denoted as E . Each event carries a value, which can be *nothing* or a value of a type (types are formally defined in the TESSLA specification, but aren't important for this thesis), a timestamp and the stream it's perceived on (e.g. a function call of a specific function or the name of an output stream).

The value of an event can be queried with the function v , its timestamp with $time$ and its stream with $stream$.

$E_e \subset E$ is the set of all external events, their stream corresponds to a specific trace. $E_o \subset E$ is the set of all output events, their stream is specified by an output name of the TESSLA specification. $E_n \subset E$ is the set of all internal events. Internal events are mostly an implementation detail, which denote steps of computation inside the runtime. The stream of internal events is implicitly given by the node that produces the stream of the event. Note that E_e, E_o, E_n are pairwise disjoint and $E_e \cup E_o \cup E_n = E$.

4.6 Streams

Streams are a collection of events with specific characteristics. While events are the atomic unit of information, streams represent the sequence of related events over time.

There are two kind of streams: signals, which carry values at all times, and eventstreams, which only hold values at specific times. Eventstreams can be described by a sequence of events. Signals can be described by a sequence of changes, where a change denotes that the value of a signal changed at a specific timestamp. The only difference between a signal and an eventstreams is that signals always have a value while an eventstream may return \perp when queried for its value at a specific time, which denotes that no event happened at that time. Based on the similarity of signals and eventstreams in the following we will mainly reason about eventstreams, but most things can also be applied to signals.

Formally a stream σ can be represented as the product of a sequence of events $\langle e_1, \dots, e_n \rangle$ where $time(e_i) < time(e_{i+1})$, $\forall i < n \in \mathbb{N}$. The set of all streams Σ is defined as all possible finite sequences of events $\Sigma = \{\sigma | \sigma \in E^*\}$. An external stream σ_e is a stream consisting only of external events, the set of all external streams is $\Sigma_e = \{\sigma_e | \sigma_e \in E_e^*\} \times \mathbb{T}$. Output and internal streams are defined analogous.

To get the event of a stream σ at a timestamp π it can be queried like a function: $\sigma(\pi) = e$ with $time(e) = \pi$. When working with signals, the function will return the latest event that happened at or before t while an eventstream may return \perp . The progress of a stream, which is the timestamp of the last event that happened on them, can be obtained with $progress(\sigma) = \pi \in \mathbb{T}$. Internal and output streams can be queried for the node that produced them with $node(\sigma) = n \in N$.

4.7 Functions

A TESSLA specification consists of functions over streams. Functions generate new streams by applying an operation on existing streams. TESSLA itself defines a syntax to write a specification, a set of types and a standard library of functions, but an implementation is free to choose the functions it supports.

An example function is $add(S_D, S_D) \rightarrow S_D$: It takes two signals, which have to hold values of some numerical type, and produces a signal which holds values of the same type. The produced stream can either be assigned to a named identifier (think: a variable) or consumed by another function (function composition).

Functions can be divided into three categories: pure, unpure and timing. Pure functions, also called stateless, are evaluated only on the values their inputs have at the timestamp they are evaluated, therefore they don't have to remember a state and will only return events. Unpure, or stateful, functions are evaluated over the values if its inputs at that timestamp and a state and will return not only new events but also an updated state. E.g. a function *eventCount* has to *remember* how many events already happened on its input stream and increment that counter on every new event. Timing functions are evaluated not only on the value of events but also on their timestamp and can also manipulate it: While non timing functions will consume events at a specific timestamp and emit events with that timestamp, timing functions can emit events with a changed timestamp. In this thesis we will only look at past time functions, meaning functions can only delay timestamps, therefore can't depend on future values.

Timing functions complicate the reasoning about schedules and causality and therefore aren't included in Section 5.1. In Section 5.4 the conclusions of earlier sections will be extended to include timing functions.

4.8 Nodes

Nodes are the atomic unit of computation for the evaluation of a TESSLA specification. A node implements a single function, e.g.: there is an *AddNode* which takes two input signals and produces a new signal. Therefore a node is the concrete implementation of a function in a runtime for TESSLA specifications. The set of all nodes is called N . The function of a node $n \in N$ is written as f_n .

Each node has a set of inputs, which are either external or internal streams, and one output, which is either an internal or an output stream. Nodes which have at least one external stream as an input are called *sources*. Nodes have a state, described in Section 4.11, which contains First In First Out (FIFO) queues, provided by the Erlang platform, which buffer events from its inputs for later computation.

Every new event added to a queue has to have a bigger timestamp than the previous event added to the queue. This means a queue has a kind of progress timestamp, which denotes the timestamp of the latest event added to it and which is strictly increasing over time. Queues support the standard operators for lists like *hd*, *tl*, *++* to respectively get the head, the tail or to append to the end.

4.9 TESSLA Evaluation Engine

Because functions in TESSLA specifications depend on other functions, and these dependencies have to be cycle free, the specification can be represented as a Directed Acyclic Graph (DAG), where the functions are vertices and the relationship between functions are edges. This is exactly how the TESSLA compiler outputs a specification. One can now use the DAG of a TESSLA specification to synthesize a system to evaluate it over inputs: The vertices of the DAG become nodes representing the functions and the edges are the input and output streams between the nodes. We will call this synthesized system an *evaluation engine*.

When fed with inputs (or *traces*) the engine will produce outputs. The relationship between inputs and outputs that is produced can be seen as a timed transducer. The input to an evaluation engine has to have strictly increasing timestamps. This is needed to have a known progress which can be distributed through the system. If inputs weren't ordered by their timestamp for example the absence of input events on a specific stream couldn't be detected because events could be present at a later position of the input trace. Especially for offline monitoring this obviously is no problem because the traces can simply be reordered into a strictly increasing sequence, except when multiple input events are at the same timestamp. This can be solved in two ways: either increase the timing precision when generating the traces

or manipulate the timestamps in the traces by adding a minimal offset to them if they are equal to another timestamp.

To evaluate a specification over traces, the evaluation engine has process the events that were traced. To do so the nodes have to run their computations until no more events are present (or the specification found an error in the trace). This leads to the question in which order nodes should be scheduled to perform their computation. We will use the term *step* to denote that one node was scheduled and performed its computation. While some schedules are simply not rational (think of unfairness and causality), there are many different schedules that are feasible. It has to be proven that a chosen schedule produces the correct conclusions for a specification, else the evaluation engine is not valid.

An evaluation engine is run inside an environment. The environment has knowledge over the state of the nodes, most important which nodes are enabled. Based on that information the environment is responsible of feeding the trace data to the engine: only when no sources are enabled the next trace is added to the queue of an input node. This ensures that the consumption of input signals is strictly ordered by timestamp, which is important as we will see in Section 5.

4.10 TESSLA Functions

TESSLA puts no restrictions on the semantics of functions other than that they have to work on streams or constants and produce streams, but allows to restrict them for evaluation approaches. We are taking advantage of that to categorize functions based on how or if at all they can be encoded in our evaluation approach. The categorization is based upon the relationship between consumption and production of input and output events that a node representing the function in an evaluation engine produces. For this we will use the terms node and functions somewhat interchangeable in the following subsections.

Nodes can only be scheduled when they are enabled, meaning they have events on their inputs buffered. When a node is scheduled, it will compute the minimal timestamp of all buffered events. The function implemented by the node is then evaluated at that timestamp, which is called the *evaluation timestamp*. This is important to understand the completeness criteria: It means that when the function is evaluated there is at least one event with that timestamp on one input. The inverse of that statement shows why this is important: a function is never evaluated at a timestamp where no event is present, therefore the system can't arbitrarily produce new timestamps.

Nodes having signals as inputs always have to remember the last occurred change of them in their state. When such a node is scheduled, the function will work on the remembered value if no new change is present at the evaluation timestamp for the signal. If a new change is present at the evaluation timestamp, it will be used and the state of the node will be updated to remember the new change.

It is important to note that all functions always have to consume an event from at least one input queue, else an evaluation engine can enter a livelock, where new events are produced forever out of nowhere. Also all functions will only produce a finite amount of events at every evaluation. Especially only timed functions can produce more than one event at an evaluation timestamp.

4.10.1 Complete Functions

Complete functions will consume one event from every input and produce one event at every timestamp they are evaluated. Most complete functions are pretty simple and often have eventstreams as inputs or only have one input. The complete functions that are present in the implemented runtime are explained in Table 4.1.

The first four functions are sources which take an external event and format them for internal use, e.g. *variable_values* takes a string containing the name and the value of a variable, casts the value to an appropriate type and produces a signal holding that produced value.

4.10.2 Output Complete Functions

Output complete functions will produce a new event everytime they are evaluated but only have to consume events from some inputs and not from all. Table 4.2 summarizes all input complete functions.

4.10.3 Input Complete Functions

Input complete function consume one events from every input but can produce zero or one output events everytime they are evaluated. Table 4.3 summarizes all input complete functions.

Name	Domain	Range	Explanation
<i>instruction_executions</i>	Events	Events	Converts a trace to an event that denotes the execution of a specific instruction in the monitored program.
<i>function_returns</i>	Events	Events	Converts a trace to an event that denotes the return from a function in the monitored program.
<i>function_calls</i>	Events	Events	Converts a trace to an event that denotes the call of a function in the monitored program.
<i>variable_values</i>	Events	Signal	Converts a trace to a change that denotes the value of a variable in a monitored program.
<i>signalAbs</i>	Signal	Signal	Computes the absolute value of a signal.
<i>eventAbs</i>	Events	Events	Computes the absolute value of an event.
<i>changeOf</i>	Signal	Events	Emits an event everytime the signal changes its value holding the new value.
<i>neg</i>	Signal	Signal	Emits the mathematical opposite of the value of a real signal.
<i>signalNot</i>	Signal	Signal	Emits the Boolean negation of a Boolean signal.
<i>eventNot</i>	Events	Events	Emits the Boolean negation of a Boolean event.
<i>eventCount</i>	Events	Signal	Emits a signal holding the number of times an event occurred on the input.
<i>timestamps</i>	Events	Events	Emits an event holding the timestamp of an input event everytime one occurs.
<i>sma</i>	Events	Events	Emits an event holding the simple moving average over the last specified number of events that occurred.

Tab. 4.1: List of complete functions

4.10.4 Incomplete Functions

Incomplete functions always consume at least one input from any input and will produce zero or one events. Table 4.4 lists all supported incomplete functions. If no explanation is given, why the function is incomplete it is the following: The function is not input complete, because it only consumes events or changes that have the timestamp at which it is evaluated, if one input only has events with bigger timestamps no event or change is removed from them and the remembered last change of them is used as a base for computation if it is a signal. Also it is

Name	Domain	Range	Explanation
<i>merge</i>	Events \times Events	Events	Merges two eventstreams. When an event is present on the first input, will emit an event with the same value, else with the value from the event on the second input. Not input complete because if an event on the second input occurs at a timestamp where no event of the first input occurs no event of the first input is removed.
<i>occurAny</i>	Events \times Events	Events	Emits an event without a value everytime an event occurs on any input. Not input complete because events are only removed from both inputs if they have the same timestamp.

Tab. 4.2: List of output complete functions

Name	Domain	Range	Explanation
<i>signalMaximum</i>	Signal	Signal	Emits a change everytime the input has a bigger value than it had anytime before.
<i>eventMaximum</i>	Events	Signal	Emits a change everytime the input has a bigger value than it had anytime before or a default value if it is the biggest value occurred yet.
<i>signalMinimum</i>	Signal	Signal	Emits a change everytime the input has a smaller value than it had anytime before.
<i>eventMinimum</i>	Events	Signal	Emits a change everytime the input has a bigger value than it had anytime before or a default value if it is the biggest value occurred yet.
<i>sum</i>	Events	Signal	Emits the summed up value of all events that happened on the input upto that point.
<i>mrw</i>	Events	Signal	Emits a change everytime the input takes a new value. Not output complete because no new change is emitted if the last value of the input was the same as the current.

Tab. 4.3: List of input complete functions

not output complete, because changes of a signal are only produced if the value actually changes. For example, if the values of the inputs of an *add* are switched at a timestamp it would not produce a new change for that timestamp but consume changes from both inputs.

Name	Domain	Range	Explanation
<i>add</i>	Signal \times Signal	Signal	Adds both inputs.
<i>and</i>	Signal \times Signal	Signal	Performs a Boolean and over both inputs.
<i>div</i>	Signal \times Signal	Signal	Divides the first input by the second input.
<i>eq</i>	Signal \times Signal	Signal	Emits if both inputs are equal.
<i>geq</i>	Signal \times Signal	Signal	Emits if the first input is greater or equal to the second input.
<i>gt</i>	Signal \times Signal	Signal	Emits if the first input is greater than the second.
<i>implies</i>	Signal \times Signal	Signal	Emits the Boolean implies relationship between both inputs.
<i>leq</i>	Signal \times Signal	Signal	Emits if the first input is smaller or equal to the second.
<i>lt</i>	Signal \times Signal	Signal	Emits if the first input is smaller than the second.
<i>max</i>	Signal \times Signal	Signal	Emits the bigger value of both inputs.
<i>min</i>	Signal \times Signal	Signal	Emits the smaller value of both inputs.
<i>mul</i>	Signal \times Signal	Signal	Multiplies the first input by the second.
<i>or</i>	Signal \times Signal	Signal	Performs a Boolean or over both inputs.
<i>sub</i>	Signal \times Signal	Signal	Subtracts the second input from the first.
<i>filter</i>	Events \times Signal	Events	Emits events whenever an event occurs on the first input with the value of that event if the second input has the value true. It is not output complete because it doesn't emit events when the second input is false.
<i>ifThen</i>	Events \times Signal	Events	Emits an event with the value of the second input everytime an event occurs on the first input. It is not output complete because it only emits outputs when an event occurred on the first input.

Name	Domain	Range	Explanation
<i>ifThenElse</i>	Signal \times Signal \times Signal	Signal	Emits the value of the second input if the first is true, else of the third input.
<i>sample</i>	Signal \times Events	Events	Same as <i>ifThen</i> with switched arguments.
<i>occurAll</i>	Events \times Events	Events	Emits an event whenever events occur on both inputs. Not output complete because it only emits events whenever events happen on both inputs.

Tab. 4.4: List of incomplete functions

4.10.5 Timing Functions

Timing functions are a bit special, therefore they are mentioned here in their own section. Basically they are also incomplete functions, but they have to buffer multiple events until they are emitted and they can produce more than one event when evaluated. TODO

4.11 State and History

All TESSLA evaluation engines have to hold a state, which encodes information necessary to continue the evaluation, and a history, which encodes what happened on all streams in the evaluation engine. The state of a whole evaluation engine is made up of the states of its nodes.

Each node has a state, which contains arbitrary information, e.g. a counter for a *CountNode*, its input queues holding the non-processed events and, if they have signals as inputs, the last changes of them.

To distinguish between the two types of states, the state of the whole engine is called the *global state* and the state of a single node the *node state*. The set of all valid node states is called \tilde{N} .

The global state of an evaluation engine at a certain step is a map from its nodes to their node state. We will denote the set of all global states as S . A global state can be queried like $s(n) = \tilde{n}$ to yield the state of the node n .

Nodes, and therefore the whole evaluation engine, change their state when they are scheduled. The transition between states is described in Section 4.12

The history of an evaluation engine is defined at every step (read: after every computation of a node) as all events that were produced by any node upto that step.

4.12 Transitions

A transition describes what happens when the evaluation engine schedules a node: Events from the the inputs of a node are removed (at least one), output events can be generated (but don't have to) and distributed and the internal state of nodes are updated. Because the function of a node is evaluated at the evaluation timestamp, which is the minimal timestamp of all events on the heads of inputs, the events which are removed are exactly the ones that have the evaluation timestamp. To look at it in another way: a transition models the computation of a node and the progressing of the stream it produces towards the evaluation timestamp, which has to be bigger than the previous progress, since input queues are strictly ordered by their timestamp and the events with the minimal timestamps are removed after the computation. Therefore when we say 'Node a is scheduled' we mean that a transition is taken which models the computation of that node.

The set of all transitions is written as T . The function $node : T \rightarrow N$ returns the node of which the transitions models the computation.

One part of a transition is a relation between two sets of events, why sometimes we write $\tau = (\{e_1, e_2\}, \{e_3\})$ to visualize a transition, but remember that there is more to a transition than that. E.g. the relation $\tau = (\{e_1, e_2\}, \{e_3\})$ means that two events were consumed by a node and one event was produced based on them.

The empty transition, meaning no input was consumed and no output produced, is labeled with λ . Note that all transitions, which produces events have to consume at least one event (therefore no events can be created from nowhere) and that it's possible that no event was produced based on the consumed events (see Section 4.10). Furthermore with timing functions it's possible to create multiple events in one transition. For example think of an *EchoNode*, which duplicates an input after a specified amount of time.

Definition 12: Application of a Transition on a State.

Given global state s_0 and transition $\tau = (\tilde{E}, \tilde{E}') = (\{e_1, e_2, \dots, e_i\}, \{e'_1, e'_2, \dots, e'_i\})$

with $n = \text{node}(\tau_1)$ and N_c the set of all nodes that are children of n . When we apply τ to s_0 , written $\text{apply}(s_0, \tau) = s_1$, we get a new global state s_1 with

```

 $\forall \tilde{n}_i = s_0(i)$ 
  if  $\text{node}(\tilde{n}_i) \notin N_c \cup \{n\}$ 
    then  $s_1(i) = \tilde{n}_i$ 
    (nothing changes for independent nodes)
  else if  $\text{node}(\tilde{n}_i) \in N_c$ 
    then append all events in  $\tilde{E}'$  to the input representing the stream from  $n$ 
  else
    remove all events in  $\tilde{E}$  from the inputs representing their streams
    and update the internal information based on the function  $n$  is modelling

```

This means that the new global state is built with the old global state by altering only the node states of the node identified by the transition and its children. The node states are altered by removing all events that were consumed by the function from the inputs of the scheduled node, updating its internal state and adding the produced events to the queues representing it in the node states of its children.

4.13 Run

A run of an evaluation engine is a sequence of transitions and states. The first element of the sequence is the empty transition and the initial state of the evaluation engine. It is a representation of the steps the engine takes to evaluate a specification over input streams. The length of a run can be retrieved with $\text{length}(r) = d \in \mathbb{N}$. A run can be queried by its index to return the element at that index: $r(i) = (\tau_i, s_i)$, $i \in [0, \text{length}(r)]$.

The run $\langle (\lambda, s_0), (\tau_1, s_1) \rangle$ means, that the engine was in its initial state, took the transition τ_1 and thereby reached the state s_1 .

Definition 13: Closeness of Runs.

The closeness δ of a run r_1 to a run r_2 is a pair $\delta(r_1, r_2) = (x, y)$, where x is the index before the first position where the two runs differ and y is the number of steps between the index of the first difference and the position where r_2 takes the transition that r_1 took after step x . The closeness of runs is ordered element-wise: $(x, y) > (x', y') \leftrightarrow ((x > x') \vee (x = x' \wedge y < y'))$. Therefore two runs with length d are equal, if their closeness is $d, 0$, which is the maximal closeness two runs of length d can have at all. Note that two runs have no closeness if they don't contain the same transitions.

Example 8: Closeness of Runs.

Let

$$r_1 = \langle (\lambda, s_0), (\tau_1, s_1), (\tau_2, s_2), (\tau_3, s_3), (\tau_4, s_4), (\tau_5, s_5), (\tau_6, s_6) \rangle$$

$$r_2 = \langle (\lambda, s_0), (\tau_1, s_1), (\tau_2, s_2), (\tau_5, s'_3), (\tau_4, s'_4), (\tau_6, s'_5), (\tau_3, s'_6) \rangle$$

$$r_3 = \langle (\lambda, s_0), (\tau_1, s_1), (\tau_2, s_2), (\tau_3, s_3), (\tau_5, s''_4), (\tau_4, s''_5), (\tau_6, s''_6) \rangle$$

Then is

- $\delta_{1,2} = \delta(r_1, r_2) = (3, 3)$ because r_1 takes τ_3 at step 3 while r_2 takes τ_5 and r_2 takes τ_3 at step $3 + 3 = 6$.
- $\delta_{1,3} = \delta(r_1, r_3) = (4, 1)$ because r_1 takes τ_4 at step 4 while r_3 takes τ_5 and r_2 takes τ_4 at step $4 + 1 = 5$.
- The explanations for the remaining cases is analogous and therefore not stated here.
- $\delta_{2,1} = \delta(r_2, r_1) = (3, 2)$, $\delta_{2,3} = \delta(r_2, r_3) = (3, 1)$
- $\delta_{3,1} = \delta(r_3, r_1) = (4, 1)$, $\delta_{3,2} = \delta(r_3, r_2) = (3, 2)$

The ordering of the distances is straightforward: $\delta_{1,3} < \delta_{2,3} < \delta_{2,1} < \delta_{3,1}$.

To reason about runs we have to restrict runs to the ones that are reasonable in the context of an evaluation engine. This means that only transitions are taken that are possible based on the global state. To do so we have to define when a node can compute based on its state. At first we will give a definition that only works for output complete TESSLA functions. Based on that definition we will see the problems that output incomplete functions have and modify the transition model to fix the problem.

Definition 14: Enabledness of a Node.

A node n with the node state \tilde{n} containing the input queues $\tilde{\sigma}$ in an evaluation engine is called enabled at a step i of a run r of that engine, if at least one input is buffered on each input queue:

$$\forall \sigma_x \in \tilde{\sigma} : \neg \text{empty}(\sigma_x)$$

Now it is possible to restrict runs to a subset where each run models a rational evaluation of a specification.

Definition 15: Valid Run.

A run r is called valid, if

- $\forall i \in [0, \text{length}(r)] : r(i) = (\tau_i, s_i) \wedge \text{node}(\tau_i) \text{ is enabled at step } i$
- and $s_i = \text{apply}(s_{i-1}, \tau_i)$

As stated, the definition for enabledness works for all output complete functions, while output incomplete functions have a problem. It is possible that an output incomplete functions will never produce a new output: E.g. think of a *FilterNode* where the second input is always *false*. Even if new input events are added to the first input and the second input is known to be *false* upto any timestamp, the children of the node will never receive new events for that input queue. Therefore no children can ever again compute, because they don't have an event buffered on all inputs.

There are two ways to fix this:

All input queues could be extended with a progress timestamp, which denotes how far the parent node has progressed. Now everytime a node n with children N_c evaluates its function at a new evaluation timestamp the inputs of all nodes in N_c representing the stream of n would be updated to have that evaluation timestamp as their new progress, even if no new event was produced at that timestamp.

The other way, which is used in this thesis, doesn't alter the input queues: First we need a new type of events: *progress events*. A progress event holds no value and exists only to notify a node that a specific input has progressed upto the timestamp of the progress event. We will write a progress event as e_π^p where π is the timestamp of it.

Now whenever a node performs its computation at an evaluation timestamp and no new event was produced a progress event with the evaluation timestamp is added to the corresponding input queue of all children. When a node is scheduled and one of the consumed events is a progress event, what happens depends on the function of the node. Some examples are:

- An *AddNode* has at an evaluation timestamp a progress event as the first input and no event as the second input, which means there is an event buffered on input queue two, else the node wouldn't have been scheduled, and that event has a bigger timestamp than the evaluation timestamp, else its timestamp would have been the evaluation timestamp. The node can't produce a new change for its output, since all new information is, that input one hasn't changed upto the new timestamp. Therefore the node will distribute a progress event with the evaluation timestamp to its children.
- If the *AddNode* would have received a progress event on input one and a new change on input two at the same timestamp, it could emit a new change. The new event would have the value of the last change on the first input, which the node has to hold in its state, added with the value of the new change on the second input.

- A *MergeNode* has a progress event as the first input at an evaluation timestamp and no event on the second input. By the same reasoning as in the first example, the node can't produce a new output and therefore will produce a progress event.

These examples show that progress events alters the categorization of TESSLA functions: output complete functions could produce no new normal output, if inputs are progress events. This is no problem, since the main reason for the categorization was to explain the necessity of progress events. With progress events all functions are output complete, because they always emit either normal events or progress events.

For the comparison of runs some more definitions are needed.

Definition 16: Independence of Nodes.

A node a is called independent of node b in an evaluation engine, if a is no descendant of b .

Definition 17: Independence of Transitions.

A transition τ_1 is called independent of another transitions τ_2 , if $node(\tau_1)$ is independent of $node(\tau_2)$.

Lemma 3: Exchange of Independent Transitions.

If a transition τ_2 is independent of a transition τ_1 , then for all runs of the evaluation engine that produces the runs the following holds:

*If $r_1 = \langle (\lambda, s_0), \dots, (\tau_1, s_i), (\tau_2, s_j), \dots, (\tau_l, s_l) \rangle$ is a valid run
then $r_2 = \langle (\lambda, s_0), \dots, (\tau_2, s'_i), (\tau_1, s_j), \dots, (\tau_l, s_l) \rangle$ is a valid run*

Proof. As a first step we will show, that the transitions τ_1 and τ_2 can be exchanged because their enabledness doesn't depend on each other.

Because $b = node(\tau_2)$ is no descendant of $a = node(\tau_1)$, the stream σ with $node(\sigma) = node(\tau_1)$ can be no input of b . Therefore the enabledness of b can't be changed by τ_1 , taken directly from the definition of enabledness. So b has to be enabled before τ_1 was taken in r_1 or else it couldn't be enabled afterwards and τ_2 couldn't be taken in the next step. Therefore r_2 also fulfills the requirements of a valid run up to and including the steps $(\tau_2, s'_i)(\tau_1, s_j)$.

As a second step we have to show, that the state s_j will stay the same, no matter the order of the two transitions and only the state s_i may be changed to s'_i in r_2 . If this holds, the subsequent states can't change either, since they are deterministically built by applying the same transitions. Now if all subsequent states stay the same, all subsequent transitions will stay enabled, since enabledness of a node only depends on the state. Figure 4.1 visualizes the exchange of the transitions and the argument why the state stays the same.




Fig. 4.1: Influence of the order of independent transitions on the global state of an evaluation engine

Remember that $s_i = \text{apply}(s_{i-1}, \tau_1)$ and $s'_i = \text{apply}(s_{i-1}, \tau_2)$. We have to show that $\text{apply}(\text{apply}(s_{i-1}, \tau_1), \tau_2) = \text{apply}(\text{apply}(s_{i-1}, \tau_2), \tau_1)$. This is straightforward when recalling what a transition does when it's being applied. $\text{apply}(\text{apply}(s_{i-1}, \tau_1), \tau_2)$ appends all events produced by τ_1 to the queue representing $n_1 = \text{node}(\tau_1)$ in the node states of the children of n_1 and updates the node state of n_1 . Then it appends all events produced by τ_2 to the queue representing $n_2 = \text{node}(\tau_2)$ in the node state of the children of n_2 and updates the node state of n_2 .

Since n_1 and n_2 are independent of each other, neither of them have input queues representing the other, therefore no event can be added to a queue of n_2 with t_1 and the same with n_1 and τ_2 . Therefore all events that were appended by applying τ_1 will still be in the queues after τ_2 was applied, since only events from the queues of n_2 can be consumed, and the same if they were applied in reversed order.

To conclude: All events appended to any node states will still be present after both transitions were applied and the node states of n_1 and n_2 will be updated based on the same events, no matter in which order the transitions are run. Therefore the state after both transitions are applied have to be the same, no matter the order in which they are applied.

□

Lemma 4: Duration of Enabledness.

A node which is enabled stays enabled at least until it is scheduled. Formally: If a node n is enabled at step i in a run r it will stay enabled at least until the first step $j > i$ with $r(j) = (\tau, s), \text{node}(\tau) = n$. Note that it doesn't have to be disabled after step j , because there could have been multiple events buffered on its inputs.

Proof. Let n be a node enabled at step i in a run r . Let $(\tau_x, s_x) = r(x), x = i + 1$. If $\text{node}(\tau_x) \neq n$, then the only influence that τ_x can have on the node state \tilde{n} of n is by appending produced events to one of the input queues, as per the definition of application of a transition. Since τ_x can't remove events from the input queues in \tilde{n} and on all input queues was at least one event buffered before the transition was

applied, else n wouldn't have been enabled at step i , \tilde{n} will have at least as many input events buffered on every input as in the step before. This means that n is still enabled at step x , and by induction at every later step until a transition with n as its node is taken. \square

Lemma 5: Finiteness of Enabledness.

Whenever a node is enabled in a run, it can only be scheduled continuously a limited number of times until becoming disabled.

Proof. Let n be an enabled node at step i with the node state \tilde{n} .

First let's assume that no parent of n are scheduled after step i and before n becomes disabled. Since the input queues of n are filled by previous transitions and only finite number of events are produced at every step, all of the queues can only have a finite number of events buffered at step i . Since no parents of n are scheduled, the queues can't get fuller. Because all nodes represent functions, and all functions have to consume at least one input event (See 4.10), everytime n is scheduled at least one input queue of n will have one event less buffered after it performed its computation. Therefore the worst case is when n only consumes one event per computation. The maximum number of time n can be scheduled is therefore bounded by the sum of the number of events on all input queues at step i .

Now let's assume that also input nodes of n can be scheduled after step i . Everytime an input node is scheduled it will add a finite amount of events to one input queue of n . This leads to a cyclic behaviour: If an input queue can be scheduled infinitely often, infinite many events will be added to the input queue and n can possibly be scheduled infinitely often. If input queues can not compute infinitely often, only a finite amount of events are added to the inputs of n and therefore n can only be scheduled a limited number of times. Closer inspection of the nature of evaluation engines show why this is no problem: Only if the external trace fed to an evaluation engine is infinite the source nodes can compute infinitely often. Hence for finite traces no node can compute infinitely often. Again the worst case is when n only consumes one event per computation. The number of times it can be scheduled is limited by the sum of the number of events buffered on the inputs at step i and the number of events that are produced by inputs of n after step i . \square

With the notion of runs and especially enabledness we can now define which schedules are seen as fair.

Definition 18: Fair Schedules.

A schedule of an evaluation engine is called fair, if for all runs r it produces the following holds:

$$\begin{array}{ll} \forall i < \text{length}(r) : & \text{if } n \text{ is enabled at step } i \\ & \text{then } \exists j \geq i \text{ such that } n \text{ is scheduled at step } j \end{array}$$

In other words, every enabled node is scheduled after a finite number of steps.

Building on this we will investigate different fair schedules in the next chapter.

Behaviours of Evaluation Engines

Based on the definitions in Chapter 4 we will now look at different schedules of evaluation engines and compare them. This is done in multiple steps: starting with a small subset of allowed schedules and functions and iteratively adding more complex cases.

For the comparison we will use timed transducers from Section 4.3. To do this the notion of a run of an evaluation engine is not sufficient: transducers describe a relationship between inputs and outputs, runs describe stepwise generation of internal and output events. Therefore the *behaviour* of a run is defined, which maps a run to relationship between inputs and outputs.

Definition 19: Behaviour of a Run.

Let r be a run of an evaluation engine. The behaviour β_r of it is a timed transducer: A set of tuples of timed sequences. It is calculated as follows:

1. Let $\beta_r = \emptyset$ and r_p an empty prefix of r .
2. Remove the prefix from r , where the first transition consumes an input upto but not including the next transition where an input is consumed, and append it to r_p .
3. Select the sequence of all output events O_p (which is possible empty) that are produced at any step in r_p .
4. Select the sequence of all input events E_p that are consumed at any step in r_p .
5. Add the tuple (E_p, O_p) to β_r .
6. Goto step 2 if r is not empty, else terminate.

Stated simple the run is chopped into pieces, where each piece begins with the consumption of an input events and ends before the next input is consumed. The pieces are then merged from left to right: First take all inputs and outputs consumed and produced upto the current piece and add them to the behaviour, then merge the current piece with the next and repeat.

Example 9: Construction of a Behaviour.

In Table 5.1 it is shown how a behaviour is built from a run. The run is denoted only by its transitions, events are labeled based on their type: e_i are external (read: input)

Run $(\{e_1\}, \{i_1\}) (\{i_1\}, \{i_2\}) (\{i_1, i_2\}, \{o_1\}) (\{e_2\}, \{i_3\}) (\{i_2, i_3\}, \{o_2\}) \leftarrow$
 $(\{i_3\}, \{o_3\}) (\{e_3\}, \{i_4\}) (\{e_4\}, \{i_5\}) (\{i_4, i_5\}, \{o_4\})$
 Tuples $(e_1, o_1) (e_1e_2, o_1o_2o_3) (e_1e_2e_3, o_1o_2o_3) (e_1e_2e_3e_4, o_1o_2o_3o_4)$

Tab. 5.1: Example how the behaviour of a run is constructed

events, i_i are internal events and o_i are output events. Parts of the run in the same color are the transitions that end up in the same piece when applying the construction algorithm. The tuples show how the sequence of input and output events from the pieces are extracted. The behaviour of the run is the set of the tuples.

The behaviour of a run is a timed transducer since all events have timestamps and all consumed events are strictly ordered by their timestamp, since inputs to an evaluation engine are required to be ordered by their timestamp.

Since the behaviour encodes the relationship between inputs and outputs a run produces it provides the foundation to reason about equivalence between different runs and whole evaluation engines.

Definition 20: Equivalence of Runs.

Two runs are called equivalent if their behaviour is observational equivalent.

Now we can define when two evaluation engines are called equivalent based on their runs

Definition 21: Equivalence of Evaluation Engines.

Two evaluation engines are called equivalent, if for every run that one can produce there is an observational equivalent run in the other.

5.1 Schedules Without Timing Functions

For a first step we specify and compare behaviours of different approaches to evaluate TESSLA specifications without timing functions. Without timing functions all nodes work only on values or the presence of events and will emit exactly one event at every computation, either a normal or a progress event. This leads to behaviours that can be easily reason about, as seen in the next sections.

All systems to evaluate TESSLA specifications we will look at are based on the described structure in Section 4.9. While there are other approaches to evaluation, a DAG based approach seems to fit most naturally and focusing on one structure makes comparing systems easier.

Let's recap and summarize how an evaluation engine performs its computation. Each evaluation engine will work in steps, where each step is synonymous with

an index in the run of the system. Therefore at each step one enabled node is scheduled to perform its operation, represented as the transition in the run. The transition will encode one of the following three things that can happen:

- The next external event (external events have to be totally ordered by their timestamp) can be consumed by a source in the DAG, which generates internal events, that are propagated to its children.
- An internal node, which has at least one new input buffered on all of its input queues, can perform its computation and generate a new internal event, which is propagated to the children of that node.
- An output node, which has at least one new input buffered on all of its input queues, can produce a new output.

Evaluation engines are free in the way they are scheduling their nodes, only limited by causality (no event can be consumed before it's produced), which is guaranteed by the enabledness criteria. In the following evaluation engines are classified by their scheduling approaches.

5.1.1 Greedy Evaluation Engines

The first class of evaluation engines are called greedy.

Definition 22: Greedy schedule.

A schedule of an evaluation engine built by the following steps is called greedy.

1. *Select all nodes that are no sources, let their count be i*
2. *Label them with unique natural numbers from $[1, i]$ in reverse topological order*
3. *Label the remaining nodes with unique natural numbers bigger than i*
4. *Schedule the enabled node with the lowest label first*

We also call an evaluation engine greedy if we mean it's run with a greedy schedule.

Obviously for many DAGs there is no unique reverse topological order, therefore one can be chosen by the evaluation engine. We will show in Section 5.2.1 that all topological orders will produce observational equivalent behaviours.

The greedy schedule ensures that no node is scheduled which has a successor that can be scheduled, therefore events are *pushed* through the DAG towards an output node as fast as possible. As shown in Section 5.1.1 any schedule built like this is fair.




Fig. 5.1: Visualization of a simple evaluation engine with a greedy schedule.

Greedy evaluation engines offer a good start to reason about behaviours and will be used as a comparison for all other evaluation engines.

Definition 23: Valid Evaluation Engines.

An evaluation engine is called valid if it is equivalent to a greedy evaluation engine.

Figure 5.1 visualizes a greedy evaluation engine. It shows two DAGs representations of an evaluation engine where the nodes a to e are labeled in a reversed topological order and o and u represents two output streams. The left system is in its initial state and an input event e_1 is present and can be consumed by the input node a . When a node is chosen to compute by the scheduler, only node a is enabled, therefore it is scheduled. The right system is the representation of the next step: node a has consumed the external event and produced an internal event n_{a1} , which is propagated to all its children: nodes b and d . In the next step node, b would be scheduled, because it has the lowest number of any node that can compute (actually it's the only node that can compute at all, because d has to wait for the event from b). After b is scheduled, it would produce the internal event n_{b1} which would then be distributed to nodes c and d .

The complete run of the greedy engine for one input is the following, where the states are omitted:

$$\begin{aligned} &\langle (\lambda, s_0), ((\{n_{a1}\}, \{n_{b1}\}), s_1), ((\{n_{b1}\}, \{o_1\}), s_2), \\ &\quad ((\{n_{a1}, n_{b1}\}, \{n_{d1}\}), s_3), ((\{n_{d1}\}, \{u_1\}), s_4) \rangle \end{aligned}$$

If there were more than one input event, at this point node a would be scheduled again. It would consume the next external event and the following nodes would be scheduled in the same order as before, extending the run in an obvious way.

Fairness of Greedy Schedules

It remains to show that greedy schedules are fair.

Lemma 6: Greedy Schedules are Fair.

Any greedy schedule is fair.

Proof. Let a be a node with the label n , which is enabled at step i and is no source. Because evaluation engines can only contain a finite number of nodes there can only be a finite number of enabled nodes with a smaller label than n . Because of Lemma 5, all nodes with a smaller label than n will become disabled after a finite number of steps. Let that number be j . The only way new events could enter the system are through sources, but they have bigger labels than n , as by the definition of the schedule, and therefore can't be scheduled before n . Because of Lemma 4, a will still be enabled after these steps. So a is the enabled node with the lowest label at step $i + j$ and therefore will be scheduled.

Now let a be a source. Sources are only scheduled when no internal node is enabled since they are labeled with higher numbers than all internal nodes. Based on the same reasoning as in the first case at some point all internal nodes will become disabled, therefore a source node has to be scheduled. This source can either be a or another source, recall that only one source is enabled at any time because of the environment of an evaluation engine. If another source was scheduled, after a finite amount of steps all internal nodes will have to become disabled again. Since finite traces are evaluated at some point either the trace will end without ever feeding an input to a , then a will never be enabled, or at some point a will receive an external event. When a receives an external event, it will be the only enabled source, else no input would be fed to an input at that step. Therefore a will be scheduled the next time no internal nodes are enabled. \square

5.1.2 Fair Evaluation Engines

Obviously greedy schedules are only a small subset of all fair schedules. As the next step we will look at the rest of them.

Definition 24: Fair Evaluation Engines.

A fair evaluation engine is one with a fair schedule.

In contrast to a greedy evaluation engine a fair one has no fixed schedule, meaning that at each step any enabled node can be scheduled. Therefore predecessors of enabled nodes can perform multiple computations before their children are scheduled and events are not *pushed* through the DAG as fast as possible.

The difference between greedy and fair schedules are similar to the ones of synchronous and asynchronous transducers: A greedy schedule will ensure that outputs are produced as fast as possible while a fair can *delay* the outputs by consuming multiple inputs first and scheduling internal nodes multiple times before scheduling an output node. But note that there is an important difference between synchronous transducers and the behaviour of a greedy evaluation engine: A greedy evaluation engine can produce multiple events at every step.

5.2 Equivalence of Different Schedules Without Timing Functions

The behaviour of a run of an evaluation engine with a given schedule allows us to reason about equivalence.

As by Definition 23 any evaluation engine has to be equivalent to a greedy one to be valid.

The equivalence is shown in two steps: first in Section 5.1.1 it is shown, that all possible greedy engines for a specification are equivalent, so there is only one valid evaluation for a specification over a fixed input. Afterwards in Section 5.2.2 it is shown that any fair evaluation engine is equivalent to a greedy one.

5.2.1 Equivalence of Greedy Systems

When given a series of input events, two greedy evaluation engines for a specification with different schedules will have different runs. But both will produce all outputs that can be produced after consuming one specific input before the next input is consumed as reasoned in Section 5.1.1. Also both runs will obviously have the same length (both engines are the same DAG, so they have the same number of nodes), let that length be l .

To proof the equivalence of both engines we can prove the equivalence of their runs. To show the equivalence it is shown that two runs r_1 and r_2 of two evaluation engines based on the same graph but with a different greedy schedule can always be reordered to become closer while preserving observational equivalence. If such

a closer run always exist, we will show that the run with closeness $(l, 0)$ to r_1 , which has to be r_1 itself, is also observational equivalent to r_2 .

Theorem 1: Equivalence of Different Greedy Evaluation Engines.

Two greedy evaluation engines for a specification with different schedules are equivalent.

Proof. Let r_1, r_2 be the runs of two greedy evaluation engines for the same specification which received the same inputs. Because each TESSLA specification contains only a finite amount of functions and works on finite traces, the runs also have to be finite.

If the two runs aren't equal, they must have a closeness which is smaller than $(l, 0)$. Let $[r_2]$ be the set of all runs that are observational equivalent to r_2 . If r_1 is in this set, we would be done. Let's assume that r_1 is not in the set. Therefore all runs in the set have to have a smaller closeness than $(l, 0)$ to r_1 , since the only run with a closeness of $(l, 0)$ to another run is the run itself. Select one run $r'_2 \in [r_2]$ which has the biggest closeness to r_1 . Let $(d, k) = \delta(r_1, r'_2)$.

This means that at step d the run r'_2 has taken a different transition than run r_1 . Let the transitions the runs have taken be τ_1 for r_1 and τ_2 for r'_2 . Run r'_2 will take transition τ_1 at step $d + k$ (as per the definition of the closeness). Obviously the two transitions have to be independent of each other, else they couldn't have been taken in different order by the two runs.

If $k > 1$ there will be a transition $\tau'_2 \neq \tau_1$ which is taken by the run r'_2 at step $d + (k - 1)$. While this transition τ'_2 must also be taken in the first run as per Lemma 4, it's not possible, that it was taken before τ_1 , because then the two runs wouldn't have been the same upto the point where τ_1 was taken. Therefore τ_1 has to be independent of τ'_2 , and because τ'_2 was scheduled by the second run before τ_1 both transitions are independent of each other.

As of Lemma 3 which one of them is taken first won't change the rest of the run at all after both transitions were applied and the run stays valid. Therefore there is a valid run r''_2 , which is equal to r'_2 , except that the transitions τ_1, τ'_2 are scheduled the other way around. See Figure 5.2 for a visualization of the runs. Dotted edges represent multiple transitions. The visualization shows how the three runs *branch* and *merge* after certain steps. The run r_1 is marked in red, r''_2 in green and r'_2 in blue. As you can see the run r''_2 has a bigger closeness to r_1 since it takes τ_1 before τ'_2 . Note that r_1 is only known upto step d , especially the transitions τ_2, τ'_2 will be taken eventually, but it isn't known when or in which order.

When two adjacent transitions are exchanged basically two things can happen:




Fig. 5.2: Visualization of three runs of an evaluation engine as explained in Proof 5.2.1.

- Outputs can be produced later, maybe moving them to the next piece in the construction of a behaviour (this happens when a transition producing outputs is exchanged with the next transition which consumes an external event)
- Outputs can be produced earlier, which can move them to an earlier piece in the construction of the behaviour (this happens in the opposite case, where a transition consuming an external event is pushed before one that produces outputs)

While the order of outputs in the tuples of the behaviour will change when two transitions producing outputs are exchanged, observational equivalence isn't influenced, since it is defined over reordered outputs.

Let's test if r_2'' is observational equivalent to r_2' : We will compare how the behaviour of the run r_2' changes when the transition τ_1 is exchanged with τ_2' . The comparison is based on the different cases the transitions can encode. Some of the cases are not feasible, they are listed for the sake of completeness and to explain why they can't happen. The cases are:

1. No inputs are consumed and no outputs produced by either transition. This obviously won't change the behaviour at all.
2. τ_2' consumes an input and doesn't produce an output and
 - a) τ_1 doesn't consume an input and doesn't produce an output. This won't change the behaviour, since τ_1 didn't add anything to it in the first place.
 - b) τ_1 doesn't consume an input, but produces one or more outputs. This changes the behaviour, since τ_1 is now part of the piece starting before τ_2' . Therefore the produced outputs are now part of one more tuple of the behaviour. But note that it still is part of all tuples built from the later pieces of the run. Therefore the two behaviours are still observational equivalent: All tuples in the behaviour are still the same, except that one or more output events are produced one step earlier, which doesn't hurt observational equivalence.
 - c) τ_1 consumes an input but doesn't produce an output. This case can't happen, since inputs to evaluation engines are totally ordered by their timestamp and therefore τ_2' couldn't have been scheduled after τ_1 in r_1 and before τ_1 in r_2 .
 - d) τ_1 consumes an input and produces outputs. This case can't happen for the same reason.
3. τ_2' produces one or more outputs and doesn't consume an input and

- a) τ_1 doesn't consume an input and doesn't produce an output. This won't change the behaviour, since τ_1 didn't added anything to it in the first place.
 - b) τ_1 doesn't consume an input, but produces one or more outputs. This only changes the order of the output events in the behaviour, but since they are reordered by their timestamp and the order of events with the same timestamp isn't important for observational equivalence, the new run is still observational equivalent.
 - c) τ_1 consumes an input but doesn't produce an output. This will *delay* the production of the outputs from τ_2' by one piece of the chopped run. While this changes the behaviour it preserves observational equivalence.
 - d) τ_1 consumes an input and produces outputs. This is kind of a combination of the previous two cases. The outputs from τ_2' are *delayed* by one piece and the outputs of τ_1 are now produced before them. But still all outputs are produced, only in different order and maybe one step later. Therefore observational equivalence holds. Also note that such a transition isn't very useful as argued in the next case.
4. τ_2' produces one or more outputs and consumes an input. First of all note that this is a rather made up combination that can only happen when a source is an output node at the same time and therefore doesn't have much of a purpose. But for the sake of completeness let's look at the cases following from this:
- a) τ_1 doesn't consume an input and doesn't produce an output. Again won't change the behaviour, as in earlier cases.
 - b) τ_1 doesn't consume an input, but produces one or more outputs. Preserves observational equivalence since the outputs of τ_1 are only produced one step earlier than before.
 - c) τ_1 consumes an input. This case can't happen as reasoned by Case 2c.

So for all cases that can happen, the run which is obtained by exchanging the two adjacent transitions is observational equivalent to the run without the change. The exchange of the two transitions brings the closeness of the new run by construction one step closer to r_1 . This means, since observational equivalence is transitive, that there is an observational equivalent run to r_2 , the run r_2'' , which has at least the closeness $(d, 1)$.

If $k = 1$ the transition τ_2' from the previous case is equal to τ_2 . The reasoning for all cases stays exactly the same, in the end we will obtain a run r_2'' which is observational equivalent to r_2' but has the closeness $(d, 0)$ to r_1 . This obviously doesn't make sense: The first element of the closeness is the last step where both runs

are equal, the second element describes how many steps afterwards the differing transition was taken. But if it was taken right in the step after the last equal step, there is no difference at that position, so the closeness of r_1 and r_2'' can be at least $(d + 1, x), x \in \mathbb{N}_{>0}$. This also contradicts our initial statement that r_2' was the run with the biggest closeness to r_1 which is observational equivalent to r_2 .

Combined we can now say, that there is no upper bound on the closeness of observational equivalent runs of r_2 to r_1 , therefore the run with the closeness $(l, 0)$ also has to be equivalent to r_2 . And as already stated, only the run r_1 can have the closeness of $(l, 0)$ to r_1 . Therefore r_1 has to be observational equivalent to r_2 . \square

This characteristic of greedy schedules gives us a baseline to compare other schedules to. Since all greedy schedules of an evaluation engine produce observational equivalent behaviours we can choose any run produced by such a schedule and compare any other run to it. In the next section we will do this for all fair schedules.

5.2.2 Equivalence of Greedy and Fair Evaluation Engines

Let's recap what fairness of a schedule means: Whenever a node becomes enabled in a run it has to be scheduled eventually. What makes fair schedules harder to reason about than greedy ones is that for one they don't have to be deterministic and furthermore that it's possible that an enabled node is not scheduled for a very long time.

Before reasoning about equivalence of greedy and fair schedules let's look at a kind of fair schedule that can be seen as worst case. Basically it is the reverse of a greedy schedule: Always schedule the enabled node that is closest to a source. Note that this schedule is not fair for infinite input traces. Stated simple this schedule will consume all input events and produce all output events per *level* of the DAG, starting at the sources and moving towards the outputs. The behaviours of runs with such a schedule are pretty special: Since no output is produced before all external events are consumed (except if a source is also an output) only one tuple of the behaviour will contain any outputs, to be specific the one which contains the sequence of all inputs as the first element. An abbreviated example of such a run is $(e_1, ()) (e_1e_2, ()) (e_1e_2e_3, ()) (e_1e_2e_3e_4, o_1o_2o_3o_4o_5)$. Such a run needs obviously more reordering than a greedy one to become observational equivalent to another greedy one since greedy schedules try to produce outputs as early as possible.

This example shows what the difference is when reordering a fair run in contrast to a greedy run: basically more transitions have to be reordered since outputs can be produced later.

Let's revisit the cases from Proof 5.2.1: Actually the cases 2b and 4b can't happen for greedy runs. If in a run r_1 a transition τ_1 , which produces an output, has to be exchanged with an earlier transition τ_2 , that consumed an external event, to become closer to a greedy run r_2 , r_1 couldn't have been greedy in the first place: Greedy schedules ensure by construction that all outputs that can be produced based on consumed external events are produced before the next external event is consumed. If τ_1 happened before τ_2 in a greedy run, the outputs from τ_1 can be produced without consuming the external event from τ_2 before, hence a run in which τ_2 happens before τ_1 couldn't be produced by a greedy schedule.

It's noteworthy that actually the whole proof of Theorem 1 doesn't depend on the fact, that the runs are produced by greedy schedules. The only real requirement is fairness, meaning that all transitions that can happen will eventually happen. Therefore the proof does hold without change for Theorem 2.

Theorem 2: Equivalence of Fair and Greedy Evaluation Engines.

Any fair evaluation engine is equivalent to a greedy evaluation engine for the same specification.

5.3 Behaviour with Timing functions

5.4 Equalitys with Timing functions

5.5 Parallel computation

Evaluation

6.1 Runtime Benchmarks

Over the course of this thesis two different approaches to handling the progress of streams in the runtime were explored, which were briefly mentioned in Chapter 4.13: progress timestamps and progress events. Progress timestamps were used first, which worked, but lead to big overhead when a great number of computations had to happen, since the whole stream had to be send to all children of the node that performed the computation. Therefore later the runtime was refactored to use progress events, so that only events had to be sent between nodes. The following benchmarks will show the performance characteristics of both approaches under different circumstances.

6.1.1 Number of Processors

Since one of the main motivations to choose erlang and BEAM as the platform for the runtime was its support for parallel execution on many processor cores, in this section the performance characteristic in regard to available cores is explored. To do so a simple TESSLA specification is given, that includes as many nodes as the maximal number of processors available. Such a spec for eighth processors is given in Listing 6.1.

The specification is evaluated over appropriate traces with a specified amount of processor cores available. The needed time of the evaluation engine from its start until it emits the conclusion, that the stream *done* is *true*, is measured for different amounts of processors. All benchmarks were performed with upto 16 cores and 50GB of RAM.

One thing that showed during benchmarking was, that the older implementation approach used big amounts of RAM and the usage grew exponential with the amount of used cores. For example with two used cores the runtime would use around 2GB RAM, with 4 Cores already over 8GB. This can be explained easily: since in the old implementation all data has to be sent between all cores while in the new implementation only generated events have to be sent around. The excessive usage of RAM even lead to timeout under some circumstances and crashes

```

1 define num_events: Signal<Int> := literal(10000)

    define add_calls: Events<Unit> := function_calls("add")
4 define add_call_sum: Signal<Int> := eventCount(add_calls)

    define overhead_0: Signal<Int> := signalAbs(add_call_sum)
7 define overhead_1: Signal<Int> := signalAbs(overhead_0)
    define overhead_2: Signal<Int> := signalAbs(overhead_1)
    define overhead_3: Signal<Int> := signalAbs(overhead_2)
10 define overhead_4: Signal<Int> := signalAbs(overhead_3)

    define done Signal<Boolean> := eq(overhead_0, num_events)

```

Listing 6.1: TESSLA specification with eight nodes on the critical path

the whole runtime. The amount of data send between processes can even lead to negative performance impacts when more cores are added.

This behaviour was one of the main reasons to switch to a new approach. To test if reduced RAM usage would lead to better performance and eliminate crashes, the old architecture was changed in a small way: Each stream was limited to only hold the last 20 events, all older events were thrown away. Since the specification used for benchmarking only works on the most recent event of each stream, this would not alter the conclusion of the engine. For example a specification computing the average of the last 21 events wouldn't work with this adaption.

While the RAM usage of the adapted runtime is not shown, the performance impact is obvious. Thus the new architecture was implemented.

TODO: describe how V2 scales with core number.

Figure ?? shows the performance characteristics of both implementations and the adapted implementation in regard to used processor cores.

6.1.2 Number of Events

TODO: also show memory consumption!

6.1.3 Number of Nodes

6.2 Practical Examples

Additionally to the theoretic benchmarks it was also important to evaluate the runtime against some more practical examples. Therefore some traces from real world examples were collected, either with the instrumentation pass or by adjusting existing logs to fit the needs of TESSLA. The programs or traces were then modified to deliberately include errors to test if appropriate TESSLA specifications will find them when evaluated in the runtime.

6.2.1 Ringbuffer Examples

The main source for tests is a ringbuffer, implemented in C, which was then modified by the instrumentation to generate traces. The code of the ringbuffer can be found in Chapter 8.1.

6.3 Instrumentation Benchmarks

After the evaluation of the runtime itself it remains to evaluate the C instrumentation program. Note that the instrumentation is mostly a proof of concept and its main purpose for now was to generate test data for the runtime. But it seems feasible that it can be optimized and extended to become a general purpose trace collection tool, therefore some benchmarks were performed.

For reliable trace collection of software the performance impact of the instrumentation is important. To measure this a trivial C program was benchmarked with and without instrumentation. The program increments each integer from 0 to 100000000 by one and, if the incremented number is divisible by a given parameter c , adds it to an intermediate result. Note that the program can and will perform some integer overflows. The code is shown in Listing 6.2.

The code is then instrumented to log each call of the add method. For each benchmark the compiled program was run 50 times. All benchmarks were performed on an Intel Core i5 with two cores and 8GB of Ram. The complete dataset can be found in Chapter 8.2.

```

#include <stdio.h>
2
int intermediate = 0;

5 int inc(int input) {
    return input + 1;
}
8
void add(int input) {
    // Instrumentation Point
11 intermediate += input;
}

14 int main() {
    for(int i = 0; i < 100000000; i++) {
        int result = inc(i);
17     if(result % 100 == 0) {
        add(result);
    }
20 }
    printf("%i", intermediate);
    return 0;
23 }
```

Listing 6.2: Example C program for benchmark purposes

6.3.1 Performance Comparison with non Instrumented Code and Compiler Optimizations

One interesting metric is, how an instrumented program performs in contrast to the same program without the added instrumentation. For this the parameter c is set to 100, so that around 1% of all function calls are instrumented.

One thing that can be recognized is that the impact of instrumentation is not predictable when using compiler optimizations. Aggressive compiler optimizations can remove function calls or inline values if it doesn't change the programs behaviour. When instrumentation calls are added no such optimization can happen, since else behaviour of the program would be altered, namely the logging would be removed when the function calls are removed.

Therefore we compare the instrumented and non-instrumented code for all optimization levels. The comparison can be seen in Figure 6.1.

The choice to instrument around 1% of function was arbitrarily, after some experimentation showed that higher amount lead to huge performance impacts as shown in Chapter 6.3.2. Intuitively it sound reasonable that in a real life example only a small subset of function calls are interesting for trace generation, since TESSLA specifications should be used to monitor critical parts of a system.

6.3.2 Performance Impact in Regard to Instrumentation Percentage

To further investigate the impact of instrumentation we will look at the performance impact in regard to the percentage of function calls that are instrumented. Therefore the variable c of the program is changed, which leads to different amount of calls to the instrumented function. For all values the program was compiled with the maximal optimization setting. In Figure 6.2 the results of changing the parameter can be seen.

Note that the x-axis is logarithmic. As expected the performance impact scales directly with the amount of calls to the instrumented function. At some point no more speed up will happen, since no calls to the instrumented functions will be made.

6.3.3 Conclusion for the Instrumentation

It can be seen, that the instrumentation adds a lot of overhead, especially when compiler optimizations are turned on and the amount of instrumented function calls are big. While the overhead is big, it is stable (see the standard deviation




Fig. 6.1: Performance of an example C program with and without instrumentation

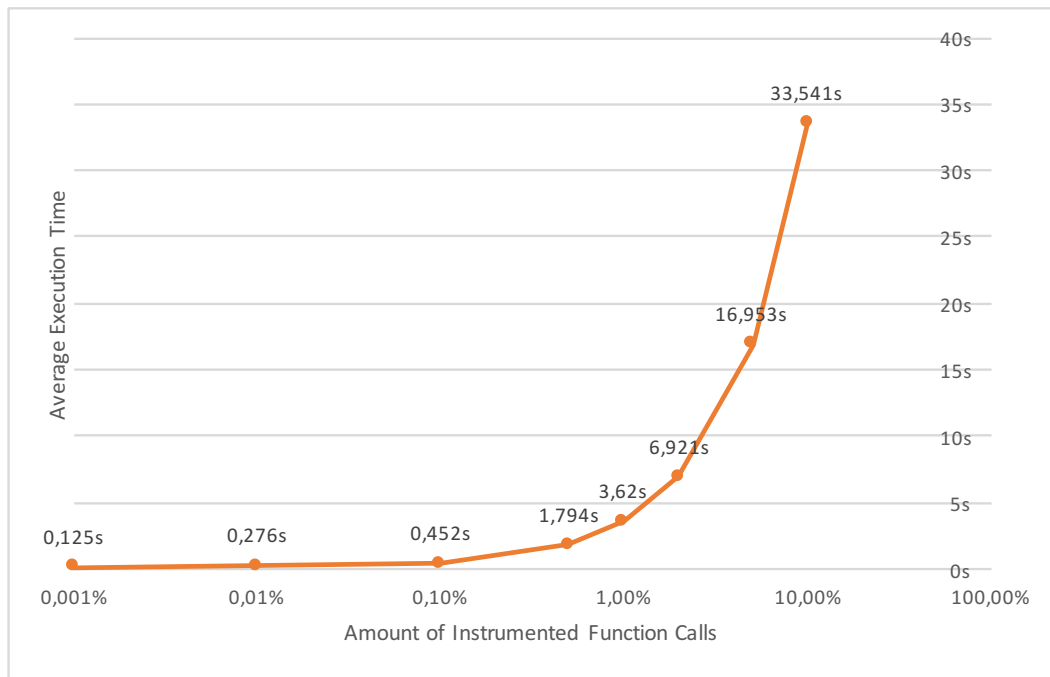


Fig. 6.2: Execution time of a program with different amount of calls to an instrumented function

in Chapter 6.3.1). This leads to the following conclusions: The generated traces should be mostly used for analysing test settings, where compiler optimizations are turned off. When used for timing specifications, the instrumented code can be benchmarked against a non instrumented version of it and the results can be used to transform the actual timing requirements to ones for the instrumented code. This will at least give an approximation of actual results. Also it is recommended to write small and specific TESSLA specifications, so that the instrumentations only have to be applied to a small amount of functions. Obviously it is strongly discouraged to use the instrumented code in any production setting.

Conclusion

7.1 Further Work

Composition of Transducers/evalEngines Port to Scala/akka Different evaluation model: Pull not push port to genstage online monitoring possible infinite traces

Architecture is similar to a vm: Tesla specs are code, compiler produces intermediate representation (json), runtime executes Ir. Therefore: Maybe define new functions (read nodes) in the spec itself and not in the runtime?

7.1.1 Error prevention

Ways of sending observations back to the program to recover from or prevent errors

Appendix

8.1 Ringbuffer Code

```
#include <stdio.h>
#include <stdatomic.h>
3 #include <pthread.h>
#include <stdlib.h>
#include <sys/time.h>
6
int BUF_SIZE = 5;

9 char* ring_buffer;
char* _Atomic read_ptr;
char* write_ptr;
12
void init_ring_buffer() {
    ring_buffer = malloc(BUF_SIZE);
15    read_ptr = ring_buffer;
    write_ptr = ring_buffer;

18 }

char* buffer_next(char* ptr) {
21    return (char*) ring_buffer + (((ptr - ring_buffer) + 1)
        % BUF_SIZE);
}

24 void process(char* data){
    struct timespec ts;
    ts.tv_sec = 0;
27    ts.tv_nsec = 1000000;
    nanosleep(&ts, NULL);
}
30
void producer_main() {
```

```

    char* next_write_ptr;
33  while (!feof(stdin)){
        next_write_ptr = buffer_next(write_ptr);
        if (next_write_ptr != read_ptr) {
36            *write_ptr = getc(stdin);
            write_ptr = next_write_ptr;

39            struct timespec ts;
            ts.tv_sec = 0;
            ts.tv_nsec = 100000000ULL * rand() / RAND_MAX ;
42            nanosleep(&ts, NULL);
        }
    }
45 }

void* consumer_main(void* thread_id) {
48  int tid = (int) thread_id;
    char* current_rptr;
    char* current_wptr;
51  char data = 0;
    char* next_read_ptr;
    char* local_read_ptr;
54
    while (1) {
        local_read_ptr = read_ptr;
57
        if (local_read_ptr != write_ptr) {
            data = *local_read_ptr;
60            next_read_ptr = buffer_next(local_read_ptr);
            if (atomic_compare_exchange_weak(&read_ptr, &
                local_read_ptr, next_read_ptr)) {
                process(&data);
63            }
            struct timespec ts;
            ts.tv_sec = 0;
66            ts.tv_nsec = 100000000ULL * rand() / RAND_MAX ;
            nanosleep(&ts, NULL);
        }
69    }
}

72

```

```

    void create_consumers(int num_consumers) {
        pthread_t threads[num_consumers];
75     int thread_ids[num_consumers];
        int rc;
        for(int t = 0; t < num_consumers; t++){
78         thread_ids[t] = t;
            rc = pthread_create(&threads[t], NULL, consumer_main,
                (void* ) t);
            if (rc){
81         pthread_exit(NULL);
            }
        }
84 }

    int main() {
87     init_ring_buffer();

        create_consumers(2);
90     producer_main();

93     return 0;
    }

```

Listing 8.1: Code of the ringbuffer example programm

8.2 Instrumentation Benchmark Data

Tab. 8.1: Performance under different optimization levels without instrumentation

O3	O2	O1	O0
0.1161370277s	0.1191740036s	0.2948842049s	0.6530430317s
0.1157708168s	0.1156070232s	0.27384305s	0.6632189751s
0.1167349815s	0.1148910522s	0.2762970924s	0.6660399437s
0.1141610146s	0.1238379478s	0.2678158283s	0.6449968815s
0.1164500713s	0.1141080856s	0.2763991356s	0.6779379845s
0.1149349213s	0.1187281609s	0.2759730816s	0.6569571495s
0.1200730801s	0.1117370129s	0.2793321609s	0.6452558041s
0.1100320816s	0.1159410477s	0.2894630432s	0.6577329636s
0.1103510857s	0.1151809692s	0.2816138268s	0.656512022s
0.1138319969s	0.1219031811s	0.2807168961s	0.6649181843s
0.1163628101s	0.1176950932s	0.2723140717s	0.6588671207s
0.1170480251s	0.1168220043s	0.2781991959s	0.6618740559s
0.1127979755s	0.1200008392s	0.2799890041s	0.6571362019s
0.1124658585s	0.1124398708s	0.2763769627s	0.6544640064s
0.1231520176s	0.1198968887s	0.2752838135s	0.6740379333s
0.1203160286s	0.1191551685s	0.2750740051s	0.6417598724s
0.1163609028s	0.1229681969s	0.2662689686s	0.6602330208s
0.1186778545s	0.1143159866s	0.2685620785s	0.6608181s
0.1182408333s	0.1160438061s	0.2935631275s	0.6540088654s
0.1166889668s	0.112831831s	0.2743690014s	0.6557939053s
0.119601965s	0.1137580872s	0.2758669853s	0.6767961979s
0.1156849861s	0.1164021492s	0.2886178493s	0.6637132168s
0.1203570366s	0.1166651249s	0.2829039097s	0.6816310883s
0.1173582077s	0.1203048229s	0.2849731445s	0.6397099495s
0.116024971s	0.1158750057s	0.2814619541s	0.6493239403s
0.1161761284s	0.1233420372s	0.2841758728s	0.6558029652s
0.1086490154s	0.1245141029s	0.2797939777s	0.6555378437s
0.1138908863s	0.1158800125s	0.2746789455s	0.6605989933s
0.1159329414s	0.11946702s	0.2767148018s	0.6677999496s
0.123156786s	0.1197969913s	0.2818350792s	0.6489400864s
0.1154050827s	0.1125380993s	0.2842049599s	0.652766943s
0.114705801s	0.1216220856s	0.2762238979s	0.6488559246s
0.1205980778s	0.1136319637s	0.2732069492s	0.6657938957s
0.1184470654s	0.1135380268s	0.2743830681s	0.6611270905s
0.114989996s	0.1160199642s	0.2773461342s	0.6637861729s
0.1183080673s	0.1157469749s	0.2722430229s	0.6638269424s
0.1138470173s	0.1105167866s	0.2787780762s	0.673830986s
0.1113619804s	0.1144788265s	0.2707090378s	0.6705200672s
0.117235899s	0.1190290451s	0.2812240124s	0.638448s
0.1177809238s	0.1132018566s	0.2821011543s	0.6629910469s
0.1113541126s	0.1121530533s	0.2777709961s	0.6451292038s
0.1174621582s	0.1164631844s	0.2755610943s	0.6338021755s
0.1120979786s	0.1176538467s	0.2741901875s	0.6301341057s
0.1208879948s	0.1189939976s	0.2752621174s	0.6448199749s
0.1128540039s	0.1147179604s	0.283716917s	0.6382200718s
0.1172318459s	0.1122090816s	0.2783620358s	0.6650280952s
0.116394043s	0.1143491268s	0.2810368538s	0.6535108089s
0.1208820343s	0.1168520451s	0.2765209675s	0.660145998s
0.1193039417s	0.1131880283s	0.2831139565s	0.6619808674s
0.1179618835s	0.1127560139s	0.2710981369s	0.6434488297s

Tab. 8.2: Performance under different optimization levels with instrumentation

O3	O2	O1	O0
3.3911721706s	3.243653059s	3.421694994s	3.869437933s
3.5288460255s	3.6410529613s	3.319204092s	3.685021162s
3.3494801521s	3.4974589348s	3.3641340733s	3.6085999012s
3.5640091896s	3.4832391739s	3.5960211754s	3.7213051319s
3.3368079662s	3.4711410999s	3.4086368084s	3.6712110043s
3.5097038746s	3.5463349819s	3.4823799133s	3.6867759228s
3.3817141056s	3.3919298649s	3.4265608788s	3.8107779026s
3.6170189381s	3.4889719486s	3.5831069946s	3.6996588707s
3.376434803s	3.340255022s	3.4814970493s	3.7145421505s
3.4867889881s	3.7752280235s	3.5054810047s	3.7900021076s
3.3393919468s	3.3320212364s	3.6760630608s	3.7082140446s
3.2769758701s	3.35531497s	3.7271540165s	3.643638134s
3.5930538177s	3.3282749653s	3.4532988071s	3.6124441624s
3.5464639664s	3.5488898754s	3.559005022s	3.7060768604s
3.6068668365s	3.529253006s	3.4408950806s	3.6440010071s
3.3833189011s	3.3312380314s	3.64407897s	3.6952431202s
4.644769907s	3.2691669464s	3.6565649509s	3.7749199867s
3.9256081581s	3.3252110481s	3.530919075s	3.6850361824s
3.3639609814s	3.3914599419s	3.6782660484s	3.6187119484s
3.6016099453s	3.3473889828s	4.5602049828s	3.6672179699s
3.2825620174s	3.3359549046s	4.2268750668s	3.598236084s
3.1969110966s	3.3892168999s	4.3686730862s	3.6702840328s
3.3095178604s	3.3870069981s	3.5713601112s	3.7751419544s
3.4130480289s	3.4193899632s	3.679762125s	3.6604380608s
3.5319731236s	3.3532409668s	3.5412201881s	3.682322979s
3.3935861588s	3.325330019s	3.5808689594s	3.6570858955s
3.3294699192s	3.3051581383s	3.4674761295s	3.6164281368s
3.453649044s	3.3509070873s	3.521504879s	3.6644599438s
3.4952099323s	3.4388210773s	3.8695850372s	3.7006070614s
3.4600970745s	3.3684821129s	4.1289849281s	3.6317489147s
3.5761289597s	3.4953501225s	4.551377058s	3.6137089729s
3.5816121101s	3.4220538139s	4.4331259727s	3.7160301208s
3.6319949627s	4.3748400211s	3.6797399521s	3.6431188583s
3.4037461281s	3.7254831791s	3.5498430729s	3.633731842s
3.5317900181s	3.3875660896s	3.9773128033s	4.0382192135s
3.6149249077s	3.2542829514s	3.6003420353s	3.6492390633s
3.3683440685s	3.1844670773s	3.6274878979s	3.6503920555s
3.3448040485s	3.3463819027s	3.6728591919s	3.6539871693s
3.3042678833s	3.2252969742s	3.4163911343s	3.6895589828s
3.5786819458s	3.2732980251s	3.5556731224s	3.602257967s
3.5094149113s	3.431540966s	3.9204399586s	3.8089568615s
3.4041252136s	3.4122869968s	4.2923779488s	3.8015639782s
3.4141039848s	3.421681881s	4.0719110966s	3.7709450722s
3.4603841305s	3.390556097s	4.224946022s	3.693472147s
3.588504076s	4.0364210606s	3.7718069553s	3.6705441475s
3.6392641068s	3.3713409901s	3.5621869564s	3.7623281479s
3.4661819935s	3.3865299225s	3.4697358608s	3.6896250248s
3.3992631435s	3.3741409779s	3.5135061741s	3.7174310684s
3.3010737896s	3.2854759693s	3.4962620735s	3.6783411503s
4.2453198433s	3.3602261543s	3.5534739494s	3.6879298687s

Tab. 8.3: Performance impact of the amount of instrumented function calls

	10%	5%	2%	1%	0.5%	0.1%	0.01%	0.001%
32.27447295	16.57465911	6.818274021	3.465317011	1.921619892	0.529150963	0.301460028	0.126204967	
32.01804805	17.24028993	6.727348089	4.798867226	1.852323055	0.514367819	0.278471947	0.124331236	
32.56871295	16.59262991	6.958166122	4.0173738	1.757158041	0.494447947	0.286482096	0.124565125	
31.72771811	16.43900895	6.566558123	3.932744026	1.735903978	0.489733934	0.275008202	0.120963097	
33.2172401	16.28655887	6.669729948	4.31479001	1.771286011	0.474808931	0.275484085	0.126640081	
32.08020306	16.41876602	6.685636044	3.900309086	1.766814947	0.455327988	0.271858931	0.131477833	
36.09543991	16.71559906	6.687566996	3.842667103	1.820206881	0.443207026	0.271209955	0.130152941	
36.94550681	16.78408003	6.593606949	4.249938011	1.746952057	0.466243029	0.27235198	0.120084047	
33.03396583	20.1691699	7.4481709	3.456017971	1.839415073	0.452888966	0.315196991	0.117549896	
35.22510219	19.88422704	7.280730963	3.801725864	1.772809029	0.446066856	0.33523798	0.118908167	
31.87948489	20.26601911	7.06736207	5.228184938	1.697596073	0.434090853	0.273921013	0.12099719	
32.9260869	19.230762	6.763454914	4.397490978	1.70266819	0.444143057	0.270895004	0.119578123	
31.93172407	17.89269495	6.711801052	3.741932869	1.754426956	0.436100006	0.275285959	0.121613026	
32.63784504	17.51772094	7.74854207	3.387830019	1.848978043	0.502735138	0.276705027	0.124044895	
32.17112899	16.38639498	6.714694023	3.371681929	2.038882017	0.438040972	0.289375067	0.126643896	
32.49857807	16.44085193	6.677272797	4.135891914	1.871739149	0.434919119	0.266953945	0.121059895	
32.13873911	16.41528106	7.115194082	3.2681849	1.825345039	0.449182034	0.271996975	0.121709108	
32.00047112	16.13716006	6.771155834	3.623764992	1.748323917	0.427417994	0.267868996	0.121583939	
32.55476999	16.43411779	7.221987963	3.762748957	1.767467022	0.455560207	0.265871048	0.12225008	
32.12196398	16.15078902	6.738729954	3.452557802	1.718863964	0.437258005	0.274132013	0.12803793	
32.70064616	16.28547597	6.58794713	3.425955057	1.727346182	0.444720984	0.273182869	0.120798826	
32.07056212	15.99900794	6.692715883	3.689862013	1.774134159	0.44466114	0.272839069	0.127115011	

Tab. 8.3: Performance impact of the amount of instrumented function calls

	10%	5%	2%	1%	0.5%	0.1%	0.01%	0.001%
	32.55583501	16.28823805	6.733635902	3.283070087	1.882145882	0.439930916	0.2744446964	0.124543905
	35.26152492	17.62537789	6.575346947	3.366142988	1.834486008	0.465766907	0.274556875	0.127908945
	46.2991581	16.66089106	6.660341024	3.311545134	1.755493164	0.443130016	0.267874002	0.126642942
	34.31453204	16.72705483	6.722349882	3.527065992	1.714565039	0.434679985	0.27280283	0.125182152
	33.36741495	15.95860291	6.73634696	3.391783953	1.731849909	0.442686081	0.27355504	0.125731945
	35.32494783	15.89461899	7.181751966	3.281201124	1.736423016	0.425590992	0.274189949	0.12746501
	37.58029795	17.04383206	8.083317041	3.951409101	1.750784159	0.461474895	0.270134926	0.125010967
	35.69649506	17.98182082	6.625492811	3.431293011	1.800821066	0.451671839	0.269315958	0.118397951
	32.39285803	16.21179414	6.822492123	3.436654091	1.785053968	0.445170164	0.271719933	0.121504068
	34.68371701	16.111503792	8.079453945	3.431628942	1.773240089	0.437523842	0.280767918	0.130489111
	32.91161203	16.68876195	7.799345016	3.613914967	1.791203976	0.446851969	0.268200874	0.124853134
	32.29311299	17.47719908	6.651239872	3.469138145	1.955821991	0.433550119	0.271760941	0.126969814
	32.55763197	16.24372602	6.682399035	3.37647891	1.843163967	0.446904182	0.271901131	0.127385139
	32.41566491	16.11199789	6.683773041	3.348512888	1.845180035	0.437231064	0.279286861	0.123465061
	32.73225403	16.19021702	6.65988493	3.529052973	1.827486038	0.458367825	0.27441597	0.125228882
	32.73954201	16.48861504	6.799346924	3.359493971	1.821255922	0.429441929	0.270982027	0.12332201
	32.71354008	16.07665801	6.868185997	3.427947998	1.838265896	0.446797132	0.274680853	0.123304844
	34.54573703	16.18055511	6.692906857	3.354235888	1.75396204	0.439625025	0.270151138	0.120908022
	32.20775294	17.64140296	6.859113932	3.479537964	1.78051281	0.451714993	0.277094126	0.129561186
	32.24934101	16.380584	6.676775932	3.383896112	1.789362192	0.447572947	0.295809984	0.17267108
	33.114712	16.323807	7.262549877	3.366688013	1.696341038	0.446156025	0.270586967	0.117249966
	32.15380812	16.76982903	7.166394949	3.322533846	1.81302619	0.440481901	0.270737886	0.129595041

Tab. 8.3: Performance impact of the amount of instrumented function calls

	10%	5%	2%	1%	0.5%	0.1%	0.01%	0.001%
	32.50190592	16.2511122	7.670663118	3.395035982	1.788846016	0.452170134	0.270417929	0.123901129
	32.13687491	16.51295781	6.750161171	3.364979982	1.706182957	0.454016924	0.267800093	0.119033098
	32.58365011	16.10383296	6.636216879	3.416320801	1.729248047	0.443239927	0.271872044	0.12143898
	34.04416585	16.26455212	6.655556202	3.349573135	1.702176094	0.448850155	0.266824961	0.12451005
	37.44555593	19.14247894	6.996401072	3.390678167	1.888344049	0.446580172	0.270505905	0.125029087
	35.42529082	20.047014	7.087264061	3.351912975	1.901141882	0.455142021	0.264132023	0.123822927

Bibliography

- [1]Rajeev Alur and D.L. Dill. „A theory of timed automata“. In: *Theoretical computer science* 126.2 (1994), pp. 183–235 (cit. on p. 16).
- [2]Ben D’Angelo, Sriram Sankaranarayanan, César Sánchez, et al. „LOLA: Runtime monitoring of synchronous systems“. In: *Proceedings of the International Workshop on Temporal Representation and Reasoning* (2005), pp. 166–175 (cit. on pp. 1, 2, 6).
- [3]Normann Decker, Daniel Thoma, and Jannis Harder. „TESSLA A Temporal Stream-based Specification Language“. 2016 (cit. on pp. 5, 6).
- [4]Klaus Havelund. „Runtime Verification of C Programs“. In: (2008) (cit. on pp. 1, 5, 8).
- [5]Menna Mostafa and Borzoo Bonakdarpour. „Decentralized Runtime Verification of LTL Specifications in Distributed Systems“. In: *2015 IEEE International Parallel and Distributed Processing Symposium* (2015), pp. 494–503 (cit. on pp. 1, 7).
- [6]George C. Necula, Scott McPeak, Shree P. Rahul, and Westley Weimer. „CIL: Intermediate language and tools for analysis and transformation of C programs“. In: *Lecture Notes in Computer Science (including subseries Lecture Notes in Artificial Intelligence and Lecture Notes in Bioinformatics)* 2304 (2002), pp. 213–228 (cit. on p. 8).
- [7]Lee Pike, Alwyn Goodloe, Robin Morisset, and Sebastian Niller. „Copilot: A hard real-time runtime monitor“. In: *Lecture Notes in Computer Science (including subseries Lecture Notes in Artificial Intelligence and Lecture Notes in Bioinformatics)* 6418 LNCS.Rv (2010), pp. 345–359 (cit. on pp. 1, 7, 8).
- [8]Margaret H. Smith and Klaus Havelund. „Requirements capture with RCAT“. In: *Proceedings of the 16th IEEE International Requirements Engineering Conference, RE’08* (2008), pp. 183–192 (cit. on p. 8).
- [9]Xi Zheng, Christine Julien, Rodion Podorozhny, and Franck Cassez. „BraceAssertion: Runtime verification of cyber-physical systems“. In: *Proceedings - 2015 IEEE 12th International Conference on Mobile Ad Hoc and Sensor Systems, MASS 2015* (2015), pp. 298–306 (cit. on p. 1).

List of Figures

2.1	Visualization of TESSLA stream model, taken from [3]	6
4.1	Influence of the order of independent transitions on the global state of an evaluation engine	35
5.1	Visualization of a simple evaluation engine with a greedy schedule.	42
5.2	Visualization of three runs of an evaluation engine as explained in Proof 5.2.1.	46
6.1	Performance of an example C program with and without instrumentation	56
6.2	Execution time of a program with different amount of calls to an instrumented function	57

List of Tables

4.1	List of complete functions	26
4.2	List of output complete functions	27
4.3	List of input complete functions	27
4.4	List of incomplete functions	29
5.1	Example how the behaviour of a run is constructed	40
8.1	Performance under different optimization levels without instrumentation	64
8.2	Performance under different optimization levels with instrumentation .	65
8.3	Performance impact of the amount of instrumented function calls . . .	66

Glossary

LOLA A specification language and algorithms for the online and offline monitoring of synchronous systems including circuits and embedded systems. 1, 6, 7

RCAT Requirements CApture Tool. 8

RMOR Requirement Monitoring and Recovery. 1, 8

TESSLA A temporal, stream based specification Language. 1, 2, 5–9, 11, 13, 21–24, 29, 32, 34, 40, 45, 51–53, 55, 57, 71

DAG Directed Acyclic Graph. 23, 40–42, 44, 49

FIFO First In First Out. 23

ISP Institute for Software Engineering and Programming Languages. 1

RV Runtime Verification. 1, 2, 7, 8

TL Temporal Logic. 1, 6

List of Theorems

1	Definition (Transducer)	13
2	Definition (Deterministic Transducer)	13
3	Definition (Synchronous Transducer)	14
4	Definition (Asynchronous Transducer)	14
5	Definition (Causal and Clairvoyant Transducers)	14
6	Definition (Asynchronous equivalence of Transducers)	14
1	Lemma (Asynchronous equivalence is an equivalence Relation) . . .	15
7	Definition (Timed Sequence)	16
8	Definition (Monotonicity of Timed Sequences)	17
9	Definition (Timed Transducer)	17
10	Definition (Boundedness of Timed Transducers)	17
11	Definition (Observational Equivalence)	18
2	Lemma (Observational Equivalence is an Equivalence Relationship for Bounded Transducers)	18
12	Definition (Application of a Transition on a State)	30
13	Definition (Closeness of Runs)	31
14	Definition (Enabledness of a Node)	32
15	Definition (Valid Run)	32
16	Definition (Independence of Nodes)	34
17	Definition (Independence of Transitions)	34
3	Lemma (Exchange of Independent Transitions)	34
4	Lemma (Duration of Enabledness)	35
5	Lemma (Finiteness of Enabledness)	36
18	Definition (Fair Schedules)	37
19	Definition (Behaviour of a Run)	39
20	Definition (Equivalence of Runs)	40
21	Definition (Equivalence of Evaluation Engines)	40
22	Definition (Greedy schedule)	41
23	Definition (Valid Evaluation Engines)	42
6	Lemma (Greedy Schedules are Fair)	43
24	Definition (Fair Evaluation Engines)	43
1	Theorem (Equivalence of Different Greedy Evaluation Engines) . . .	45
2	Theorem (Equivalence of Fair and Greedy Evaluation Engines) . . .	50

